
Amazon EC2 Auto Scaling

User Guide



Amazon EC2 Auto Scaling: User Guide

Copyright © 2019 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

The AWS Documentation website is getting a new look!

Try it now and let us know what you think. [Switch to the new look >>](#)

You can return to the original look by selecting English in the language selector above.

Table of Contents

What Is Amazon EC2 Auto Scaling?	1
Auto Scaling Components	1
Getting Started	2
Accessing Amazon EC2 Auto Scaling	2
Pricing for Amazon EC2 Auto Scaling	3
PCI DSS Compliance	3
Related Services	3
Benefits of Auto Scaling	3
Example: Covering Variable Demand	4
Example: Web App Architecture	5
Example: Distributing Instances Across Availability Zones	6
Auto Scaling Lifecycle	7
Scale Out	8
Instances In Service	8
Scale In	8
Attach an Instance	9
Detach an Instance	9
Lifecycle Hooks	9
Enter and Exit Standby	9
Amazon EC2 Auto Scaling Limits	9
Setting Up	11
Sign Up for AWS	11
Prepare to Use Amazon EC2	11
Getting Started	12
Step 1: Create a Launch Template	12
Step 2: Create an Auto Scaling Group	14
Step 3: Verify Your Auto Scaling Group	15
(Optional) Terminate an Instance in Your Auto Scaling Group	16
Step 4: (Optional) Delete Your Scaling Infrastructure	17
Tutorial: Set Up a Scaled and Load-Balanced Application	18
Prerequisites	18
Configure Scaling and Load Balancing (Console)	18
Create or Select a Launch Configuration	19
Create or Select a Launch Template	20
Create an Auto Scaling Group	21
(Optional) Verify That Your Load Balancer Is Attached to Your Auto Scaling Group	21
Configure Scaling and Load Balancing (AWS CLI)	22
Create a Launch Configuration	22
Create a Launch Template	22
Create an Auto Scaling Group with a Load Balancer	22
Cleaning Up Your AWS Resources	23
Launch Templates	24
Creating a Launch Template for an Auto Scaling Group	24
Copying a Launch Configuration to a Launch Template	29
Replacing a Launch Configuration with a Launch Template	30
Launch Configurations	32
Creating a Launch Configuration	32
Creating a Launch Configuration Using an EC2 Instance	33
Create a Launch Configuration Using an EC2 Instance	34
Create a Launch Configuration from an Instance and Override the Block Devices (AWS CLI)	35
Create a Launch Configuration and Override the Instance Type (AWS CLI)	36
Changing a Launch Configuration	37
Launching Auto Scaling Instances in a VPC	38
Default VPC	38

IP Addressing in a VPC	39
Instance Placement Tenancy	39
Linking EC2-Classic Instances to a VPC	40
Auto Scaling Groups	42
Using Multiple Instance Types and Purchase Options	43
Allocation Strategies	43
Controlling the Proportion of On-Demand Instances	44
Best Practices for Spot Instances	46
Prerequisites	46
Creating an Auto Scaling Group with Multiple Purchase Options	46
Creating a Group Using a Launch Template	51
Creating a Group Using a Launch Configuration	53
Creating a Group Using an EC2 Instance	54
Create an Auto Scaling Group from an EC2 Instance (Console)	55
Create an Auto Scaling Group from an EC2 Instance (AWS CLI)	55
Creating a Group Using the Launch Wizard	56
Tagging Auto Scaling Groups and Instances	57
Tag Restrictions	58
Tagging Lifecycle	58
Add or Modify Tags for Your Auto Scaling Group	58
Delete Tags	61
Using a Load Balancer with an Auto Scaling Group	61
Elastic Load Balancing Types	62
Attaching a Load Balancer	63
Adding ELB Health Checks	65
Adding an Availability Zone	66
Launching Spot Instances in Your Auto Scaling Group	68
Merging Auto Scaling Groups	69
Merge Zones (AWS CLI)	69
Deleting Your Auto Scaling Infrastructure	70
Delete Your Auto Scaling Group	71
(Optional) Delete the Launch Configuration	71
(Optional) Delete the Launch Template	71
(Optional) Delete the Load Balancer	72
(Optional) Delete CloudWatch Alarms	73
Scaling Your Group	74
Scaling Options	74
Maintaining the Size of Your Auto Scaling Group	75
Manual Scaling for Amazon EC2 Auto Scaling	75
Changing the Size of Your Auto Scaling Group (Console)	76
Changing the Size of Your Auto Scaling Group (AWS CLI)	76
Attach EC2 Instances to Your Auto Scaling Group	78
Detach EC2 Instances from Your Auto Scaling Group	82
Scheduled Scaling	85
Considerations	85
Create and Manage Scheduled Actions (Console)	85
Create and Manage Scheduled Actions (AWS CLI)	86
Dynamic Scaling	87
Scaling Policy Types	88
Multiple Scaling Policies	88
Target Tracking Scaling Policies	89
Simple and Step Scaling Policies	93
Add a Scaling Policy to an Existing Auto Scaling Group	101
Scaling Based on Amazon SQS	102
Deleting a Scaling Policy	106
Scaling Cooldowns	107
Example: Cooldowns	107

Default Cooldowns	108
Scaling-Specific Cooldowns	109
Cooldowns and Multiple Instances	109
Cooldowns and Lifecycle Hooks	109
Auto Scaling Instance Termination	110
Default Termination Policy	110
Customizing the Termination Policy	112
Instance Protection	113
Lifecycle Hooks	115
How Lifecycle Hooks Work	116
Considerations When Using Lifecycle Hooks	117
Prepare for Notifications	118
Add Lifecycle Hooks	121
Complete the Lifecycle Hook	122
Test the Notification	122
Temporarily Removing Instances	123
How the Standby State Works	123
Health Status of an Instance in a Standby State	124
Temporarily Remove an Instance (Console)	124
Temporarily Remove an Instance (AWS CLI)	124
Suspending Scaling	127
Scaling Processes	127
Choosing to Suspend	128
Suspend and Resume Scaling Processes (Console)	130
Suspend and Resume Scaling Processes (AWS CLI)	130
Monitoring Your Auto Scaling Instances and Groups	132
Health Checks	132
Instance Health Status	133
Determining Instance Health	133
Health Check Grace Period	133
Replacing Unhealthy Instances	134
Custom Health Checks	134
Amazon CloudWatch Metrics	135
Auto Scaling Group Metrics	135
Dimensions for Auto Scaling Group Metrics	136
Enable Auto Scaling Group Metrics	136
Configure Monitoring for Auto Scaling Instances	137
View CloudWatch Metrics	138
Create Amazon CloudWatch Alarms	139
Amazon CloudWatch Events	140
Auto Scaling Events	140
Create a Lambda Function	144
Route Events to Your Lambda Function	144
Amazon SNS Notifications	145
SNS Notifications	146
Configure Amazon SNS	147
Configure Your Auto Scaling Group to Send Notifications	147
Test the Notification Configuration	148
Verify That You Received Notification of the Scaling Event	148
Delete the Notification Configuration	150
AWS CloudTrail Logging	150
Amazon EC2 Auto Scaling Information in CloudTrail	150
Understanding Amazon EC2 Auto Scaling Log File Entries	151
Controlling Access to Your Resources	153
Specifying Actions in a Policy	153
Additional IAM Permissions	154
Specifying the Resource	154

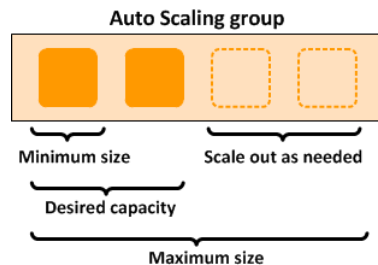
Specifying Conditions in a Policy	156
Predefined AWS Managed Policies	156
Customer Managed Policy Examples	157
Example: Restricting Which Service-Linked Role Can Be Passed (Using PassRole)	157
Example: Require a Launch Template	158
Example: Create and Manage Launch Configurations	158
Example: Create and Manage Auto Scaling Groups and Scaling Policies	159
Example: Control Access Using Tags	160
Example: Change the Capacity of Auto Scaling Groups	162
Service-Linked Roles	163
Overview	163
Permissions Granted by the Service-Linked Role	164
Create a Service-Linked Role (Automatic)	164
Create a Service-Linked Role (Manual)	164
Edit the Service-Linked Role	165
Delete the Service-Linked Role	166
Supported Regions for Amazon EC2 Auto Scaling Service-Linked Roles	166
Required CMK Key Policy for Use with Encrypted Volumes	166
Configuring Key Policies	166
Example: CMK Key Policy Sections That Allow Access to the CMK	167
Example: CMK Key Policy Sections That Allow Cross-Account Access to the CMK	168
IAM Role for Applications That Run on Amazon EC2 Instances	169
Prerequisites	170
Create a Launch Configuration	170
Create a Launch Template	170
Troubleshooting	172
Retrieving an Error Message	172
Instance Launch Failure	174
The security group <name of the security group> does not exist. Launching EC2 instance failed.	175
The key pair <key pair associated with your EC2 instance> does not exist. Launching EC2 instance failed.	175
The requested configuration is currently not supported.	175
AutoScalingGroup <Auto Scaling group name> not found.	175
The requested Availability Zone is no longer supported. Please retry your request... ..	176
Your requested instance type (<instance type>) is not supported in your requested Availability Zone (<instance Availability Zone>)... ..	176
You are not subscribed to this service. Please see https://aws.amazon.com/	176
Invalid device name upload. Launching EC2 instance failed.	176
Value (<name associated with the instance storage device>) for parameter virtualName is invalid... ..	177
EBS block device mappings not supported for instance-store AMIs.	177
Placement groups may not be used with instances of type 'm1.large'. Launching EC2 instance failed.	177
Client.InternalError: Client error on launch.	177
AMI Issues	178
The AMI ID <ID of your AMI> does not exist. Launching EC2 instance failed.	179
AMI <AMI ID> is pending, and cannot be run. Launching EC2 instance failed.	179
Value (<ami ID>) for parameter virtualName is invalid.	179
The requested instance type's architecture (i386) does not match the architecture in the manifest for ami-6622f00f (x86_64). Launching ec2 instance failed.	179
Load Balancer Issues	180
Cannot find Load Balancer <your launch environment>. Validating load balancer configuration failed.	180
There is no ACTIVE Load Balancer named <load balancer name>. Updating load balancer configuration failed.	180
EC2 instance <instance ID> is not in VPC. Updating load balancer configuration failed.	181
EC2 instance <instance ID> is in VPC. Updating load balancer configuration failed.	181

The security token included in the request is invalid. Validating load balancer configuration failed.	181
Capacity Limits	181
We currently do not have sufficient <instance type> capacity in the Availability Zone you requested (<requested Availability Zone>)... ..	181
<number of instances> instance(s) are already running. Launching EC2 instance failed.	182
Resources	183
Document History	184

What Is Amazon EC2 Auto Scaling?

Amazon EC2 Auto Scaling helps you ensure that you have the correct number of Amazon EC2 instances available to handle the load for your application. You create collections of EC2 instances, called *Auto Scaling groups*. You can specify the minimum number of instances in each Auto Scaling group, and Amazon EC2 Auto Scaling ensures that your group never goes below this size. You can specify the maximum number of instances in each Auto Scaling group, and Amazon EC2 Auto Scaling ensures that your group never goes above this size. If you specify the desired capacity, either when you create the group or at any time thereafter, Amazon EC2 Auto Scaling ensures that your group has this many instances. If you specify scaling policies, then Amazon EC2 Auto Scaling can launch or terminate instances as demand on your application increases or decreases.



For example, the following Auto Scaling group has a minimum size of one instance, a desired capacity of two instances, and a maximum size of four instances. The scaling policies that you define adjust the number of instances, within your minimum and maximum number of instances, based on the criteria that you specify.



For more information about the benefits of Amazon EC2 Auto Scaling, see [Benefits of Auto Scaling](#) (p. 3).

Auto Scaling Components

The following table describes the key components of Amazon EC2 Auto Scaling.

	<p>Groups</p> <p>Your EC2 instances are organized into <i>groups</i> so that they can be treated as a logical unit for the purposes of scaling and management. When you create a group, you can specify its minimum, maximum, and, desired number of EC2 instances. For more information, see Auto Scaling Groups (p. 42).</p>
	<p>Configuration templates</p> <p>Your group uses a <i>launch template</i> or a <i>launch configuration</i> as a configuration template for its EC2 instances. You can specify information such as the AMI ID, instance type, key pair, security groups, and block device mapping for your instances. For more information, see Launch Templates (p. 24) and Launch Configurations (p. 32).</p>



Scaling options

Amazon EC2 Auto Scaling provides several ways for you to scale your Auto Scaling groups. For example, you can configure a group to scale based on the occurrence of specified conditions (dynamic scaling) or on a schedule. For more information, see [Scaling Options](#) (p. 74).

Getting Started

If you're new to Amazon EC2 Auto Scaling, we recommend that you review [Auto Scaling Lifecycle](#) (p. 7) before you begin.

To begin, complete the [Getting Started with Amazon EC2 Auto Scaling](#) (p. 12) tutorial to create an Auto Scaling group and see how it responds when an instance in that group terminates. If you already have running EC2 instances, you can create an Auto Scaling group using an existing EC2 instance, and remove the instance from the group at any time.

Accessing Amazon EC2 Auto Scaling

If you've signed up for an AWS account, you can access Amazon EC2 Auto Scaling by signing into the AWS Management Console, choosing **EC2** from the console home page, and then choosing **Auto Scaling Groups** from the navigation pane.

You can also access Amazon EC2 Auto Scaling using the [Amazon EC2 Auto Scaling API](#). Amazon EC2 Auto Scaling provides a Query API. These requests are HTTP or HTTPS requests that use the HTTP verbs GET or POST and a Query parameter named *Action*. For more information about the API actions for Amazon EC2 Auto Scaling, see [Actions](#) in the *Amazon EC2 Auto Scaling API Reference*.

If you prefer to build applications using language-specific APIs instead of submitting a request over HTTP or HTTPS, AWS provides libraries, sample code, tutorials, and other resources for software developers. These libraries provide basic functions that automate tasks such as cryptographically signing your requests, retrying requests, and handling error responses, making it is easier for you to get started. For more information, see [AWS SDKs and Tools](#).

If you prefer to use a command line interface, you have the following options:

AWS Command Line Interface (CLI)

Provides commands for a broad set of AWS products, and is supported on Windows, macOS, and Linux. To get started, see [AWS Command Line Interface User Guide](#). For more information, see [autoscaling](#) in the *AWS CLI Command Reference*.

AWS Tools for Windows PowerShell

Provides commands for a broad set of AWS products for those who script in the PowerShell environment. To get started, see the [AWS Tools for Windows PowerShell User Guide](#). For more information, see the [AWS Tools for PowerShell Cmdlet Reference](#).

For information about your credentials for accessing AWS, see [AWS Security Credentials](#) in the *Amazon Web Services General Reference*. For information about regions and endpoints for calls to Amazon EC2 Auto Scaling, visit [AWS Regions and Endpoints](#) in the *AWS General Reference*.

Pricing for Amazon EC2 Auto Scaling

There are no additional fees with Amazon EC2 Auto Scaling, so it's easy to try it out and see how it can benefit your AWS architecture.

PCI DSS Compliance

Auto Scaling supports the processing, storage, and transmission of credit card data by a merchant or service provider, and has been validated as being compliant with Payment Card Industry (PCI) Data Security Standard (DSS). For more information about PCI DSS, including how to request a copy of the AWS PCI Compliance Package, see [PCI DSS Level 1](#).

Related Services

To configure automatic scaling for all of the scalable AWS resources for your application, use AWS Auto Scaling. With AWS Auto Scaling, you can also simplify the process of defining dynamic scaling policies for your Auto Scaling groups and use predictive scaling to scale your Amazon EC2 capacity in advance of predicted traffic changes. For more information, see the [AWS Auto Scaling User Guide](#).

To automatically distribute incoming application traffic across multiple instances in your Auto Scaling group, use Elastic Load Balancing. For more information, see the [Elastic Load Balancing User Guide](#).

To monitor basic statistics for your instances and Amazon EBS volumes, use Amazon CloudWatch. For more information, see the [Amazon CloudWatch User Guide](#).

To monitor the calls made to the Amazon EC2 Auto Scaling API for your account, use AWS CloudTrail. The data logged includes calls made by the AWS Management Console, command line tools, and other services. For more information, see the [AWS CloudTrail User Guide](#).

Benefits of Auto Scaling

Adding Amazon EC2 Auto Scaling to your application architecture is one way to maximize the benefits of the AWS Cloud. When you use Amazon EC2 Auto Scaling, your applications gain the following benefits:

- Better fault tolerance. Amazon EC2 Auto Scaling can detect when an instance is unhealthy, terminate it, and launch an instance to replace it. You can also configure Amazon EC2 Auto Scaling to use multiple Availability Zones. If one Availability Zone becomes unavailable, Amazon EC2 Auto Scaling can launch instances in another one to compensate.
- Better availability. Amazon EC2 Auto Scaling helps ensure that your application always has the right amount of capacity to handle the current traffic demand.
- Better cost management. Amazon EC2 Auto Scaling can dynamically increase and decrease capacity as needed. Because you pay for the EC2 instances you use, you save money by launching instances when they are needed and terminating them when they aren't.

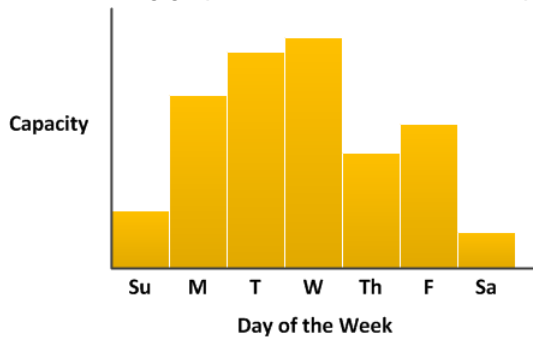
Contents

- [Example: Covering Variable Demand \(p. 4\)](#)
- [Example: Web App Architecture \(p. 5\)](#)
- [Example: Distributing Instances Across Availability Zones \(p. 6\)](#)

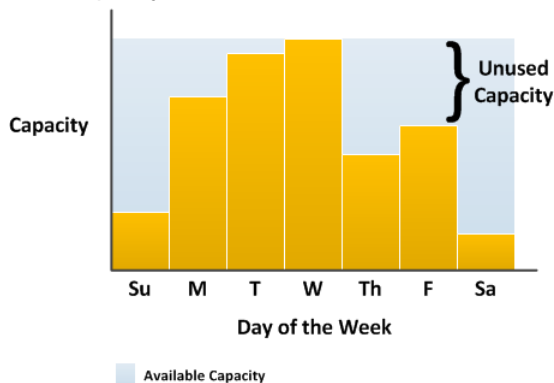
Example: Covering Variable Demand

To demonstrate some of the benefits of Amazon EC2 Auto Scaling, consider a basic web application running on AWS. This application allows employees to search for conference rooms that they might want to use for meetings. During the beginning and end of the week, usage of this application is minimal. During the middle of the week, more employees are scheduling meetings, so the demand on the application increases significantly.

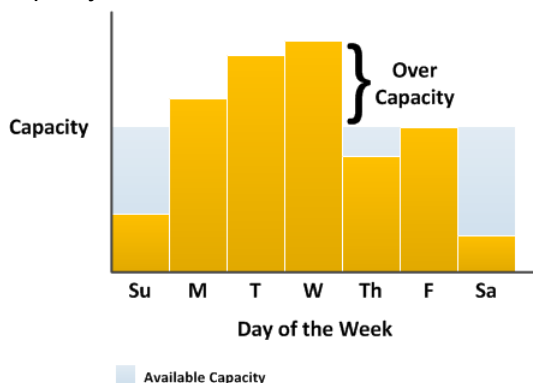
The following graph shows how much of the application's capacity is used over the course of a week.



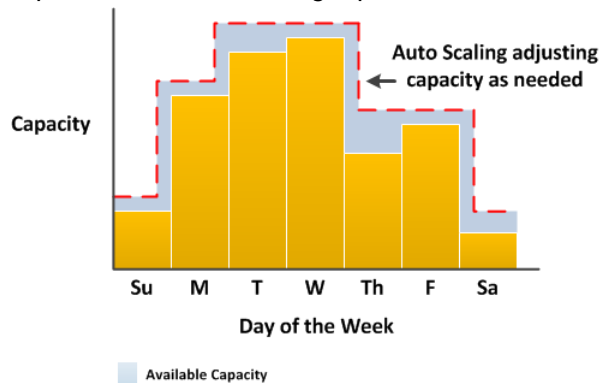
Traditionally, there are two ways to plan for these changes in capacity. The first option is to add enough servers so that the application always has enough capacity to meet demand. The downside of this option, however, is that there are days in which the application doesn't need this much capacity. The extra capacity remains unused and, in essence, raises the cost of keeping the application running.



The second option is to have enough capacity to handle the average demand on the application. This option is less expensive, because you aren't purchasing equipment that you'll only use occasionally. However, you risk creating a poor customer experience when the demand on the application exceeds its capacity.

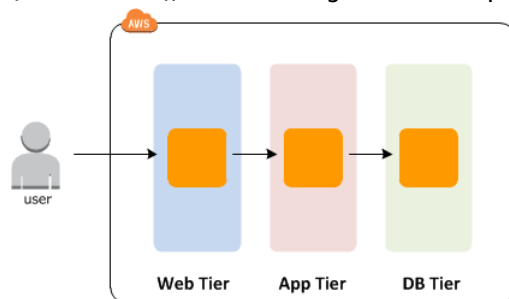


By adding Amazon EC2 Auto Scaling to this application, you have a third option available. You can add new instances to the application only when necessary, and terminate them when they're no longer needed. Because Amazon EC2 Auto Scaling uses EC2 instances, you only have to pay for the instances you use, when you use them. You now have a cost-effective architecture that provides the best customer experience while minimizing expenses.

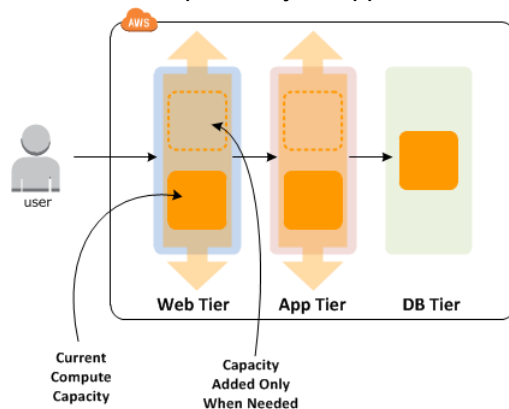


Example: Web App Architecture

In a common web app scenario, you run multiple copies of your app simultaneously to cover the volume of your customer traffic. These multiple copies of your application are hosted on identical EC2 instances (cloud servers), each handling customer requests.



Amazon EC2 Auto Scaling manages the launch and termination of these EC2 instances on your behalf. You define a set of criteria (such as an Amazon CloudWatch alarm) that determines when the Auto Scaling group launches or terminates EC2 instances. Adding Auto Scaling groups to your network architecture helps make your application more highly available and fault tolerant.



You can create as many Auto Scaling groups as you need. For example, you can create an Auto Scaling group for each tier.

To distribute traffic between the instances in your Auto Scaling groups, you can introduce a load balancer into your architecture. For more information, see [Using a Load Balancer with an Auto Scaling Group](#) (p. 61).

Example: Distributing Instances Across Availability Zones

AWS resources, such as EC2 instances, are housed in highly available data centers. To provide additional scalability and reliability, these data centers are in different physical locations. *Regions* are large and widely dispersed geographic locations. Each Region contains multiple distinct locations, called *Availability Zones*, which are engineered to be isolated from failures in other Availability Zones. They provide inexpensive, low-latency network connectivity to other Availability Zones in the same Region. For more information, see [Regions and Endpoints: Amazon EC2 Auto Scaling](#) in the *Amazon Web Services General Reference*.

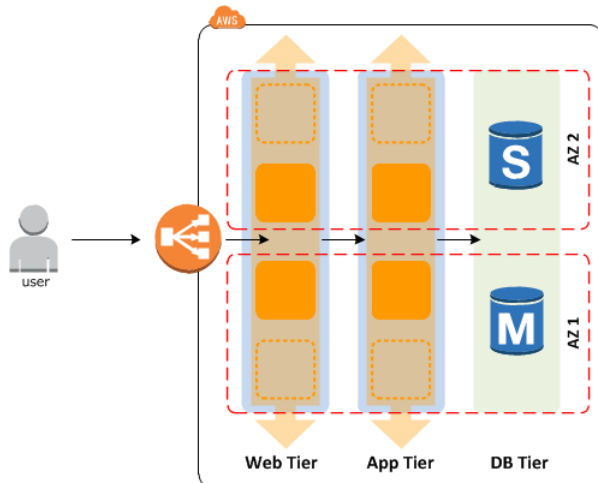
Amazon EC2 Auto Scaling enables you to take advantage of the safety and reliability of geographic redundancy by spanning Auto Scaling groups across multiple Availability Zones within a Region. When one Availability Zone becomes unhealthy or unavailable, Auto Scaling launches new instances in an unaffected Availability Zone. When the unhealthy Availability Zone returns to a healthy state, Auto Scaling automatically redistributes the application instances evenly across all of the designated Availability Zones.

An Auto Scaling group can contain EC2 instances in one or more Availability Zones within the same Region. However, Auto Scaling groups cannot span multiple Regions.

For Auto Scaling groups in a VPC, the EC2 instances are launched in subnets. You select the subnets for your EC2 instances when you create or update the Auto Scaling group. You can select one or more subnets per Availability Zone. For more information, see [VPCs and Subnets](#) in the *Amazon VPC User Guide*.

Instance Distribution

Amazon EC2 Auto Scaling attempts to distribute instances evenly between the Availability Zones that are enabled for your Auto Scaling group. Amazon EC2 Auto Scaling does this by attempting to launch new instances in the Availability Zone with the fewest instances. If the attempt fails, however, Amazon EC2 Auto Scaling attempts to launch the instances in another Availability Zone until it succeeds. For Auto Scaling groups in a VPC, if there are multiple subnets in an Availability Zone, Amazon EC2 Auto Scaling selects a subnet from the Availability Zone at random.



Rebalancing Activities

After certain actions occur, your Auto Scaling group can become unbalanced between Availability Zones. Amazon EC2 Auto Scaling compensates by rebalancing the Availability Zones. The following actions can lead to rebalancing activity:

- You change the Availability Zones for your group.
- You explicitly terminate or detach instances and the group becomes unbalanced.
- An Availability Zone that previously had insufficient capacity recovers and has additional capacity available.
- An Availability Zone that previously had a Spot price above your maximum price now has a Spot price below your maximum price.

When rebalancing, Amazon EC2 Auto Scaling launches new instances before terminating the old ones, so that rebalancing does not compromise the performance or availability of your application.

Because Amazon EC2 Auto Scaling attempts to launch new instances before terminating the old ones, being at or near the specified maximum capacity could impede or completely halt rebalancing activities. To avoid this problem, the system can temporarily exceed the specified maximum capacity of a group by a 10 percent margin (or by a 1-instance margin, whichever is greater) during a rebalancing activity. The margin is extended only if the group is at or near maximum capacity and needs rebalancing, either because of user-requested rezoning or to compensate for zone availability issues. The extension lasts only as long as needed to rebalance the group typically a few minutes.

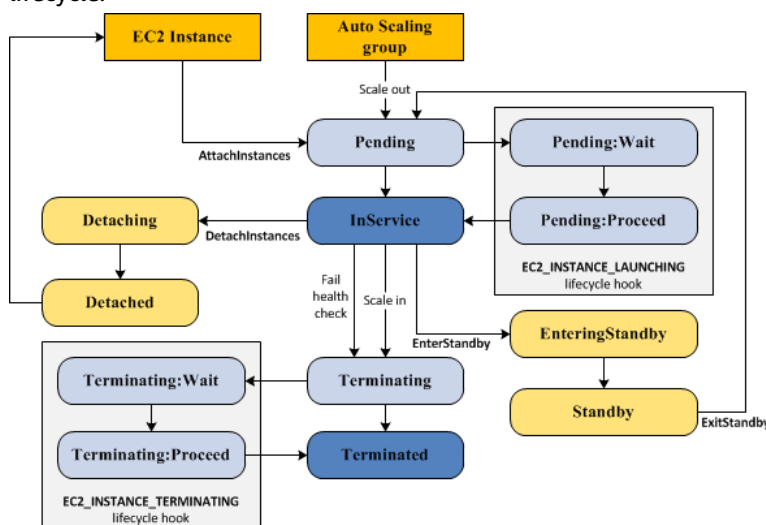
Auto Scaling Lifecycle

The EC2 instances in an Auto Scaling group have a path, or lifecycle, that differs from that of other EC2 instances. The lifecycle starts when the Auto Scaling group launches an instance and puts it into service. The lifecycle ends when you terminate the instance, or the Auto Scaling group takes the instance out of service and terminates it.

Note

You are billed for instances as soon as they are launched, including the time that they are not yet in service.

The following illustration shows the transitions between instance states in the Amazon EC2 Auto Scaling lifecycle.



Scale Out

The following scale out events direct the Auto Scaling group to launch EC2 instances and attach them to the group:

- You manually increase the size of the group. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#).
- You create a scaling policy to automatically increase the size of the group based on a specified increase in demand. For more information, see [Dynamic Scaling for Amazon EC2 Auto Scaling \(p. 87\)](#).
- You set up scaling by schedule to increase the size of the group at a specific time. For more information, see [Scheduled Scaling for Amazon EC2 Auto Scaling \(p. 85\)](#).

When a scale out event occurs, the Auto Scaling group launches the required number of EC2 instances, using its assigned launch configuration. These instances start in the `Pending` state. If you add a lifecycle hook to your Auto Scaling group, you can perform a custom action here. For more information, see [Lifecycle Hooks \(p. 9\)](#).

When each instance is fully configured and passes the Amazon EC2 health checks, it is attached to the Auto Scaling group and it enters the `InService` state. The instance is counted against the desired capacity of the Auto Scaling group.

Instances In Service

Instances remain in the `InService` state until one of the following occurs:

- A scale in event occurs, and Amazon EC2 Auto Scaling chooses to terminate this instance in order to reduce the size of the Auto Scaling group. For more information, see [Controlling Which Auto Scaling Instances Terminate During Scale In \(p. 110\)](#).
- You put the instance into a `Standby` state. For more information, see [Enter and Exit Standby \(p. 9\)](#).
- You detach the instance from the Auto Scaling group. For more information, see [Detach an Instance \(p. 9\)](#).
- The instance fails a required number of health checks, so it is removed from the Auto Scaling group, terminated, and replaced. For more information, see [Health Checks for Auto Scaling Instances \(p. 132\)](#).

Scale In

The following scale in events direct the Auto Scaling group to detach EC2 instances from the group and terminate them:

- You manually decrease the size of the group. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#).
- You create a scaling policy to automatically decrease the size of the group based on a specified decrease in demand. For more information, see [Dynamic Scaling for Amazon EC2 Auto Scaling \(p. 87\)](#).
- You set up scaling by schedule to decrease the size of the group at a specific time. For more information, see [Scheduled Scaling for Amazon EC2 Auto Scaling \(p. 85\)](#).

It is important that you create a corresponding scale in event for each scale out event that you create. This helps ensure that the resources assigned to your application match the demand for those resources as closely as possible.

When a scale in event occurs, the Auto Scaling group detaches one or more instances. The Auto Scaling group uses its termination policy to determine which instances to terminate. Instances that are in the process of detaching from the Auto Scaling group and shutting down enter the `Terminating` state, and can't be put back into service. If you add a lifecycle hook to your Auto Scaling group, you can perform a custom action here. Finally, the instances are completely terminated and enter the `Terminated` state.

Attach an Instance

You can attach a running EC2 instance that meets certain criteria to your Auto Scaling group. After the instance is attached, it is managed as part of the Auto Scaling group.

For more information, see [Attach EC2 Instances to Your Auto Scaling Group \(p. 78\)](#).

Detach an Instance

You can detach an instance from your Auto Scaling group. After the instance is detached, you can manage it separately from the Auto Scaling group or attach it to a different Auto Scaling group.

For more information, see [Detach EC2 Instances from Your Auto Scaling Group \(p. 82\)](#).

Lifecycle Hooks

You can add a lifecycle hook to your Auto Scaling group so that you can perform custom actions when instances launch or terminate.

When Amazon EC2 Auto Scaling responds to a scale out event, it launches one or more instances. These instances start in the `Pending` state. If you added an `autoscaling:EC2_INSTANCE_LAUNCHING` lifecycle hook to your Auto Scaling group, the instances move from the `Pending` state to the `Pending:Wait` state. After you complete the lifecycle action, the instances enter the `Pending:Proceed` state. When the instances are fully configured, they are attached to the Auto Scaling group and they enter the `InService` state.

When Amazon EC2 Auto Scaling responds to a scale in event, it terminates one or more instances. These instances are detached from the Auto Scaling group and enter the `Terminating` state. If you added an `autoscaling:EC2_INSTANCE_TERMINATING` lifecycle hook to your Auto Scaling group, the instances move from the `Terminating` state to the `Terminating:Wait` state. After you complete the lifecycle action, the instances enter the `Terminating:Proceed` state. When the instances are fully terminated, they enter the `Terminated` state.

For more information, see [Amazon EC2 Auto Scaling Lifecycle Hooks \(p. 115\)](#).

Enter and Exit Standby

You can put any instance that is in an `InService` state into a `Standby` state. This enables you to remove the instance from service, troubleshoot or make changes to it, and then put it back into service.

Instances in a `Standby` state continue to be managed by the Auto Scaling group. However, they are not an active part of your application until you put them back into service.

For more information, see [Temporarily Removing Instances from Your Auto Scaling Group \(p. 123\)](#).

Amazon EC2 Auto Scaling Limits

Your AWS account has the following limits related to Amazon EC2 Auto Scaling. For information about the service limits for other services, see [AWS Service Limits](#) in the *Amazon Web Services General Reference*.

Default Limits

- Launch configurations per Region: 200
- Auto Scaling groups per Region: 200

To view the current limits for your account, use the **Limits** page of the Amazon EC2 console or the [describe-account-limits](#) (AWS CLI) command. To request a limit increase, use the [Auto Scaling Limits form](#).

Auto Scaling Group Limits

- Scaling policies per Auto Scaling group: 50
- Scheduled actions per Auto Scaling group: 125
- Lifecycle hooks per Auto Scaling group: 50
- SNS topics per Auto Scaling group: 10
- Classic Load Balancers per Auto Scaling group: 50
- Target groups per Auto Scaling group: 50

Scaling Policy Limits

- Step adjustments per scaling policy: 20

Auto Scaling API Limits

- You can use [AttachInstances](#), [DetachInstances](#), [EnterStandby](#), and [ExitStandby](#) with at most 20 instance IDs at a time.
- You can use [AttachLoadBalancers](#) and [DetachLoadBalancers](#) with at most 10 load balancers at a time.
- You can use [AttachLoadBalancerTargetGroups](#) and [DetachLoadBalancerTargetGroups](#) with at most 10 target groups at a time.

Setting Up Amazon EC2 Auto Scaling

Before you start using Amazon EC2 Auto Scaling, complete the following tasks.

Tasks

- [Sign Up for AWS](#) (p. 11)
- [Prepare to Use Amazon EC2](#) (p. 11)

Sign Up for AWS

When you create an AWS account, we automatically sign up your account for all AWS services. You pay only for the services that you use. You can use Amazon EC2 Auto Scaling at no additional charge beyond what you are paying for your EC2 instances.

If you don't have an AWS account, sign up for AWS as follows.

To sign up for an AWS account

1. Open <https://portal.aws.amazon.com/billing/signup>.
2. Follow the online instructions.

Part of the sign-up procedure involves receiving a phone call and entering a verification code on the phone keypad.

AWS sends you a confirmation email after the sign-up process is complete.

Prepare to Use Amazon EC2

If you haven't used Amazon EC2 before, complete the tasks described in the Amazon EC2 documentation. For more information, see [Setting Up with Amazon EC2](#) in the *Amazon EC2 User Guide for Linux Instances* or [Setting Up with Amazon EC2](#) in the *Amazon EC2 User Guide for Windows Instances*.

Getting Started with Amazon EC2 Auto Scaling

When you use Amazon EC2 Auto Scaling, you must use certain building blocks to get started. This tutorial walks you through the process for setting up the basic infrastructure for Amazon EC2 Auto Scaling.

Before you create an Auto Scaling group for use with your application, review your application thoroughly as it runs in the AWS Cloud. Take note of the following:

- How long it takes to launch and configure a server.
- What metrics have the most relevance to your application's performance.
- How many Availability Zones the Auto Scaling group should span.
- What existing resources can be used, such as security groups or Amazon Machine Images (AMIs).
- Do you want to scale to increase or decrease capacity, or do you just want to ensure that a specific number of servers are always running? Keep in mind that Amazon EC2 Auto Scaling can do both simultaneously.

The better you understand your application, the more effective you can make your Auto Scaling architecture.

The following instructions are for a configuration template that defines your EC2 instances, creates an Auto Scaling group to maintain a fixed number of instances even if an instance becomes unhealthy, and optionally deletes this basic infrastructure.

This tutorial assumes that you are familiar with launching EC2 instances and have already created a key pair and a security group. For more information, see [Setting Up with Amazon EC2](#) in the *Amazon EC2 User Guide for Linux Instances*.

Tasks

- [Step 1: Create a Launch Template](#) (p. 12)
- [Step 2: Create an Auto Scaling Group](#) (p. 14)
- [Step 3: Verify Your Auto Scaling Group](#) (p. 15)
- [Step 4: \(Optional\) Delete Your Scaling Infrastructure](#) (p. 17)

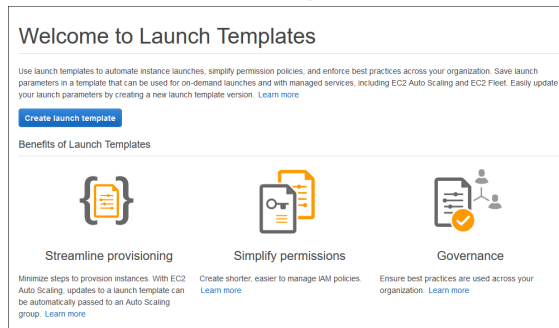
Step 1: Create a Launch Template

For this step, you create a launch template that specifies the type of EC2 instance that Amazon EC2 Auto Scaling creates for you. Include information such as the ID of the Amazon Machine Image (AMI) to use, the instance type, key pairs, security groups, and block device mappings.

To create a launch template for an Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, select an AWS Region. The Amazon EC2 Auto Scaling resources that you create are tied to the region you specify.

3. On the navigation pane, choose **Launch Templates**.
4. Choose **Create launch template**.



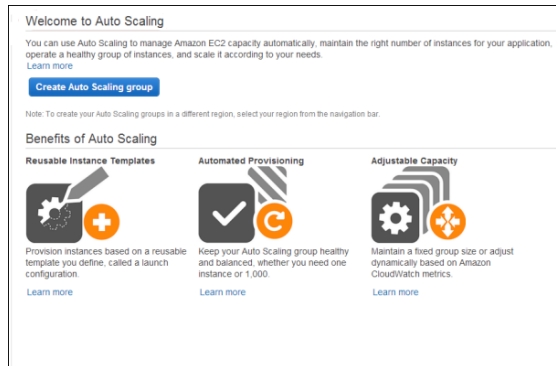
5. Choose **Create a new template**. For **Launch template name**, enter a name (for example, `my_template`).
6. For **Template version description**, enter a description for the initial version of the launch template to help you remember what this template is used for (for example, `test launch template for an Auto Scaling group`).
7. For **AMI ID**, choose a version of Amazon Linux 2 (HVM) from the **Quick Start** list. The Amazon Machine Image (AMI) serves as basic configuration templates for your instances.
8. For **Instance type**, choose a hardware configuration that is compatible with the AMI that you specified. Note that the free tier Linux server is a `t2.micro` instance.
Note
If your account is less than 12 months old, you can use a `t2.micro` instance for free within certain usage limits. For more information, see [AWS Free Tier](#).
9. (Optional) For **Key pair name**, type the name of the key pair to use when connecting to your instances.
10. (Optional) For **Network type**, choose **VPC**.
11. Skip **Security Groups** to configure a security group in the next step. When a network interface is specified, the security group must be a part of it.
12. For **Network interfaces**, do the following to specify the primary network interface:
 - a. Choose **Add network interface**.
 - b. (Optional) To assign public IP addresses to instances in a nondefault VPC, for **Auto-assign public IP**, choose **Enable**. This allows your instances to communicate with the internet and other services in AWS.
 - c. For **Security group ID**, specify a security group for the network interface.
 - d. For **Delete on termination**, choose whether the network interface is deleted when the Auto Scaling group scales in and terminates the instance to which the network interface is attached.
13. (Optional) For **Storage (Volumes)**, specify volumes to attach to the instances in addition to the volumes specified by the AMI you specified.
14. (Optional) For **Instance Tags**, specify one or more tags to associate with the instances and volumes.
15. Choose **Create launch template**.
16. On the confirmation page, choose **Create Auto Scaling group**.

If you are not currently using launch templates, you can create a launch configuration instead.

A launch configuration is similar to a launch template, in that it specifies the type of EC2 instance that Amazon EC2 Auto Scaling creates for you. Create the launch configuration by including information such as the ID of the Amazon Machine Image (AMI) to use, the instance type, the key pair, security groups, and block device mapping.

To create a launch configuration

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar, select a Region. The Auto Scaling resources that you create are tied to the Region that you specify.
3. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
4. On the **Welcome to Auto Scaling** page, choose **Create Auto Scaling group**.



5. On the **Create Auto Scaling Group** page, choose **Launch Configuration**, **Create a new launch configuration**, and then choose **Next Step**.
6. For the **Choose AMI** step, there is a list of basic configurations, called Amazon Machine Images (AMIs), that serve as templates for your instances. Choose **Select** for the Amazon Linux 2 AMI.
7. For the **Choose Instance Type** step, select a hardware configuration for your instances. We recommend that you keep the default, a `t2.micro` instance. Choose **Next: Configure details**.
8. For the **Configure details** step, do the following:
 - a. For **Name**, type a name for your launch configuration (for example, `my-first-lc`).
 - b. For **Advanced Details**, select an IP address type. If you want to provide internet connectivity to instances in a VPC, you must select an option that assigns a public IP address. If you want to provide internet connectivity to your instances but aren't sure whether you have a default VPC, select **Assign a public IP address to every instance**.
 - c. Choose **Skip to review**.
9. For the **Review** step, choose **Edit security groups**. Follow the instructions to choose an existing security group, and then choose **Review**.
10. For the **Review** step, choose **Create launch configuration**.
11. Complete the **Select an existing key pair or create a new key pair** step as instructed. You won't connect to your instance as part of this tutorial. Therefore, you can select **Proceed without a key pair** unless you intend to connect to your instance.
12. Choose **Create launch configuration**. The launch configuration is created and the wizard to create an Auto Scaling group is displayed.

Step 2: Create an Auto Scaling Group

An Auto Scaling group is a collection of EC2 instances, and the core of Amazon EC2 Auto Scaling. When you create an Auto Scaling group, you include information such as the subnets for the instances and the initial number of instances to start with.

Use the following procedure to continue where you left off after creating the launch configuration or template.

To create an Auto Scaling group

1. For the **Configure Auto Scaling group details** step, do the following:
 - a. For **Group name**, type a name for your Auto Scaling group (for example, `my-first-asg`).

The screenshot shows the 'Create Auto Scaling Group' page in the AWS Management Console. The page has a progress bar at the top with five steps: 1. Configure Auto Scaling group details (active), 2. Configure scaling policies, 3. Configure Notifications, 4. Configure Tags, and 5. Review. The main form includes the following fields and options:

- Group name:** `my-first-asg`
- Launch Template:** `it-08af5e8475e7ad30f`
- Launch Template Version:** `Default` (with a 'Create new launch template' link)
- Launch Template Description:** `-`
- Fleet Composition:** ☒ **Adhere to the launch template** (with a note: 'This launch template determines the instance type and purchase option (On-Demand or Spot). Choose a mix of On-Demand instances and Spot instances and multiple instance types. Spot instances are automatically launched at the lowest price available.') and ☐ **Combine purchase options and instances**
- Group size:** `Start with 1 instances`
- Network:** `vpc-a5f6eadd (172.31.0.0/16) (default)` (with a 'Create new VPC' link)
- Subnet:** (empty field with a 'Create new subnet' link)

At the bottom left, there is a link for 'Advanced Details'. At the top right, there is a 'Cancel and Exit' link.

- b. [Launch template] For **Launch template version**, choose whether the Auto Scaling group uses the default, the latest, or a specific version of the launch template when scaling out.
 - c. [Launch template] For **Fleet Composition**, choose **Adhere to the launch template**.
 - d. Keep **Group size** set to the default value of 1 instance for this tutorial.
 - e. Keep **Network** set to the default VPC for your chosen AWS Region, or select your own VPC.
 - f. For **Subnet**, choose a subnet for the VPC.
- Note**
You can choose the Availability Zone for your instance by choosing its corresponding default subnet.
- g. Choose **Next: Configure scaling policies**.
2. On the **Configure scaling policies** page, select **Keep this group at its initial size** and choose **Review**.
 3. On the **Review** page, choose **Create Auto Scaling group**.
 4. On the **Auto Scaling group creation status** page, choose **Close**.

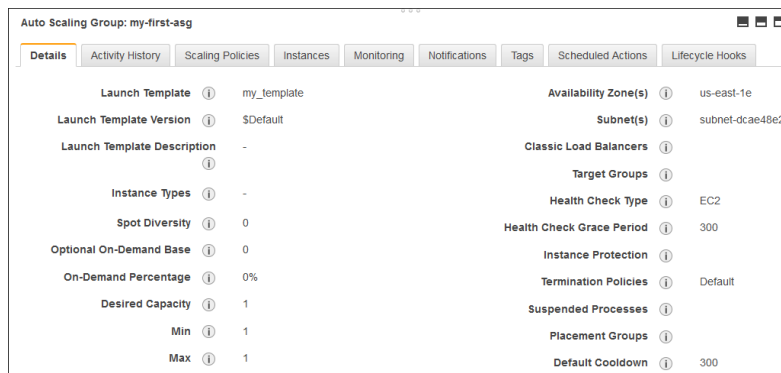
Step 3: Verify Your Auto Scaling Group

Now that you have created your Auto Scaling group, you are ready to verify that the group has launched an EC2 instance.


To verify that your Auto Scaling group has launched an EC2 instance

1. On the **Auto Scaling Groups** page, select the Auto Scaling group that you just created.
2. The **Details** tab provides information about the Auto Scaling group.

Amazon EC2 Auto Scaling User Guide (Optional) Terminate an Instance in Your Auto Scaling Group




3. On the **Activity History** tab, the **Status** column shows the current status of your instance. While your instance is launching, the status column shows **In progress**. The status changes to **Successful** after the instance is launched. You can also use the refresh button to see the current status of your instance.
4. On the **Instances** tab, the **Lifecycle** column shows the state of your instance. You can see that your Auto Scaling group has launched your EC2 instance, and that it is in the **InService** lifecycle state. The **Health Status** column shows the result of the EC2 instance health check on your instance.



(Optional) Terminate an Instance in Your Auto Scaling Group

If you want, you can try the following experiment to learn more about Amazon EC2 Auto Scaling. The minimum size for your Auto Scaling group is one instance. Therefore, if you terminate the running instance, Amazon EC2 Auto Scaling must launch a new instance to replace it.

1. On the **Instances** tab, select the ID of the instance. This shows you the instance on the **Instances** page.
2. Choose **Actions, Instance State, Terminate**. When prompted for confirmation, choose **Yes, Terminate**.
3. On the navigation pane, choose **Auto Scaling Groups**. Select your Auto Scaling group and choose the **Activity History** tab. The default cooldown for the Auto Scaling group is 300 seconds (5 minutes), so it takes about 5 minutes until you see the scaling activity. When the scaling activity starts, you see an entry for the termination of the first instance and an entry for the launch of a new instance. The **Instances** tab shows the new instance only.



4. On the navigation pane, choose **Instances**. This page shows both the terminated instance and the running instance.

Go to the next step if you would like to delete your basic infrastructure for automatic scaling. Otherwise, you can use this infrastructure as your base and try one or more of the following:

- [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#)
- [Dynamic Scaling for Amazon EC2 Auto Scaling \(p. 87\)](#)
- [Getting Amazon SNS Notifications When Your Auto Scaling Group Scales \(p. 145\)](#)

Step 4: (Optional) Delete Your Scaling Infrastructure

You can either delete your scaling infrastructure or delete just your Auto Scaling group and keep your launch template to use later.

If you launched an instance that is not within the [AWS Free Tier](#), you should terminate your instances to prevent additional charges. The EC2 instance and the data associated will be deleted.

To delete your Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. Choose **Actions**, **Delete**. When prompted for confirmation, choose **Yes, Delete**.

The **Name** column indicates that the Auto Scaling group is being deleted. The **Desired**, **Min**, and **Max** columns show 0 instances for the Auto Scaling group.

Skip this procedure if you would like to keep your launch template.

To delete your launch template

1. On the navigation pane, choose **Launch Templates**.
2. Select your launch template.
3. Choose **Actions**, **Delete template**. When prompted for confirmation, choose **Delete launch template**.

Skip this procedure if you would like to keep your launch configuration.

To delete your launch configuration

1. On the navigation pane, under **Auto Scaling**, choose **Launch Configurations**.
2. Select your launch configuration.
3. Choose **Actions**, **Delete launch configuration**. When prompted for confirmation, choose **Yes, Delete**.

Tutorial: Set Up a Scaled and Load-Balanced Application

You can attach a load balancer to your Auto Scaling group. The load balancer automatically distributes incoming traffic across the instances in the group.

This tutorial attaches a load balancer to an Auto Scaling group when you create the group. To attach a load balancer to an existing Auto Scaling group, see [Attaching a Load Balancer to Your Auto Scaling Group](#) (p. 63).

Before you explore this tutorial, we recommend that you first review the following introductory topic: [Using a Load Balancer with an Auto Scaling Group](#) (p. 61).

Contents

- [Prerequisites](#) (p. 18)
- [Configure Scaling and Load Balancing \(Console\)](#) (p. 18)
- [Configure Scaling and Load Balancing \(AWS CLI\)](#) (p. 22)
- [Cleaning Up Your AWS Resources](#) (p. 23)

Prerequisites

- (Optional) Create an IAM role that grants your application the access to AWS that it needs.
- Launch an instance. Be sure to specify the IAM role (if you created one), and specify any configuration scripts that you need as user data. Connect to the instance and customize it. For example, you can install software and applications and copy data. Test your application on your instance to ensure that your instance is configured correctly. Create a custom Amazon Machine Image (AMI) from your instance. You can terminate the instance if you no longer need it.
- Create a load balancer. Elastic Load Balancing supports three types of load balancers: Application Load Balancers, Network Load Balancers, and Classic Load Balancers. You can attach any of these types of load balancers to your Auto Scaling group. For more information, see [Elastic Load Balancing Types](#) (p. 62).
- Make sure that you choose the same Availability Zones for the load balancer that you plan to enable for your Auto Scaling group.

Configure Scaling and Load Balancing (Console)

Complete the following tasks to set up a scaled and load-balanced application when you create your Auto Scaling group.

Tasks

- [Create or Select a Launch Configuration](#) (p. 19)
- [Create or Select a Launch Template](#) (p. 20)
- [Create an Auto Scaling Group](#) (p. 21)

- (Optional) [Verify That Your Load Balancer Is Attached to Your Auto Scaling Group \(p. 21\)](#)

Create or Select a Launch Configuration

A launch configuration specifies the type of EC2 instance that Amazon EC2 Auto Scaling creates for you. When you create a launch configuration, you include information such as the ID of the Amazon Machine Image (AMI) to use, the instance type, key pair, and block device mapping. If you created a launch template, you can use your launch template to create an Auto Scaling group instead of using a launch configuration. For more information, see [Create or Select a Launch Template \(p. 20\)](#).

If you already have a launch configuration that you'd like to use, select it by using the following procedure.

To select an existing launch configuration

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, select the AWS Region that you used when creating your load balancer.
3. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
4. On the next page, choose **Create Auto Scaling group**.
5. On the **Create Auto Scaling Group** page, choose **Launch Configuration**, select an existing launch configuration, and then choose **Next Step**.

To create a new launch configuration, use the following procedure:

To create a launch configuration

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, select the AWS Region that you used when creating your load balancer.
3. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
4. On the next page, choose **Create Auto Scaling group**.
5. On the **Create Auto Scaling Group** page, choose **Launch Configuration**, **Create a new launch configuration**, and then choose **Next Step**.
6. On the **Choose AMI** page, select your custom AMI.
7. On the **Choose Instance Type** page, select a hardware configuration for your instance, and then choose **Next: Configure details**.
8. On the **Configure Details** page, do the following:
 - a. For **Name**, enter a name for your launch configuration.
 - b. (Optional) To securely distribute credentials to your EC2 instance, select your IAM role.
 - c. (Optional) If you need to connect to instances in a nondefault VPC, for **Advanced Details**, **IP Address Type**, choose **Assign a public IP address to every instance**.
 - d. (Optional) To specify user data or a configuration script for your instance, for **Advanced Details**, **User data**, paste your configuration script.
 - e. Choose **Skip to review**.
9. On the **Review** page, choose **Edit security groups**. Follow the instructions to choose an existing security group, and then choose **Review**.
10. On the **Review** page, choose **Create launch configuration**.
11. On the **Select an existing key pair or create a new key pair** page, select one of the listed options. Select the acknowledgment check box, and then choose **Create launch configuration**.

Warning

Do not choose **Proceed without a key pair** if you need to connect to your instance.

After completing the instructions above, you're ready to proceed with the wizard to create an Auto Scaling group.

Create or Select a Launch Template

If you already have a launch template that you'd like to use, select it using the following procedure.

To select an existing launch template

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, select the AWS Region that you used when creating your load balancer.
3. On the navigation pane, choose **Launch Templates**.
4. Select a launch template.
5. Choose **Actions, Create Auto Scaling group**.

Alternatively, to create a new launch template, use the following procedure.

To create a launch template

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, select the AWS Region that you used when creating your load balancer.
3. On the navigation pane, choose **Launch Templates**.
4. Choose **Create launch template**.
5. Provide a name and description for the launch template. The name of a launch template must have 3 to 125 characters, and can include the following characters: ()./_.
6. For **AMI ID**, enter the ID of the AMI for your instances.
7. For **Instance type**, select a hardware configuration for your instances that is compatible with the AMI that you specified.
8. (Optional) For **Key pair name**, enter the name of the key pair to use when connecting to your instances.
9. For **Network interfaces**, do the following:
 - a. Choose **Add network interface**.
 - b. (Optional) To connect to instances in a nondefault VPC, for **Auto-assign public IP**, choose **Enable**.
 - c. For **Security group ID**, specify a security group for your instances.
 - d. For **Delete on termination**, choose whether the network interface is deleted when the Auto Scaling group scales in and terminates the instance to which the network interface is attached.
10. (Optional) To securely distribute credentials to your instances, for **Advanced details, IAM instance profile**, enter the Amazon Resource Name (ARN) of your IAM role.
11. (Optional) To specify user data or a configuration script for your instances, paste it into **Advanced details, User data**.
12. Choose **Create launch template**.
13. On the confirmation page, choose **Create Auto Scaling group**.

Create an Auto Scaling Group

Use the following procedure to continue where you left off after selecting or creating your launch configuration or template.

To create an Auto Scaling group

1. On the **Configure Auto Scaling group details** page, do the following:
 - a. For **Group name**, enter a name for your Auto Scaling group.
 - b. [Launch template] For **Launch template version**, choose whether the Auto Scaling group uses the default, the latest, or a specific version of the launch template when scaling out.
 - c. [Launch template] For **Fleet Composition**, choose **Adhere to the launch template** to use the EC2 instance type and purchase option that are specified in the launch template.
 - d. For **Group size**, enter the initial number of instances for your Auto Scaling group.
 - e. If you selected an instance type for your launch configuration or template that requires a VPC, such as a T2 instance, you must select a VPC for **Network**. Otherwise, if your account supports EC2-Classic and you selected an instance type that doesn't require a VPC, you can select either **Launch into EC2-Classic** or a VPC.
 - f. If you selected a VPC in the previous step, select one or more subnets from **Subnet**. If you selected EC2-Classic, select one or more Availability Zones from **Availability Zone(s)**.
 - g. For **Advanced Details**, select **Receive traffic from Elastic Load Balancer(s)** and then do one of the following:
 - [Classic Load Balancers] Select your load balancer from **Load Balancers**.
 - [Application/Network Load Balancers] Select your target group from **Target Groups**.
 - h. (Optional) To use Elastic Load Balancing health checks, choose **ELB** for **Advanced Details, Health Check Type**.
 - i. Choose **Next: Configure scaling policies**.
2. On the **Configure scaling policies** page, select **Keep this group at its initial size**, and then choose **Review**.

To configure scaling policies for your Auto Scaling group, see [Configure Scaling Policies \(p. 91\)](#).
3. Review the details of your Auto Scaling group. You can choose **Edit** to make changes. When you are finished, choose **Create Auto Scaling group**.

(Optional) Verify That Your Load Balancer Is Attached to Your Auto Scaling Group

To verify that your load balancer is attached to your Auto Scaling group

1. Select your Auto Scaling group.
2. On the **Details** tab, **Load Balancers** shows any attached load balancers and **Target Groups** shows any attached target groups.
3. On the **Activity History** tab, the **Status** column shows you the status of your Auto Scaling instances. While an instance is launching, its status is `In progress`. The status changes to `Successful` after the instance is launched.
4. On the **Instances** tab, the **Lifecycle** column shows the state of your Auto Scaling instances. After an instance is ready to receive traffic, its state is `InService`.

The **Health Status** column shows the result of the health checks on your instances.

Configure Scaling and Load Balancing (AWS CLI)

Complete the following tasks to set up a scaled and load-balanced application.

Tasks

- [Create a Launch Configuration \(p. 22\)](#)
- [Create a Launch Template \(p. 22\)](#)
- [Create an Auto Scaling Group with a Load Balancer \(p. 22\)](#)

Create a Launch Configuration

If you already have a launch configuration that you'd like to use, skip this step.

To create the launch configuration

Use the following [create-launch-configuration](#) command. To assign public IP addresses to instances in a nondefault VPC, use the `--associate-public-ip-address` option as shown.

```
aws autoscaling create-launch-configuration --launch-configuration-name my-launch-config \
--image-id ami-01e24be29428c15b2 --instance-type t2.micro --associate-public-ip-address \
--security-groups sg-eb2af88e
```

Create a Launch Template

If you already have a launch template that you'd like to use, skip this step.

To create the launch template

Use the following [create-launch-template](#) command and pass a JSON file that contains the information for creating the launch template. To assign public IP addresses to instances in a nondefault VPC, add the network interface attribute and set it to "NetworkInterfaces": {"AssociatePublicIpAddress":true} as shown.

```
aws ec2 create-launch-template --launch-template-name my-launch-template --version-
description my-version-description \
--launch-template-data '{"NetworkInterfaces":
[{"DeviceIndex":0,"AssociatePublicIpAddress":true,"Groups":
["sg-903004f8"],"DeleteOnTermination":true}],"ImageId":"ami-01e24be29428c15b2","InstanceType":"t2.micro"
```

Create an Auto Scaling Group with a Load Balancer

You can attach an existing load balancer to an Auto Scaling group when you create the group. You can use either a launch configuration or a launch template to automatically configure the instances that your Auto Scaling group launches.

To create an Auto Scaling group with an attached Classic Load Balancer

Use the following [create-auto-scaling-group](#) command with the `--load-balancer-names` option to create an Auto Scaling group with an attached Classic Load Balancer.

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg \
--launch-configuration-name my-launch-config \
--vpc-zone-identifier "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782" \
```

```
--load-balancer-names "my-load-balancer" \  
--max-size 5 --min-size 1 --desired-capacity 2
```

To create an Auto Scaling group with an attached target group for an Application Load Balancer or Network Load Balancer

Use the following [create-auto-scaling-group](#) command with the `--target-group-arns` option to create an Auto Scaling group with an attached target group.

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg \  
--launch-template "LaunchTemplateName=my-launch-template,Version=1" \  
--vpc-zone-identifier "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782" \  
--target-group-arns "arn:aws:elasticloadbalancing:region:123456789012:targetgroup/my-  
targets/1234567890123456" \  
--max-size 5 --min-size 1 --desired-capacity 2
```

Cleaning Up Your AWS Resources

You've now successfully completed the tutorial. To avoid ongoing charges for resources that you aren't using, you should clean up the resources that you created just for this tutorial. For step-by-step instructions, see [Deleting Your Auto Scaling Infrastructure](#) (p. 70).

Launch Templates

A launch template is similar to a [launch configuration \(p. 32\)](#), in that it specifies instance configuration information. Included are the ID of the Amazon Machine Image (AMI), the instance type, a key pair, security groups, and the other parameters that you use to launch EC2 instances. However, defining a launch template instead of a launch configuration allows you to have multiple versions of a template. With versioning, you can create a subset of the full set of parameters and then reuse it to create other templates or template versions. For example, you can create a default template that defines common configuration parameters such as tags or network configurations, and allow the other parameters to be specified as part of another version of the same template.

We recommend that you use launch templates instead of launch configurations to ensure that you can use the latest features of Amazon EC2, such as T2 Unlimited instances. For more information, see [Unlimited Mode for Burstable Performance Instances](#) and [Using an Auto Scaling Group to Launch a Burstable Performance Instance as Unlimited](#) in the *Amazon EC2 User Guide for Linux Instances*.

With launch templates, you can also provision capacity across multiple instance types using both On-Demand Instances and Spot Instances to achieve the desired scale, performance, and cost. For more information, see [Auto Scaling Groups with Multiple Instance Types and Purchase Options \(p. 43\)](#).

If you currently use launch configurations, you can specify a launch template when you update an Auto Scaling group that was created using a launch configuration.

To create a launch template to use with an Auto Scaling group, create the template from scratch, create a new version of an existing template, or copy the parameters from a launch configuration, running instance, or other template.

The following topics describe the most common procedures for creating and working with launch templates for use with your Auto Scaling groups. For more information about launch templates, see the [launch templates](#) section of the *Amazon EC2 User Guide for Linux Instances*.

Contents

- [Creating a Launch Template for an Auto Scaling Group \(p. 24\)](#)
- [Copying a Launch Configuration to a Launch Template \(p. 29\)](#)
- [Replacing a Launch Configuration with a Launch Template \(p. 30\)](#)

Creating a Launch Template for an Auto Scaling Group

Before you can create an Auto Scaling group using a launch template, you must create a launch template that includes the parameters required to launch an EC2 instance, such as the ID of the Amazon Machine Image (AMI) and an instance type.

The following procedure works for creating a new launch template. The new template uses parameters that you define (starting from scratch) or an existing launch template. After you create your launch template, you can create the Auto Scaling group by following the instructions in [Creating an Auto Scaling Group Using a Launch Template \(p. 51\)](#).

Prerequisites

For information about the required IAM permissions, see [Controlling the Use of Launch Templates](#) in the *Amazon EC2 User Guide for Linux Instances*.

Considerations

Keep the following considerations in mind when creating a launch template for use with an Auto Scaling group:

- Launch templates give you the flexibility of launching one type of instance, or a combination of instance types and On-Demand and Spot purchase options. For more information, see [Auto Scaling Groups with Multiple Instance Types and Purchase Options](#) (p. 43). Launching instances with such a combination is not supported:
 - If you specify a Spot Instance request in **Additional Details**
 - In EC2-Classic
- If you configure a network type (VPC or EC2-Classic), subnet, and Availability Zone for your template, these settings are ignored in favor of what is specified in the Auto Scaling group.
- If you specify a network interface, you must configure the security group as part of the network interface, and not in the **Security Groups** section of the template.
- You cannot specify multiple network interfaces.
- You cannot assign specific private IP addresses. When an instance launches, a private address is allocated from the CIDR range of the subnet in which the instance is launched. For more information on specifying CIDR ranges for your VPC or subnet, see the [Amazon VPC User Guide](#).
- To specify an existing network interface to use, its device index must be 0 (eth0). For this scenario, you must use the CLI or API to create the Auto Scaling group. When you create the group using the CLI create-auto-scaling-group command or API CreateAutoScalingGroup action, you must specify the Availability Zones parameter instead of the subnet (VPC zone identifier) parameter.
- You cannot use host placement affinity.

To create a new launch template for an Auto Scaling group using the console

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Launch Templates**.
3. Choose **Create a new template**. Enter a name and provide a description for the initial version of the launch template.
4. If you choose to create the new template based on another template:
 - a. For **Source template**, choose an existing launch template.
 - b. Choose the launch template version on which to base the new launch template from **Source template version**.
5. Under **Launch template contents**, provide the following information.
 - a. **AMI ID**: Choose the ID of the Amazon Machine Image (AMI) from which to launch the instances. You can search through all available AMIs using the **Search for AMI** dialog. From the **Quick Start** list, select from one of the commonly used AMIs in the list. If you don't see the AMI that you need, select the **AWS Marketplace** or **Community AMIs** list to [find a suitable AMI](#).
 - b. **Instance type**: Choose the [instance type](#). Ensure that the instance type is compatible with the AMI you've specified.
 - c. **Key pair name**: Specify the [key pair](#) to use when connecting to your instances.
 - d. **Network type**: You can choose to specify whether to launch instances into a VPC or EC2-Classic, if applicable. However, the network type and Availability Zone settings of the launch template are ignored for Amazon EC2 Auto Scaling in favor of the settings of the Auto Scaling group.

- e. **Security Groups:** Choose one or more [security groups](#), or leave blank to configure the security group as part of the network interface. You cannot specify security groups in both places. If you're using EC2-Classic, you must use security groups created specifically for EC2-Classic.
6. Under **Network interfaces**, choose **Add network interface** and provide the following optional information. You can only specify one network interface. Pay attention to the following fields:
 - a. **Device:** Specify `eth0` as the device name (the device for which the device index is 0).
 - b. **Network interface:** Leave blank to let AWS create a new network interface when an instance is launched, or enter the ID of an existing network interface. If you specify an ID, this limits your Auto Scaling group to one instance.
 - c. **Description:** Enter a descriptive name.
 - d. **Subnet:** While you may choose to specify a subnet, it is ignored for Amazon EC2 Auto Scaling in favor of the settings of the Auto Scaling group.
 - e. **Auto-assign public IP:** Specify whether to automatically assign a public IP address to the network interface with the device index of `eth0`. This setting can only be enabled for a single, new network interface.
 - f. **Security group ID:** Enter the IDs of one or more [security groups](#) with which to associate the primary network interface (`eth0`). Each security group must be configured for the VPC that your Auto Scaling group will launch instances into. Separate the entries with commas.
 - g. **Delete on termination:** Choose whether the network interface is deleted when the Auto Scaling group scales in and terminates the instance to which the network interface is attached.
7. If you choose to specify volumes to attach to the instances not including the volumes specified by the AMI, under **Storage (Volumes)**, provide the following information:
 - a. **Volume type:** Choose instance store or Amazon EBS volumes. The type of volume depends on the instance type that you've chosen. For more information, see [Amazon EC2 Instance Store](#) and [Amazon EBS Volumes](#) in the *Amazon EC2 User Guide for Linux Instances*.
 - b. **Device name:** Specify a device name for the volume.
 - c. **Snapshot:** Enter the ID of the snapshot from which to create the volume.
 - d. **Size:** For Amazon EBS-backed volumes, specify a storage size. If you're creating the volume from a snapshot and don't specify a volume size, the default is the snapshot size.
 - e. **Volume type:** For Amazon EBS volumes, choose the [volume type](#).
 - f. **IOPS:** With a Provisioned IOPS SSD volume, enter the maximum number of input/output operations per second (IOPS) that the volume should support.
 - g. **Delete on termination:** For Amazon EBS volumes, choose whether to delete the volume when the associated instance is terminated.
 - h. **Encrypted:** Choose **Yes** to change the encryption state of an Amazon EBS volume. The default effect of setting this parameter varies with the choice of volume source, as described in the table below. You must in all cases have permission to use the specified CMK. For more information about specifying encrypted volumes, see [Amazon EBS Encryption](#) in the *Amazon EC2 User Guide for Linux Instances*.

Encryption Outcomes

If Encrypted parameter is set to...	And if source of volume is...	Then the default encryption state is...	Notes
No	New (empty) volume	Unencrypted*	N/A
	Unencrypted snapshot that you own	Unencrypted*	

If Encrypted parameter is set to...	And if source of volume is...	Then the default encryption state is...	Notes
	Encrypted snapshot that you own	Encrypted by same key	
	Unencrypted snapshot that is shared with you	Unencrypted*	
	Encrypted snapshot that is shared with you	Encrypted by default CMK	
Yes	New volume	Encrypted by default CMK	To use a non-default CMK, specify a value for the Key parameter.
	Unencrypted snapshot that you own	Encrypted by default CMK	
	Encrypted snapshot that you own	Encrypted by same key	
	Unencrypted snapshot that is shared with you	Encrypted by default CMK	
	Encrypted snapshot that is shared with you	Encrypted by default CMK	

* If [encryption by default](#) is enabled, all newly created volumes (whether or not the **Encrypted** parameter is set to **Yes**) are encrypted using the default CMK. Setting both the **Encrypted** and **Key** parameters allows you to specify a non-default CMK.

- i. [Optional] **Key:** If you chose **Yes** in the previous step, enter the customer master key (CMK) you want to use when encrypting the volumes. Enter any CMK that you previously created using the AWS Key Management Service. You can paste the full ARN of any key that you have access to. For more information, see the [AWS Key Management Service Developer Guide](#) and the [Required CMK Key Policy for Use with Encrypted Volumes \(p. 166\)](#) topic in this guide.

Note

Providing a CMK without also setting the **Encrypted** parameter results in an error.

8. For **Instance tags**, specify tags by providing key and value combinations. You can tag the instances, the volumes, or both.
 9. For **Advanced Details**, expand the section to view the fields and specify any additional parameters for the instances. For more information about this section and how each parameter can be used, see [Creating a Launch Template](#) in the *Amazon EC2 User Guide for Linux Instances*.
- **Purchasing option:** You have the option to request Spot Instances and specify the maximum price you are willing to pay per instance hour. For this to work with an Auto Scaling group, you must specify a one-time request with no end date. For more information, see [Launching Spot Instances in Your Auto Scaling Group \(p. 68\)](#).

Important

If you plan to specify multiple instance types and purchase options when you configure your Auto Scaling group, leave these fields empty. For more information, see [Auto Scaling Groups with Multiple Instance Types and Purchase Options \(p. 43\)](#).

- **IAM instance profile:** Specify an AWS Identity and Access Management (IAM) instance profile to associate with the instances. For more information, see [IAM Role for Applications That Run on Amazon EC2 Instances \(p. 169\)](#).

- **Shutdown behavior:** You can leave this field blank because it is ignored for Amazon EC2 Auto Scaling. The default behavior of Amazon EC2 Auto Scaling is to terminate the instance.
- **Termination protection:** Provides additional termination protection but is ignored for Amazon EC2 Auto Scaling when scaling in your Auto Scaling group. To control whether an Auto Scaling group can terminate a particular instance when scaling in, use [Instance Protection](#) (p. 113).
- **Monitoring:** Choose whether to enable detailed monitoring of the instances using Amazon CloudWatch. Additional charges apply. For more information, see [Monitoring Your Auto Scaling Groups and Instances Using Amazon CloudWatch](#) (p. 135).
- **T2/T3 Unlimited:** (Only valid for T2 and T3 instances) Choose whether to enable applications to burst beyond the baseline for as long as needed. Additional charges may apply. For more information, see [Using an Auto Scaling Group to Launch a Burstable Performance Instance as Unlimited](#) in the *Amazon EC2 User Guide for Linux Instances*.
- **Placement group name:** Specify a placement group in which to launch the instances. Not all instance types can be launched in a placement group. If you configure an Auto Scaling group using a CLI command that specifies a different placement group, the setting is ignored in favor of the one specified for the Auto Scaling group.
- **EBS-optimized instance:** Provides additional, dedicated capacity for Amazon EBS I/O. Not all instance types support this feature, and additional charges apply.
- **Tenancy:** You can choose to run your instances either on shared hardware (**Shared**), or on isolated, dedicated hardware (**Dedicated**). Additional charges may apply. You cannot choose **Dedicated host**.
- **RAM disk ID:** The ID of the RAM disk associated with the AMI. Only valid for paravirtual (PV) AMIs.
- **Kernel ID:** The ID of the kernel associated with the AMI. Only valid for paravirtual (PV) AMIs.
- **User data:** You can specify user data to configure an instance during launch, or to run a configuration script.

10. Choose **Create launch template**.

To create a launch template from an existing instance

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Instances**.
3. Select the instance and choose **Actions, Create template from instance**.
4. Provide a name and description. Adjust any other launch parameters as required, and choose **Create Launch Template**.

To create a launch template using the command line

You can use one of the following commands:

- [create-launch-template](#) (AWS CLI)
- [New-EC2LaunchTemplate](#) (AWS Tools for Windows PowerShell)

Create a launch template using the [create-launch-template](#) command as follows. Specify a value for Groups that corresponds to security groups for the VPC that your Auto Scaling group will launch instances into. Specify the VPC and subnets as properties of the Auto Scaling group.

```
aws ec2 create-launch-template --launch-template-name my-template-for-auto-scaling --
version-description version1 \
--launch-template-data '{"NetworkInterfaces":
[{"DeviceIndex":0,"AssociatePublicIpAddress":true,"Groups":
["sg-7c227019"],"DeleteOnTermination":true},"ImageId":"ami-01e24be29428c15b2","InstanceType":"t2.micro",
[{"ResourceType":"instance","Tags":[{"Key":"purpose","Value":"webserver"}]}]}'
```

Use the following `describe-launch-templates` command to describe the launch template `my-template-for-auto-scaling`.

```
aws ec2 describe-launch-templates --launch-template-names my-template-for-auto-scaling
```

Use the following `describe-launch-template-versions` command to describe the versions of the specified launch template `my-template-for-auto-scaling`.

```
aws ec2 describe-launch-template-versions --launch-template-id lt-068f72b72934aff71
```

The following is an example response:

```
{
  "LaunchTemplateVersions": [
    {
      "VersionDescription": "version1",
      "LaunchTemplateId": "lt-068f72b72934aff71",
      "LaunchTemplateName": "my-template-for-auto-scaling",
      "VersionNumber": 1,
      "CreatedBy": "arn:aws:iam::123456789012:user/Bob",
      "LaunchTemplateData": {
        "TagSpecifications": [
          {
            "ResourceType": "instance",
            "Tags": [
              {
                "Key": "purpose",
                "Value": "webserver"
              }
            ]
          }
        ],
        "ImageId": "ami-01e24be29428c15b2",
        "InstanceType": "t2.micro",
        "NetworkInterfaces": [
          {
            "DeviceIndex": 0,
            "DeleteOnTermination": true,
            "Groups": [
              "sg-7c227019"
            ],
            "AssociatePublicIpAddress": true
          }
        ],
        "DefaultVersion": true,
        "CreateTime": "2019-02-28T19:52:27.000Z"
      }
    ]
  }
}
```

Copying a Launch Configuration to a Launch Template

Use the following procedure to copy the options from an existing launch configuration to create a new launch template.

You can create launch templates from existing launch configurations to make it easy for you to update your Auto Scaling groups to use launch templates. Like launch configurations, launch templates can contain all or some of the parameters to launch an instance. With launch templates, you can also create multiple versions of a template to make it faster and easier to launch new instances.

To create a launch template from a launch configuration

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Launch Configurations**.
3. Select the launch configuration you want to copy and choose **Copy to launch template**. This sets up a new launch template with the same name and options as the launch configuration you selected.
4. For **New launch template name**, you can use the name of the launch configuration (the default) or enter a new name. Launch template names must be unique.
5. (Optional) To create an Auto Scaling group using the new launch template, select **Create an Auto Scaling group using the new template**. For more information, see [Creating an Auto Scaling Group Using a Launch Template \(p. 51\)](#).
6. Choose **Submit**.

After creating your launch template, you can update your existing Auto Scaling groups, as needed, with the launch template that you created. For more information, see [Replacing a Launch Configuration with a Launch Template \(p. 30\)](#).

Replacing a Launch Configuration with a Launch Template

When you edit an Auto Scaling group that has an existing launch configuration, you have the option of replacing the launch configuration with a launch template. This lets you use launch templates with any Auto Scaling groups that you currently use. In doing so, you can take advantage of the versioning and other features of launch templates.

After you replace the launch configuration for an Auto Scaling group, any new instances are launched using the new launch template, but existing instances are not affected.

Prerequisites

Before you can replace a launch configuration in an Auto Scaling group, you must first create your launch template. The easiest way to create a launch template is to copy it from the launch configuration. For more information, see [Copying a Launch Configuration to a Launch Template \(p. 29\)](#).

To replace the launch configuration for an Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group and choose **Details, Edit**.
4. For **Launch Instances Using**, choose the **Launch Template** option.
5. For **Launch Template**, select your launch template.
6. For **Launch Template Version**, select the launch template version, as needed. After you create versions of a launch template, you can choose whether the Auto Scaling group uses the default or the latest version of the launch template when scaling out.
7. When you have finished, choose **Save**.

To replace a launch configuration using the command line

You can use one of the following commands:

- [update-auto-scaling-group](#) (AWS CLI)
- [Update-ASAutoScalingGroup](#) (AWS Tools for Windows PowerShell)

Launch Configurations

A *launch configuration* is an instance configuration template that an Auto Scaling group uses to launch EC2 instances. When you create a launch configuration, you specify information for the instances. Include the ID of the Amazon Machine Image (AMI), the instance type, a key pair, one or more security groups, and a block device mapping. If you've launched an EC2 instance before, you specified the same information in order to launch the instance.

You can specify your launch configuration with multiple Auto Scaling groups. However, you can only specify one launch configuration for an Auto Scaling group at a time, and you can't modify a launch configuration after you've created it. To change the launch configuration for an Auto Scaling group, you must create a launch configuration and then update your Auto Scaling group with it.

Keep in mind that whenever you create an Auto Scaling group, you must specify a launch configuration, a launch template, or an EC2 instance. When you create an Auto Scaling group using an EC2 instance, Amazon EC2 Auto Scaling automatically creates a launch configuration for you and associates it with the Auto Scaling group. For more information, see [Creating an Auto Scaling Group Using an EC2 Instance](#) (p. 54). Alternatively, if you are using launch templates, you can specify a launch template instead of a launch configuration or an EC2 instance. For more information, see [Launch Templates](#) (p. 24).

Contents

- [Creating a Launch Configuration](#) (p. 32)
- [Creating a Launch Configuration Using an EC2 Instance](#) (p. 33)
- [Changing the Launch Configuration for an Auto Scaling Group](#) (p. 37)
- [Launching Auto Scaling Instances in a VPC](#) (p. 38)

Creating a Launch Configuration

When you create a launch configuration, you must specify information about the EC2 instances to launch. Include the Amazon Machine Image (AMI), instance type, key pair, security groups, and block device mapping. Alternatively, you can create a launch configuration using attributes from a running EC2 instance. For more information, see [Creating a Launch Configuration Using an EC2 Instance](#) (p. 33).

Note

When you create an Auto Scaling group, you can specify a launch template, launch configuration, or an EC2 instance. We recommend that you use a launch template to ensure that you can use the latest features of Amazon EC2. For more information, see [Launch Templates](#) (p. 24).

After you create a launch configuration, you can create an Auto Scaling group. For more information, see [Creating an Auto Scaling Group Using a Launch Configuration](#) (p. 53).

An Auto Scaling group is associated with one launch configuration at a time, and you can't modify a launch configuration after you've created it. Therefore, if you want to change the launch configuration for an existing Auto Scaling group, you must update it with the new launch configuration. For more information, see [Changing the Launch Configuration for an Auto Scaling Group](#) (p. 37).

To create a launch configuration using the console

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, the current region is displayed. Select a region for your Auto Scaling group that meets your needs.

3. On the navigation pane, under **Auto Scaling**, choose **Launch Configurations**.
4. On the next page, choose **Create launch configuration**.
5. On the **Choose AMI** page, select an AMI.
6. On the **Choose Instance Type** page, select a hardware configuration for your instance. Choose **Next: Configure details**.

Note

T2 instances must be launched into a subnet of a VPC. If you select a `t2.micro` instance but don't have a VPC, one is created for you. This VPC includes a public subnet in each Availability Zone in the region.

7. On the **Configure Details** page, do the following:
 - a. For **Name**, type a name for your launch configuration.
 - b. (Optional) For **Purchasing option**, you may request Spot Instances and specify the maximum price you are willing to pay per instance hour. For more information, see [Launching Spot Instances in Your Auto Scaling Group](#) (p. 68).
 - c. (Optional) For **IAM role**, select a role to associate with the instances. For more information, see [IAM Role for Applications That Run on Amazon EC2 Instances](#) (p. 169).
 - d. (Optional) By default, basic monitoring is enabled for your Auto Scaling instances. To enable detailed monitoring for your Auto Scaling instances, select **Enable CloudWatch detailed monitoring**.
 - e. For **Advanced Details, IP Address Type**, select an option. To connect to instances in a VPC, you must select an option that assigns a public IP address. If you want to connect to your instances but aren't sure whether you have a default VPC, select **Assign a public IP address to every instance**.
 - f. Choose **Skip to review**.
8. On the **Review** page, choose **Edit security groups**. Follow the instructions to choose an existing security group, and then choose **Review**.
9. On the **Review** page, choose **Create launch configuration**.
10. For **Select an existing key pair or create a new key pair**, select one of the listed options. Select the acknowledgment check box, and then choose **Create launch configuration**.

Warning

If you need to connect to your instance, do not select **Proceed without a key pair**.

To create a launch configuration using the command line

You can use one of the following commands:

- [create-launch-configuration](#) (AWS CLI)
- [New-ASLaunchConfiguration](#) (AWS Tools for Windows PowerShell)

Creating a Launch Configuration Using an EC2 Instance

Amazon EC2 Auto Scaling provides you with an option to create a launch configuration using the attributes from a running EC2 instance.

Tip

You can [create an Auto Scaling group directly from an EC2 instance](#) (p. 54). When you use this feature, Amazon EC2 Auto Scaling automatically creates a launch configuration for you as well.

If the specified instance has properties that are not currently supported by launch configurations, the instances launched by the Auto Scaling group might not be identical to the original EC2 instance.

There are differences between creating a launch configuration from scratch and creating a launch configuration from an existing EC2 instance. When you create a launch configuration from scratch, you specify the image ID, instance type, optional resources (such as storage devices), and optional settings (like monitoring). When you create a launch configuration from a running instance, Amazon EC2 Auto Scaling derives attributes for the launch configuration from the specified instance. Attributes are also derived from the block device mapping for the AMI from which the instance was launched, ignoring any additional block devices that were added after launch.

When you create a launch configuration using a running instance, you can override the following attributes by specifying them as part of the same request: AMI, block devices, key pair, instance profile, instance type, kernel, monitoring, placement tenancy, ramdisk, security groups, Spot price, user data, whether the instance has a public IP address is associated, and whether the instance is EBS-optimized.

The following examples show you to create a launch configuration from an EC2 instance.

Examples

- [Create a Launch Configuration Using an EC2 Instance \(p. 34\)](#)
- [Create a Launch Configuration from an Instance and Override the Block Devices \(AWS CLI\) \(p. 35\)](#)
- [Create a Launch Configuration and Override the Instance Type \(AWS CLI\) \(p. 36\)](#)

Create a Launch Configuration Using an EC2 Instance

To create a launch configuration using the attributes of an existing EC2 instance, specify the ID of the instance.

Important

The AMI used to launch the specified instance must still exist.

Create a Launch Configuration from an EC2 Instance (Console)

You can use the console to create a launch configuration and an Auto Scaling group from a running EC2 instance and add the instance to the new Auto Scaling group. For more information, see [Attach EC2 Instances to Your Auto Scaling Group \(p. 78\)](#).

Create a Launch Configuration from an EC2 Instance (AWS CLI)

Use the following `create-launch-configuration` command to create a launch configuration from an instance using the same attributes as the instance. Any block devices added after launch are ignored.

```
aws autoscaling create-launch-configuration --launch-configuration-name my-lc-from-instance
--instance-id i-a8e09d9c
```

You can use the following `describe-launch-configurations` command to describe the launch configuration and verify that its attributes match those of the instance:

```
aws autoscaling describe-launch-configurations --launch-configuration-names my-lc-from-instance
```

The following is an example response:

```
{
  "LaunchConfigurations": [
    {
```

```
    "UserData": null,
    "EbsOptimized": false,
    "LaunchConfigurationARN": "arn",
    "InstanceMonitoring": {
      "Enabled": false
    },
    "ImageId": "ami-05355a6c",
    "CreatedTime": "2014-12-29T16:14:50.382Z",
    "BlockDeviceMappings": [],
    "KeyName": "my-key-pair",
    "SecurityGroups": [
      "sg-8422d1eb"
    ],
    "LaunchConfigurationName": "my-lc-from-instance",
    "KernelId": "null",
    "RamdiskId": null,
    "InstanceType": "t1.micro",
    "AssociatePublicIpAddress": true
  }
]
```

Create a Launch Configuration from an Instance and Override the Block Devices (AWS CLI)

By default, Amazon EC2 Auto Scaling uses the attributes from the EC2 instance that you specify to create the launch configuration. However, the block devices come from the AMI used to launch the instance, not the instance. To add block devices to the launch configuration, override the block device mapping for the launch configuration.

Important

The AMI used to launch the specified instance must still exist.

Create a Launch Configuration and Override the Block Devices

Use the following [create-launch-configuration](#) command to create a launch configuration using an EC2 instance but with a custom block device mapping:

```
aws autoscaling create-launch-configuration --launch-configuration-name my-lc-from-
instance-bdm --instance-id i-a8e09d9c \
  --block-device-mappings "[{\"DeviceName\":\"/dev/sda1\",\"Ebs\":{\"SnapshotId\":
\"snap-3decf207\"}}, {\"DeviceName\":\"/dev/sdf\",\"Ebs\":{\"SnapshotId\":\"snap-
eed6ac86\"}}]"
```

Use the following [describe-launch-configurations](#) command to describe the launch configuration and verify that it uses your custom block device mapping:

```
aws autoscaling describe-launch-configurations --launch-configuration-names my-lc-from-
instance-bdm
```

The following example response describes the launch configuration:

```
{
  "LaunchConfigurations": [
    {
      "UserData": null,
      "EbsOptimized": false,
      "LaunchConfigurationARN": "arn",
```

```
    "InstanceMonitoring": {
      "Enabled": false
    },
    "ImageId": "ami-c49c0dac",
    "CreatedTime": "2015-01-07T14:51:26.065Z",
    "BlockDeviceMappings": [
      {
        "DeviceName": "/dev/sda1",
        "Ebs": {
          "SnapshotId": "snap-3decf207"
        }
      },
      {
        "DeviceName": "/dev/sdf",
        "Ebs": {
          "SnapshotId": "snap-eed6ac86"
        }
      }
    ],
    "KeyName": "my-key-pair",
    "SecurityGroups": [
      "sg-8637d3e3"
    ],
    "LaunchConfigurationName": "my-lc-from-instance-bdm",
    "KernelId": null,
    "RamdiskId": null,
    "InstanceType": "t1.micro",
    "AssociatePublicIpAddress": true
  }
}
```

Create a Launch Configuration and Override the Instance Type (AWS CLI)

By default, Amazon EC2 Auto Scaling uses the attributes from the EC2 instance you specify to create the launch configuration. Depending on your requirements, you might want to override attributes from the instance and use the values that you need. For example, you can override the instance type.

Important

The AMI used to launch the specified instance must still exist.

Create a Launch Configuration and Override the Instance Type

Use the following [create-launch-configuration](#) command to create a launch configuration using an EC2 instance but with a different instance type (for example `t2.medium`) than the instance (for example `t2.micro`):

```
aws autoscaling create-launch-configuration --launch-configuration-name my-lc-from-
instance-change-type \
  --instance-id i-a8e09d9c --instance-type t2.medium
```

Use the following [describe-launch-configurations](#) command to describe the launch configuration and verify that the instance type was overridden:

```
aws autoscaling describe-launch-configurations --launch-configuration-names my-lc-from-
instance-change-type
```

The following example response describes the launch configuration:

```
{
  "LaunchConfigurations": [
    {
      "UserData": null,
      "EbsOptimized": false,
      "LaunchConfigurationARN": "arn",
      "InstanceMonitoring": {
        "Enabled": false
      },
      "ImageId": "ami-05355a6c",
      "CreatedTime": "2014-12-29T16:14:50.382Z",
      "BlockDeviceMappings": [],
      "KeyName": "my-key-pair",
      "SecurityGroups": [
        "sg-8422d1eb"
      ],
      "LaunchConfigurationName": "my-lc-from-instance-changetype",
      "KernelId": "null",
      "RamdiskId": null,
      "InstanceType": "t2.medium",
      "AssociatePublicIpAddress": true
    }
  ]
}
```

Changing the Launch Configuration for an Auto Scaling Group

An Auto Scaling group is associated with one launch configuration at a time, and you can't modify a launch configuration after you've created it. To change the launch configuration for an Auto Scaling group, use an existing launch configuration as the basis for a new launch configuration. Then, update the Auto Scaling group to use the new launch configuration.

After you change the launch configuration for an Auto Scaling group, any new instances are launched using the new configuration options, but existing instances are not affected.

To change the launch configuration for an Auto Scaling group (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Launch Configurations**.
3. Select the launch configuration and choose **Actions, Copy launch configuration**. This sets up a new launch configuration with the same options as the original, but with "Copy" added to the name.
4. On the **Copy Launch Configuration** page, edit the configuration options as needed and choose **Create launch configuration**.
5. On the confirmation page, choose **View your Auto Scaling groups**.
6. Select the Auto Scaling group and choose **Details, Edit**.
7. Select the new launch configuration from **Launch Configuration** and choose **Save**.

To change the launch configuration for an Auto Scaling group (AWS CLI)

1. Describe the current launch configuration using the [describe-launch-configurations](#) command.
2. Create a new launch configuration using the [create-launch-configuration](#) command.
3. Update the launch configuration for the Auto Scaling group using the [update-auto-scaling-group](#) command with the `--launch-configuration-names` parameter.

To change the launch configuration for an Auto Scaling group (Tools for Windows PowerShell)

1. Describe the current launch configuration using the [Get-ASLaunchConfiguration](#) command.
2. Create a new launch configuration using the [New-ASLaunchConfiguration](#) command.
3. Update the launch configuration for the Auto Scaling group using the [Update-ASAutoScalingGroup](#) command with the `-LaunchConfigurationName` parameter.

Launching Auto Scaling Instances in a VPC

Amazon Virtual Private Cloud (Amazon VPC) enables you to define a virtual networking environment in a private, isolated section of the AWS Cloud. You have complete control over your virtual networking environment. For more information, see the [Amazon VPC User Guide](#).

Within a virtual private cloud (VPC), you can launch AWS resources such as an Auto Scaling group. An Auto Scaling group in a VPC works essentially the same way as it does on Amazon EC2 and supports the same set of features.

A subnet in Amazon VPC is a subdivision within an Availability Zone defined by a segment of the IP address range of the VPC. Using subnets, you can group your instances based on your security and operational needs. A subnet resides entirely within the Availability Zone it was created in. You launch Auto Scaling instances within the subnets.

To enable communication between the internet and the instances in your subnets, you must create an internet gateway and attach it to your VPC. An internet gateway enables your resources within the subnets to connect to the internet through the Amazon EC2 network edge. If a subnet's traffic is routed to an internet gateway, the subnet is known as a *public* subnet. If a subnet's traffic is not routed to an internet gateway, the subnet is known as a *private* subnet. Use a public subnet for resources that must be connected to the internet, and a private subnet for resources that need not be connected to the internet.

Prerequisites

Before you can launch your Auto Scaling instances in a VPC, you must first create your VPC environment. After you create your VPC and subnets, you launch Auto Scaling instances within the subnets. The easiest way to create a VPC with one public subnet is to use the VPC wizard. For more information, see the [Amazon VPC Getting Started Guide](#).

Contents

- [Default VPC \(p. 38\)](#)
- [IP Addressing in a VPC \(p. 39\)](#)
- [Instance Placement Tenancy \(p. 39\)](#)
- [Linking EC2-Classical Instances to a VPC \(p. 40\)](#)

Default VPC

If you have created your AWS account after 2013-12-04 or you are creating your Auto Scaling group in a new region, we create a default VPC for you. Your default VPC comes with a default subnet in each Availability Zone. If you have a default VPC, your Auto Scaling group is created in the default VPC by default.

For information about default VPCs and checking whether your account comes with a default VPC, see [Your Default VPC and Subnets](#) in the *Amazon VPC Developer Guide*.

IP Addressing in a VPC

When you launch your Auto Scaling instances in a VPC, your instances are automatically assigned a private IP address in the address range of the subnet. This enables your instances to communicate with other instances in the VPC.

You can configure your launch configuration to assign public IP addresses to your instances. Assigning public IP addresses to your instances enables them to communicate with the internet or other services in AWS.

When you enable public IP addresses for your instances and launch them into a subnet that is configured to automatically assign IPv6 addresses, they receive both IPv4 and IPv6 addresses. Otherwise, they receive only IPv4 addresses. For more information, see [IPv6 Addresses](#) in the *Amazon EC2 User Guide for Linux Instances*.

Instance Placement Tenancy

Dedicated Instances are physically isolated at the host hardware level from instances that aren't dedicated and from instances that belong to other AWS accounts. When you create a VPC, by default its tenancy attribute is set to `default`. In such a VPC, you can launch instances with a tenancy value of `dedicated` so that they run as single-tenancy instances. Otherwise, they run as shared-tenancy instances by default. If you set the tenancy attribute of a VPC to `dedicated`, all instances launched in the VPC run as single-tenancy instances. For more information, see [Dedicated Instances](#) in the *Amazon VPC User Guide*. For pricing information, see the [Amazon EC2 Dedicated Instances](#) product page.

When you create a launch configuration, the default value for the instance placement tenancy is `null` and the instance tenancy is controlled by the tenancy attribute of the VPC. The following table summarizes the instance placement tenancy of the Auto Scaling instances launched in a VPC.

Launch Configuration Tenancy	VPC Tenancy = default	VPC Tenancy = dedicated
not specified	shared-tenancy instance	Dedicated Instance
default	shared-tenancy instance	Dedicated Instance
dedicated	Dedicated Instance	Dedicated Instance

You can specify the instance placement tenancy for your launch configuration as `default` or `dedicated` using the [create-launch-configuration](#) command with the `--placement-tenancy` option. For example, the following command sets the launch configuration tenancy to `dedicated`:

```
aws autoscaling create-launch-configuration --launch-configuration-name my-launch-config --placement-tenancy dedicated --image-id ...
```

You can use the following [describe-launch-configurations](#) command to verify the instance placement tenancy of the launch configuration:

```
aws autoscaling describe-launch-configurations --launch-configuration-names my-launch-config
```

The following is example output for a launch configuration that creates Dedicated Instances. The `PlacementTenancy` parameter is only part of the output for this command when you explicitly set the instance placement tenancy.

```
{
```

```
"LaunchConfigurations": [  
  {  
    "UserData": null,  
    "EbsOptimized": false,  
    "PlacementTenancy": "dedicated",  
    "LaunchConfigurationARN": "arn",  
    "InstanceMonitoring": {  
      "Enabled": true  
    },  
    "ImageId": "ami-b5a7ea85",  
    "CreatedTime": "2015-03-08T23:39:49.011Z",  
    "BlockDeviceMappings": [],  
    "KeyName": null,  
    "SecurityGroups": [],  
    "LaunchConfigurationName": "my-launch-config",  
    "KernelId": null,  
    "RamdiskId": null,  
    "InstanceType": "m3.medium"  
  }  
]
```

Linking EC2-Classic Instances to a VPC

If you are launching the instances in your Auto Scaling group in EC2-Classic, you can link them to a VPC using *ClassicLink*. ClassicLink enables you to associate one or more security groups for the VPC with the EC2-Classic instances in your Auto Scaling group. It enables communication between these linked EC2-Classic instances and instances in the VPC using private IP addresses. For more information, see [ClassicLink](#) in the *Amazon EC2 User Guide for Linux Instances*.

If you have running EC2-Classic instances in your Auto Scaling group, you can link them to a VPC with ClassicLink enabled. For more information, see [Linking an Instance to a VPC](#) in the *Amazon EC2 User Guide for Linux Instances*. Alternatively, you can update the Auto Scaling group to use a launch configuration that automatically links the EC2-Classic instances to a VPC at launch. Then, terminate the running instances and let the Auto Scaling group launch new instances that are linked to the VPC.

Link to a VPC (Console)

Use the following procedure to create a launch configuration that links EC2-Classic instances to the specified VPC and update an existing Auto Scaling group to use the launch configuration.

To link EC2-Classic instances in an Auto Scaling group to a VPC

1. Verify that the VPC has ClassicLink enabled. For more information, see [Viewing Your ClassicLink-Enabled VPCs](#) in the *Amazon EC2 User Guide for Linux Instances*.
2. Create a security group for the VPC that you are going to link EC2-Classic instances to. Add rules to control communication between the linked EC2-Classic instances and instances in the VPC.
3. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
4. On the navigation pane, choose **Launch Configurations**.
5. Choose **Create launch configuration**.
6. On the **Choose AMI** page, select an AMI.
7. On the **Choose an Instance Type** page, select an instance type, and then choose **Next: Configure details**.
8. On the **Configure details** page, do the following:
 - a. Type a name for your launch configuration.
 - b. Expand **Advanced Details**, select the **IP Address Type** that you need, and then select **Link to VPC**.

- c. For **VPC**, select the VPC with ClassicLink enabled from step 1.
 - d. For **Security Groups**, select the security group from step 2.
 - e. Choose **Skip to review**.
9. On the **Review** page, make any changes that you need, and then choose **Create launch configuration**. For **Select an existing key pair or create a new key pair**, select an option, select the acknowledgment check box (if present), and then choose **Create launch configuration**.
10. When prompted, follow the directions to create an Auto Scaling group that uses the new launch configuration. Be sure to select **Launch into EC2-Classic** for **Network**. Otherwise, choose **Cancel** and then add your launch configuration to an existing Auto Scaling group as follows:
 - a. On the navigation pane, choose **Auto Scaling Groups**.
 - b. Select your Auto Scaling group, choose **Actions**, **Edit**.
 - c. For **Launch Configuration**, select your new launch configuration and then choose **Save**.

Link to a VPC (AWS CLI)

Use the following procedure to create a launch configuration that links EC2-Classic instances to the specified VPC and update an existing Auto Scaling group to use the launch configuration.

To link EC2-Classic instances in an Auto Scaling group to a VPC

1. Verify that the VPC has ClassicLink enabled. For more information, see [Viewing Your ClassicLink-Enabled VPCs](#) in the *Amazon EC2 User Guide for Linux Instances*.
2. Create a security group for the VPC that you are going to link EC2-Classic instances to. Add rules to control communication between the linked EC2-Classic instances and instances in the VPC.
3. Create a launch configuration using the [create-launch-configuration](#) command as follows. Specify a value for *vpc_id* as the ID of the VPC with ClassicLink enabled from step 1 and for *group_id* as the security group from step 2:

```
aws autoscaling create-launch-configuration --launch-configuration-name classiclink-config \
  --image-id ami_id --instance-type instance_type \
  --classic-link-vpc-id vpc_id --classic-link-vpc-security-groups group_id
```

4. Update your existing Auto Scaling group, for example *my-asg*, with the launch configuration that you created in the previous step. Any new EC2-Classic instances launched in this Auto Scaling group are linked EC2-Classic instances. Use the [update-auto-scaling-group](#) command as follows:

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg \
  --launch-configuration-name classiclink-config
```

Alternatively, you can use this launch configuration with a new Auto Scaling group that you create using [create-auto-scaling-group](#).

Auto Scaling Groups

An *Auto Scaling group* contains a collection of Amazon EC2 instances that are treated as a logical grouping for the purposes of automatic scaling and management. An Auto Scaling group also enables you to use Amazon EC2 Auto Scaling features such as health check replacements and scaling policies. Both maintaining the number of instances in an Auto Scaling group and automatic scaling are the core functionality of the Amazon EC2 Auto Scaling service.

The size of an Auto Scaling group depends on the number of instances you set as the desired capacity. You can adjust its size to meet demand, either manually or by using automatic scaling.

An Auto Scaling group starts by launching enough instances to meet its desired capacity. It maintains this number of instances by performing periodic health checks on the instances in the group. The Auto Scaling group continues to maintain a fixed number of instances even if an instance becomes unhealthy. If an instance becomes unhealthy, the group terminates the unhealthy instance and launches another instance to replace it. For more information, see [Health Checks for Auto Scaling Instances \(p. 132\)](#).

You can use scaling policies to increase or decrease the number of instances in your group dynamically to meet changing conditions. When the scaling policy is in effect, the Auto Scaling group adjusts the desired capacity of the group, between the minimum and maximum capacity values that you specify, and launches or terminates the instances as needed. You can also scale on a schedule. For more information, see [Scaling the Size of Your Auto Scaling Group \(p. 74\)](#).

An Auto Scaling group can launch On-Demand Instances, Spot Instances, or both. You can specify multiple purchase options for your Auto Scaling group only when you configure the group to use a launch template. (We recommend that you use launch templates instead of launch configurations to make sure that you can use the latest features of Amazon EC2.)

Spot Instances provide you with access to unused Amazon EC2 capacity at steep discounts relative to On-Demand prices. For more information, see [Amazon EC2 Spot Instances](#). The key differences between Spot Instances and On-Demand Instances are that the price for Spot Instances varies based on demand, and Amazon EC2 can terminate an individual Spot Instance as the availability of, or price for, Spot Instances changes. When a Spot Instance is terminated, the Auto Scaling group attempts to launch a replacement instance to maintain the desired capacity for the group.

When instances are launched, if you specified multiple Availability Zones, the desired capacity is distributed across these Availability Zones. If a scaling action occurs, Amazon EC2 Auto Scaling automatically maintains balance across all of the Availability Zones that you specify.

This section provides detailed steps for creating an Auto Scaling group. If you're new to Auto Scaling groups, start with [Getting Started with Amazon EC2 Auto Scaling \(p. 12\)](#) to learn about the various building blocks that are used in Amazon EC2 Auto Scaling.

Contents

- [Auto Scaling Groups with Multiple Instance Types and Purchase Options \(p. 43\)](#)
- [Creating an Auto Scaling Group Using a Launch Template \(p. 51\)](#)
- [Creating an Auto Scaling Group Using a Launch Configuration \(p. 53\)](#)
- [Creating an Auto Scaling Group Using an EC2 Instance \(p. 54\)](#)
- [Creating an Auto Scaling Group Using the Amazon EC2 Launch Wizard \(p. 56\)](#)
- [Tagging Auto Scaling Groups and Instances \(p. 57\)](#)
- [Using a Load Balancer with an Auto Scaling Group \(p. 61\)](#)
- [Launching Spot Instances in Your Auto Scaling Group \(p. 68\)](#)
- [Merging Your Auto Scaling Groups into a Single Multi-Zone Group \(p. 69\)](#)
- [Deleting Your Auto Scaling Infrastructure \(p. 70\)](#)

Auto Scaling Groups with Multiple Instance Types and Purchase Options

You can launch and automatically scale a fleet of On-Demand Instances and Spot Instances within a single Auto Scaling group. In addition to receiving discounts for using Spot Instances, if you specify instance types for which you have matching Reserved Instances, your discounted rate of the regular On-Demand Instance pricing also applies. The only difference between On-Demand Instances and Reserved Instances is that you must purchase the Reserved Instances in advance. All of these factors combined help you to optimize your cost savings for Amazon EC2 instances, while making sure that you obtain the desired scale and performance for your application.

You first specify the common configuration parameters in a launch template, and choose it when you create an Auto Scaling group. When you configure the Auto Scaling group, you can:

- Specify how much On-Demand and Spot capacity to launch
- Choose one or more instance types to use (optionally overriding the default instance type specified by the launch template)
- Define how Amazon EC2 Auto Scaling should distribute your capacity across instance types within each purchasing option

You enhance availability by deploying your application across multiple instance types running in multiple Availability Zones. You can use just one instance type, but it is a best practice to use a few instance types to avoid trying to launch instances from instance pools with insufficient capacity. If the Auto Scaling group's request for Spot Instances cannot be fulfilled in one Spot Instance pool, it keeps trying in other Spot Instance pools rather than launching On-Demand Instances, so that you can leverage the cost savings of Spot Instances.

Allocation Strategies

The following allocation strategies determine how the Auto Scaling group fulfills On-Demand and Spot capacity from the possible instance types.

In each case, Amazon EC2 Auto Scaling first tries to ensure that your instances are evenly balanced across the Availability Zones that you specified. Then, it launches instance types according to the allocation strategy that is specified.

On-Demand Instances

The allocation strategy for On-Demand Instances is **prioritized**. Amazon EC2 Auto Scaling uses the order of instance types in the list of launch template overrides to determine which instance type to use first when fulfilling On-Demand capacity.

For example, you specified three launch template overrides in the following order: `c5.large`, `c4.large`, and `c3.large`. When your On-Demand Instances are launched, the Auto Scaling group fulfills On-Demand capacity by starting with `c5.large`, then `c4.large`, and then `c3.large`. If you have unused Reserved Instances for `c4.large`, you can set the priority of your instance types to give the highest priority to your Reserved Instances by specifying `c4.large` as the highest priority instance type. When a `c4.large` instance launches, you receive the Reserved Instance pricing.

Spot Instances

Amazon EC2 Auto Scaling provides two types of allocation strategies that can be used for Spot Instances:

capacity-optimized

Amazon EC2 Auto Scaling allocates your instances from the Spot Instance pool with optimal capacity for the number of instances that are launching. Deploying in this way helps you make the most efficient use of spare EC2 capacity.

With Spot Instances, pricing changes slowly over time based on long-term trends in supply and demand, but capacity fluctuates in real time. The `capacity-optimized` strategy automatically launches Spot Instances into the most available pools by looking at real-time capacity data and predicting which are the most available. This works well for workloads such as big data and analytics, image and media rendering, machine learning, and high performance computing that may have a higher cost of interruption associated with restarting work and checkpointing. By offering the possibility of fewer interruptions, the `capacity-optimized` strategy can lower the overall cost of your workload.

lowest-price

Amazon EC2 Auto Scaling allocates your instances from the number (N) of Spot Instance pools that you specify and from the pools with the lowest price per unit at the time of fulfillment.

For example, if you specify 4 instance types and 4 Availability Zones, your Auto Scaling group can potentially draw from as many as 16 different Spot Instance pools. If you specify 2 Spot pools (N=2) for the allocation strategy, your Auto Scaling group can fulfill Spot capacity from a minimum of 8 different Spot pools where the price per unit is the lowest.

To get started, we recommend choosing the `capacity-optimized` allocation strategy and specifying a few instance types that are appropriate for your application. In addition, you can define a range of Availability Zones for Amazon EC2 Auto Scaling to choose from when launching instances.

Optionally, you can specify a maximum price for your Spot Instances. If you don't specify a maximum price, the default maximum price is the On-Demand price, but you still receive the steep discounts provided by Spot Instances. These discounts are possible because of the stable Spot pricing that is made available using the new [Spot pricing model](#).

For more information about the allocation strategies for Spot Instances, see [Introducing the capacity-optimized allocation strategy for Amazon EC2 Spot Instances](#) in the AWS blog.

Controlling the Proportion of On-Demand Instances

You have full control over the proportion of instances in the Auto Scaling group that are launched as On-Demand Instances. To ensure that you always have instance capacity, you can designate a percentage of the group to launch as On-Demand Instances and, optionally, a base number of On-Demand Instances to start with. If you choose to specify a base capacity of On-Demand Instances, the Auto Scaling group ensures that this base capacity of On-Demand Instances is launched first when the group scales out. Anything beyond the base capacity uses the On-Demand percentage to determine how many On-Demand Instances and Spot Instances to launch. You can specify any number from 0 to 100 for the On-Demand percentage.

The behavior of the Auto Scaling group as it increases in size is as follows:

Example: Scaling behavior

Instances Distribution	Total Number of Running Instances Across Purchase Options			
	10	20	30	40
Example 1				

Instances Distribution	Total Number of Running Instances Across Purchase Options			
	On-Demand base	On-Demand percentage above base	Spot percentage	Total
On-Demand base: 10	10	10	10	10
On-Demand percentage above base: 50%	0	5	10	15
Spot percentage: 50%	0	5	10	15
Example 2				
On-Demand base: 0	0	0	0	0
On-Demand percentage above base: 0%	0	0	0	0
Spot percentage: 100%	10	20	30	40
Example 3				
On-Demand base: 0	0	0	0	0
On-Demand percentage above base: 60%	6	12	18	24
Spot percentage: 40%	4	8	12	16
Example 4				
On-Demand base: 0	0	0	0	0
On-Demand percentage above base: 100%	10	20	30	40
Spot percentage: 0%	0	0	0	0
Example 5				
On-Demand base: 12	10	12	12	12
On-Demand percentage above base: 0%	0	0	0	0
Spot percentage: 100%	0	8	18	28

Best Practices for Spot Instances

Before you create your Auto Scaling group to request Spot Instances, review [Spot Best Practices](#). Use these best practices when you plan your request so that you can provision the type of instances you want at the lowest possible price. We also recommend that you do the following:

- Use the default maximum price, which is the On-Demand price. You pay only the Spot price for the Spot Instances that you launch. If the Spot price is within your maximum price, whether your request is fulfilled depends on availability. For more information, see [Pricing and Savings](#) in the *Amazon EC2 User Guide for Linux Instances*.
- Create your Auto Scaling group with multiple instance types. Because prices fluctuate independently for each instance type in an Availability Zone, you can often get more compute capacity for the same price when you have instance type flexibility.
- Similarly, don't limit yourself to only the most popular instance types. Because prices adjust based on long-term demand, popular instance types (such as recently launched instance families), tend to have more price adjustments. Picking older-generation instance types that are less popular tends to result in lower costs and fewer interruptions.
- If you chose the lowest-price allocation strategy and you run a web service, specify a high number of Spot pools, for example, N=10, to reduce the impact of Spot Instance interruptions if a pool in one of the Availability Zones becomes temporarily unavailable. If you run batch processing or other non-mission critical applications, you can specify a lower number of Spot pools, for example, N=2. This helps to ensure that you provision Spot Instances from only the very lowest priced Spot pools available per Availability Zone.
- Use Spot Instance interruption notices to monitor the status of your Spot Instances. For example, you can set up a rule in Amazon CloudWatch Events that automatically sends the EC2 Spot two-minute warning to an Amazon SNS topic, an AWS Lambda function, or another target. For more information, see [Spot Instance Interruption Notices](#) in the *Amazon EC2 User Guide for Linux Instances* and the [Amazon CloudWatch Events User Guide](#).

Prerequisites

Your launch template is configured for use with an Auto Scaling group. For information, see [Creating a Launch Template for an Auto Scaling Group](#) (p. 24).

IAM users can create an Auto Scaling group using a launch template only if they have permissions to call the `ec2:RunInstances` action. For information about configuring user permissions, see [Controlling Access to Your Amazon EC2 Auto Scaling Resources](#) (p. 153).

Creating an Auto Scaling Group with Multiple Purchase Options

Follow these steps to create a fleet of On-Demand Instances and Spot Instances that you can scale. You can quickly create your fleet by using the default settings for **Instances Distribution**. Otherwise, you can modify any of the default settings. You have full control over scaling the group. You can set desired capacity to a specific number of instances, or you can take advantage of automatic scaling to adjust capacity based on actual demand.

To create an Auto Scaling group with multiple purchase options (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, choose the same AWS Region that you used when you created the launch template.
3. In the navigation pane, choose **Auto Scaling Groups**.

4. Choose **Create Auto Scaling group**.
5. Choose **Launch Template**, choose your launch template, and then choose **Next Step**.
6. On the **Configure Auto Scaling group details** page, for **Group name**, enter a name for your Auto Scaling group.
7. For **Launch template version**, choose whether the Auto Scaling group uses the default, the latest, or a specific version of the launch template when scaling out.
8. For **Fleet Composition**, choose **Combine purchase options and instances** to launch instances across multiple instance types using both On-Demand and Spot purchase options.
9. For **Instance Types**, choose the types of instances that can be launched. By default, a single instance type is specified by the chosen launch template, but it can be changed if you want to specify a different instance type. When choosing multiple instance types, the order in which you add instance types sets their priority for On-Demand Instances. The instance type at the top of the list is prioritized the highest when the Auto Scaling group launches your On-Demand capacity.
10. For **Instances Distribution**, choose whether to keep or edit the default settings. Clearing the box for **Use the default settings to get started quickly** allows you to edit the default settings.
11. If you chose to edit the default settings, provide the following information.
 - For **Maximum Spot Price**, choose **Use default** to cap your maximum Spot price at the On-Demand price. Or choose **Set your maximum price** and enter a value to specify the maximum price you are willing to pay per instance per hour for Spot Instances.
 - For **Spot Allocation Strategy**, the default is **Launch Spot Instances optimally based on the available Spot capacity per Availability Zone**. To use the lowest price strategy instead, choose **Diversify Spot Instances across your N lowest priced instance types per Availability Zone**, and then enter the number of Spot pools to use, or accept the default (2).
 - For **Optional On-Demand Base**, specify the minimum number of instances for the Auto Scaling group's initial capacity that must be fulfilled by On-Demand Instances.
 - For **On-Demand Percentage Above Base**, specify the percentages of On-Demand Instances and Spot Instances for your additional capacity beyond the optional On-Demand base amount.
12. For **Group size**, enter the initial number of instances for your Auto Scaling group.
13. For **Network**, choose a VPC for your Auto Scaling group.

Note
Launching instances using a combination of instance types and purchase options is not supported in EC2-Classic.
14. For **Subnet**, choose one or more subnets in the specified VPC. Use subnets in multiple Availability Zones for high availability. For more information about high availability with Amazon EC2 Auto Scaling, see [Distributing Instances Across Availability Zones \(p. 6\)](#).
15. (Optional) To register your Amazon EC2 instances with a load balancer, choose **Advanced Details**, choose **Receive traffic from one or more load balancers**, and choose one or more Classic Load Balancers or target groups.
16. Choose **Next: Configure scaling policies**.
17. On the **Configure scaling policies** page, choose one of the following options, and then choose **Next: Configure Notifications**:
 - To manually adjust the size of the Auto Scaling group as needed, choose **Keep this group at its initial size**. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#).
 - To automatically adjust the size of the Auto Scaling group based on criteria that you specify, choose **Use scaling policies to adjust the capacity of this group** and follow the directions. For more information, see [Configure Scaling Policies \(p. 91\)](#).
18. (Optional) To receive notifications, choose **Add notification**, configure the notification, and then choose **Next: Configure Tags**.
19. (Optional) To add tags, choose **Edit tags**, provide a tag key and value for each tag, and then choose **Review**.

Alternatively, you can add tags later on. For more information, see [Tagging Auto Scaling Groups and Instances \(p. 57\)](#).

20. On the **Review** page, choose **Create Auto Scaling group**.

21. On the **Auto Scaling group creation status** page, choose **Close**.

To create an Auto Scaling group with multiple purchase options using the command line

You can use one of the following commands:

- [create-auto-scaling-group](#) (AWS CLI)
- [New-ASAutoScalingGroup](#) (AWS Tools for Windows PowerShell)

Example: Launch Spot Instances using the capacity-optimized allocation strategy

The following [create-auto-scaling-group](#) command creates an Auto Scaling group that specifies the following:

- The percentage of the group to launch as On-Demand Instances (0) and a base number of On-Demand Instances to start with (1)
- The instance types to launch in priority order (c3.large, c4.large, c5.large)
- The subnets in which to launch the instances (subnet-5ea0c127, subnet-6194ea3b, subnet-c934b782), each corresponding to a different Availability Zone
- The launch template (my-launch-template) and the launch template version (\$Default)

When Amazon EC2 Auto Scaling attempts to fulfill your On-Demand capacity, it launches the c3.large instance type first. The Spot Instances come from the optimal Spot pool in each Availability Zone based on Spot Instance capacity.

```
aws autoscaling create-auto-scaling-group --cli-input-json file://~/config.json
```

The following is an example config.json file.

```
{
  "AutoScalingGroupName": "my-asg",
  "MixedInstancesPolicy": {
    "LaunchTemplate": {
      "LaunchTemplateSpecification": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "$Default"
      },
      "Overrides": [
        {
          "InstanceType": "c3.large"
        },
        {
          "InstanceType": "c4.large"
        },
        {
          "InstanceType": "c5.large"
        }
      ]
    },
    "InstancesDistribution": {
      "OnDemandBaseCapacity": 1,
      "OnDemandPercentageAboveBaseCapacity": 0,
      "SpotAllocationStrategy": "capacity-optimized"
    }
  }
}
```

```
    },  
    "MinSize": 1,  
    "MaxSize": 5,  
    "DesiredCapacity": 3,  
    "VPCZoneIdentifier": "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782",  
    "Tags": []  
  }  
}
```

Example: Launch Spot Instances using the lowest-price allocation strategy diversified over two pools

The following `create-auto-scaling-group` command creates an Auto Scaling group that specifies the following:

- The percentage of the group to launch as On-Demand Instances (50) without also specifying a base number of On-Demand Instances to start with
- The instance types to launch in priority order (c3.large, c4.large, c5.large)
- The subnets in which to launch the instances (subnet-5ea0c127, subnet-6194ea3b, subnet-c934b782), each corresponding to a different Availability Zone
- The launch template (my-launch-template) and the launch template version (\$Latest)

When Amazon EC2 Auto Scaling attempts to fulfill your On-Demand capacity, it launches the c3.large instance type first. For your Spot capacity, Amazon EC2 Auto Scaling attempts to launch the Spot Instances evenly across the two lowest-priced pools in each Availability Zone.

```
aws autoscaling create-auto-scaling-group --cli-input-json file://~/config.json
```

The following is an example config.json file.

```
{  
  "AutoScalingGroupName": "my-asg",  
  "MixedInstancesPolicy": {  
    "LaunchTemplate": {  
      "LaunchTemplateSpecification": {  
        "LaunchTemplateName": "my-launch-template",  
        "Version": "$Latest"  
      },  
      "Overrides": [  
        {  
          "InstanceType": "c3.large"  
        },  
        {  
          "InstanceType": "c4.large"  
        },  
        {  
          "InstanceType": "c5.large"  
        }  
      ]  
    },  
    "InstancesDistribution": {  
      "OnDemandPercentageAboveBaseCapacity": 50,  
      "SpotAllocationStrategy": "lowest-price",  
      "SpotInstancePools": 2  
    }  
  },  
  "MinSize": 1,  
  "MaxSize": 5,  
  "DesiredCapacity": 3,  
  "VPCZoneIdentifier": "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782",  
}
```



```
}
  "Tags": []
}
```

To verify that the group has launched instances

Use the following [describe-auto-scaling-groups](#) command.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-names my-asg
```

The following example response shows that the desired capacity is 3 and that the group has three running instances.

```
{
  "AutoScalingGroups": [
    {
      "AutoScalingGroupARN": "arn",
      "ServiceLinkedRoleARN": "arn",
      "TargetGroupARNs": [],
      "SuspendedProcesses": [],
      "DesiredCapacity": 3,
      "MixedInstancesPolicy": {
        "InstancesDistribution": {
          "SpotAllocationStrategy": "lowest-price",
          "OnDemandPercentageAboveBaseCapacity": 50,
          "OnDemandAllocationStrategy": "prioritized",
          "SpotInstancePools": 2,
          "OnDemandBaseCapacity": 0
        },
        "LaunchTemplate": {
          "LaunchTemplateSpecification": {
            "LaunchTemplateName": "my-launch-template",
            "Version": "$Latest",
            "LaunchTemplateId": "lt-050555ad16a3f9c7f"
          },
          "Overrides": [
            {
              "InstanceType": "c3.large"
            },
            {
              "InstanceType": "c4.large"
            },
            {
              "InstanceType": "c5.large"
            }
          ]
        }
      },
      "EnabledMetrics": [],
      "Tags": [],
      "AutoScalingGroupName": "my-asg",
      "DefaultCooldown": 300,
      "MinSize": 1,
      "Instances": [
        {
          "ProtectedFromScaleIn": false,
          "AvailabilityZone": "us-west-2a",
          "LaunchTemplate": {
            "LaunchTemplateName": "my-launch-template",
            "Version": "1",
            "LaunchTemplateId": "lt-050555ad16a3f9c7f"
          },
          "InstanceId": "i-0aae8709d49eeba4f",
          "HealthStatus": "Healthy",

```

```
        "LifecycleState": "InService"
      },
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2b",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-0c43f6003841d2d2b",
        "HealthStatus": "Healthy",
        "LifecycleState": "InService"
      },
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2c",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-0feb4cd6677d39903",
        "HealthStatus": "Healthy",
        "LifecycleState": "InService"
      }
    ],
    "MaxSize": 5,
    "VPCZoneIdentifier": "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782",
    "HealthCheckGracePeriod": 0,
    "TerminationPolicies": [
      "Default"
    ],
    "LoadBalancerNames": [],
    "CreatedTime": "2019-02-17T02:29:12.853Z",
    "AvailabilityZones": [
      "us-west-2a",
      "us-west-2b",
      "us-west-2c"
    ],
    "HealthCheckType": "EC2",
    "NewInstancesProtectedFromScaleIn": false
  }
}
```

Creating an Auto Scaling Group Using a Launch Template

When you create an Auto Scaling group, you must specify the necessary information to configure the Amazon EC2 instances, the subnets for the instances, and the initial number of instances.

To configure Amazon EC2 instances, you can specify a launch configuration, a launch template, or an EC2 instance. The following procedure demonstrates how to create an Auto Scaling group using a launch template.

With launch templates, you choose the launch template and which specific version of the launch template the group uses to launch EC2 instances. You can change these selections anytime by updating the group. Alternatively, you can configure the Auto Scaling group to choose either the default version or the latest version of the launch template dynamically when a scale-out event occurs. For example, you

configure your Auto Scaling group to choose the default version of a launch template dynamically. To change the configuration of the EC2 instances to be launched by the group, create or designate a new default version of the launch template.

Prerequisites

- You must have created a launch template that includes the parameters required to launch an EC2 instance. For information about these parameters and the limitations that apply when creating a launch template for use with an Auto Scaling group, see [Creating a Launch Template for an Auto Scaling Group](#) (p. 24).
- You must have IAM permissions to create an Auto Scaling group using a launch template and also to create EC2 resources for the instances. For more information, see [Controlling Access to Your Amazon EC2 Auto Scaling Resources](#) (p. 153).

To create an Auto Scaling group using a launch template (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, choose the same AWS Region that you used when you created the launch template.
3. In the navigation pane, choose **Auto Scaling Groups**.
4. Choose **Create Auto Scaling group**.
5. Choose **Launch Template**, choose your launch template, and then choose **Next Step**.
6. On the **Configure Auto Scaling group details** page, for **Group name**, enter a name for your Auto Scaling group.
7. For **Launch template version**, choose whether the Auto Scaling group uses the default, the latest, or a specific version of the launch template when scaling out.
8. For **Fleet Composition**, choose **Adhere to the launch template** to use the EC2 instance type and purchase option specified in the launch template.

Note

Alternatively, to launch instances across multiple instance types using both On-Demand and Spot purchase options, choose **Combine purchase options and instances**. For more information, see [Auto Scaling Groups with Multiple Instance Types and Purchase Options](#) (p. 43).

9. For **Group size**, enter the initial number of instances for your Auto Scaling group.
10. For **Network**, choose a VPC for your Auto Scaling group.
11. For **Subnet**, choose one or more subnets in the specified VPC. Use subnets in multiple Availability Zones for high availability. For more information about high availability with Amazon EC2 Auto Scaling, see [Distributing Instances Across Availability Zones](#) (p. 6).
12. (Optional) To register your Amazon EC2 instances with a load balancer, choose **Advanced Details**, choose **Receive traffic from one or more load balancers**, and choose one or more Classic Load Balancers or target groups.
13. Choose **Next: Configure scaling policies**.
14. On the **Configure scaling policies** page, choose one of the following options, and then choose **Next: Configure Notifications**:
 - To manually adjust the size of the Auto Scaling group as needed, choose **Keep this group at its initial size**. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling](#) (p. 75).
 - To automatically adjust the size of the Auto Scaling group based on criteria that you specify, choose **Use scaling policies to adjust the capacity of this group** and follow the directions. For more information, see [Configure Scaling Policies](#) (p. 91).
15. (Optional) To receive notifications, choose **Add notification**, configure the notification, and then choose **Next: Configure Tags**.

16. (Optional) To add tags, choose **Edit tags**, provide a tag key and value for each tag, and then choose **Review**.

Alternatively, you can add tags later on. For more information, see [Tagging Auto Scaling Groups and Instances \(p. 57\)](#).

17. On the **Review** page, choose **Create Auto Scaling group**.
18. On the **Auto Scaling group creation status** page, choose **Close**.

To create an Auto Scaling group using the command line

You can use one of the following commands:

- [create-auto-scaling-group](#) (AWS CLI)
- [New-ASAutoScalingGroup](#) (AWS Tools for Windows PowerShell)

Creating an Auto Scaling Group Using a Launch Configuration

When you create an Auto Scaling group, you must specify the necessary information to configure the Amazon EC2 instances, the subnets for the instances, and the initial number of instances.

Important

To configure the Amazon EC2 instances, you can specify a launch template, a launch configuration, or an EC2 instance. We recommend that you use a launch template to make sure that you can use the latest features of Amazon EC2. For more information, see [Launch Templates \(p. 24\)](#).

The following procedure demonstrates how to create an Auto Scaling group using a launch configuration. You cannot modify a launch configuration after it is created, but you can replace the launch configuration for an Auto Scaling group. For more information, see [Changing the Launch Configuration for an Auto Scaling Group \(p. 37\)](#).

Prerequisites

Create a launch configuration. For more information, see [Creating a Launch Configuration \(p. 32\)](#).

To create an Auto Scaling group using a launch configuration (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation bar at the top of the screen, choose the same AWS Region that you used when you created the launch configuration.
3. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
4. Choose **Create Auto Scaling group**.
5. On the **Create Auto Scaling Group** page, choose **Launch Configuration**, choose an existing launch configuration, and then choose **Next Step**.

Note

If you do not have any launch configurations, you're first prompted to create one before you can continue with the steps to create an Auto Scaling group.

6. On the **Configure Auto Scaling group details** page, do the following:
 - a. For **Group name**, enter a name for your Auto Scaling group.
 - b. For **Group size**, enter the initial number of instances for your Auto Scaling group.

- c. For **Network**, choose a VPC for your Auto Scaling group.
 - d. For **Subnet**, choose one or more subnets in your VPC. Use subnets in multiple Availability Zones for high availability. For more information about high availability with Amazon EC2 Auto Scaling, see [Distributing Instances Across Availability Zones \(p. 6\)](#).
 - e. (Optional) To register your Amazon EC2 instances with a load balancer, choose **Advanced Details**, choose **Receive traffic from one or more load balancers**, and choose one or more Classic Load Balancers or target groups.
 - f. Choose **Next: Configure scaling policies**.
7. On the **Configure scaling policies** page, choose one of the following options, and then choose **Next: Configure Notifications**:
- To manually adjust the size of the Auto Scaling group as needed, choose **Keep this group at its initial size**. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#).
 - To automatically adjust the size of the Auto Scaling group based on criteria that you specify, choose **Use scaling policies to adjust the capacity of this group** and follow the directions. For more information, see [Configure Scaling Policies \(p. 91\)](#).
8. (Optional) To receive notifications, choose **Add notification**, configure the notification, and then choose **Next: Configure Tags**.
9. (Optional) To add tags, choose **Edit tags**, provide a tag key and value for each tag, and then choose **Review**.
- Alternatively, you can add tags later on. For more information, see [Tagging Auto Scaling Groups and Instances \(p. 57\)](#).
10. On the **Review** page, choose **Create Auto Scaling group**.
11. On the **Auto Scaling group creation status** page, choose **Close**.

To create an Auto Scaling group using the command line

You can use one of the following commands:

- [create-auto-scaling-group](#) (AWS CLI)
- [New-ASAutoScalingGroup](#) (AWS Tools for Windows PowerShell)

Creating an Auto Scaling Group Using an EC2 Instance

When you create an Auto Scaling group, you must specify the necessary information to configure the Amazon EC2 instances, the subnets for the instances, and the initial number of instances.

To configure Amazon EC2 instances, you can specify a launch configuration, a launch template, or an EC2 instance. The following procedure demonstrates how to create an Auto Scaling group using an EC2 instance. To use a launch configuration or a launch template, see [Creating an Auto Scaling Group Using a Launch Configuration \(p. 53\)](#) or [Creating an Auto Scaling Group Using a Launch Template \(p. 51\)](#).

When you create an Auto Scaling group using an EC2 instance, Amazon EC2 Auto Scaling creates a launch configuration for you and associates it with the Auto Scaling group. This launch configuration has the same name as the Auto Scaling group, and it derives its attributes from the specified instance, such as AMI ID, instance type, and Availability Zone.

Limitations

The following are limitations when creating an Auto Scaling group from an EC2 instance:

- If the identified instance has tags, the tags are not copied to the `Tags` attribute of the new Auto Scaling group.
- The Auto Scaling group includes the block device mapping from the AMI used to launch the instance; it does not include any block devices attached after instance launch.
- If the identified instance is registered with one or more load balancers, the information about the load balancer is not copied to the load balancer or target group attribute of the new Auto Scaling group.

Prerequisites

Before you begin, find the ID of the EC2 instance using the Amazon EC2 console or the [describe-instances](#) command (AWS CLI). The EC2 instance must meet the following criteria:

- The instance is in the Availability Zone in which to create the Auto Scaling group.
- The instance is not a member of another Auto Scaling group.
- The instance is in the `running` state.
- The AMI used to launch the instance must still exist.

Contents

- [Create an Auto Scaling Group from an EC2 Instance \(Console\)](#) (p. 55)
- [Create an Auto Scaling Group from an EC2 Instance \(AWS CLI\)](#) (p. 55)

Create an Auto Scaling Group from an EC2 Instance (Console)

You can use the console to create an Auto Scaling group from a running EC2 instance and add the instance to the new Auto Scaling group. For more information, see [Attach EC2 Instances to Your Auto Scaling Group](#) (p. 78).

Create an Auto Scaling Group from an EC2 Instance (AWS CLI)

Use the following [create-auto-scaling-group](#) command to create an Auto Scaling group, *my-asg-from-instance*, from the EC2 instance `i-7f12e649`.

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg-from-instance \
  --instance-id i-7f12e649 --min-size 1 --max-size 2 --desired-capacity 2
```

Use the following [describe-auto-scaling-groups](#) command to describe the Auto Scaling group.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg-from-instance
```

The following example response shows that the desired capacity of the group is 2, the group has 2 running instances, and the launch configuration is also named *my-asg-from-instance*.

```
{
  "AutoScalingGroups": [
    {
      "AutoScalingGroupARN": "arn",
      "HealthCheckGracePeriod": 0,
      "SuspendedProcesses": [],
      "DesiredCapacity": 2,
```

```
{
  "Tags": [],
  "EnabledMetrics": [],
  "LoadBalancerNames": [],
  "AutoScalingGroupName": "my-asg-from-instance",
  "DefaultCooldown": 300,
  "MinSize": 1,
  "Instances": [
    {
      "InstanceId": "i-6bd79d87",
      "AvailabilityZone": "us-west-2a",
      "HealthStatus": "Healthy",
      "LifecycleState": "InService",
      "LaunchConfigurationName": "my-asg-from-instance"
    },
    {
      "InstanceId": "i-6cd79d80",
      "AvailabilityZone": "us-west-2a",
      "HealthStatus": "Healthy",
      "LifecycleState": "InService",
      "LaunchConfigurationName": "my-asg-from-instance"
    }
  ],
  "MaxSize": 2,
  "VPCZoneIdentifier": "subnet-6bea5f06",
  "TerminationPolicies": [
    "Default"
  ],
  "LaunchConfigurationName": "my-asg-from-instance",
  "CreatedTime": "2014-12-29T16:14:50.397Z",
  "AvailabilityZones": [
    "us-west-2a"
  ],
  "HealthCheckType": "EC2"
}
```

Use the following [describe-launch-configurations](#) command to describe the launch configuration *my-asg-from-instance*.

```
aws autoscaling describe-launch-configurations --launch-configuration-names my-asg-from-instance
```

Creating an Auto Scaling Group Using the Amazon EC2 Launch Wizard

You can create a launch configuration and an Auto Scaling group in a single procedure by using the Amazon EC2 launch wizard. This is useful if you're launching more than one instance, and want to create a new launch configuration and Auto Scaling group from settings you've already selected in the Amazon EC2 launch wizard. You cannot use this option to create an Auto Scaling group using an existing launch configuration.

To create a launch configuration and Auto Scaling group using the launch wizard

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. From the dashboard, choose **Launch Instance**.
3. Choose an AMI, then choose an instance type on the next page, and then choose **Next: Configure Instance Details**.

4. In **Number of instances**, enter the number of instances that you want to launch, and then choose **Launch into Auto Scaling Group**. You do not need to add any other configuration details on the page.
5. On the confirmation page, choose **Create Launch Configuration**.
6. You are switched to step 3 of the launch configuration wizard. The AMI and instance type are already selected based on the selection you made in the Amazon EC2 launch wizard. Enter a name for the launch configuration, configure any other settings as required, and then choose **Next: Add Storage**.
7. Configure any additional volumes, and then choose **Next: Configure Security Group**.
8. Create a new security group, or choose an existing group, and then choose **Review**.
9. Review the details of the launch configuration, and then choose **Create launch configuration** to choose a key pair and create the launch configuration.
10. On the **Configure Auto Scaling group details** page, the launch configuration you created is already selected for you, and the number of instances you specified in the Amazon EC2 launch wizard is populated for **Group size**. Enter a name for the group, specify a VPC and subnet (if required), and then choose **Next: Configure scaling policies**.
11. On the **Configure scaling policies** page, choose one of the following options, and then choose **Review**:
 - To manually adjust the size of the Auto Scaling group as needed, choose **Keep this group at its initial size**. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#).
 - To automatically adjust the size of the Auto Scaling group based on criteria that you specify, choose **Use scaling policies to adjust the capacity of this group** and follow the directions. For more information, see [Configure Scaling Policies \(p. 91\)](#).
12. On the **Review** page, you can optionally add tags or notifications, and edit other configuration details. When you have finished, choose **Create Auto Scaling group**.

Tagging Auto Scaling Groups and Instances

You can organize and manage your Auto Scaling groups by assigning your own metadata to each group as *tags*. You specify a *key* and a *value* for each tag. A key can be a general category, such as "project," "owner," or "environment," with specific associated values. For example, to differentiate between your testing and production environments, you could assign each Auto Scaling group a tag with a key of "environment." Use a value of "test" to indicate your test environment or "production" to indicate your production environment. We recommend that you use a consistent set of tags to assist you in tracking your Auto Scaling groups.

Additionally, you can propagate the tags from the Auto Scaling group to the Amazon EC2 instances it launches. Tagging your instances enables you to see instance cost allocation by tag in your AWS bill. For more information, see [Using Cost Allocation Tags](#) in the *AWS Billing and Cost Management User Guide*.

You can add one or more tags to an Auto Scaling group when you create it. You can also add, list, edit, or delete tags for existing Auto Scaling groups.

You can control which IAM users and groups in your account have permission to create, edit, or delete tags. For more information, see [Controlling Access to Your Amazon EC2 Auto Scaling Resources \(p. 153\)](#). Keep in mind, however, that a policy that restricts your users from performing a tagging operation on an Auto Scaling group does not prevent them from manually changing the tags on the instances after they have launched. For information about IAM policies for Amazon EC2, see [Controlling Access to Amazon EC2 Resources](#) in the *Amazon EC2 User Guide for Linux Instances*.

Important

You can also add tags to instances by specifying the tags in your launch template. However, use caution and ensure that you do not use duplicate keys for instance tags. If you do so, Amazon EC2 Auto Scaling overrides the tag value from your launch template with the value for the

same key specified by the Auto Scaling group. For information about specifying tags in a launch template, see [Creating a Launch Template for an Auto Scaling Group \(p. 24\)](#). By specifying the tags in a launch template, you can also add tags to Amazon EBS volumes on creation.

Contents

- [Tag Restrictions \(p. 58\)](#)
- [Tagging Lifecycle \(p. 58\)](#)
- [Add or Modify Tags for Your Auto Scaling Group \(p. 58\)](#)
- [Delete Tags \(p. 61\)](#)

Tag Restrictions

The following basic restrictions apply to tags:

- The maximum number of tags per resource is 50.
- The maximum number of tags that you can add or remove using a single call is 25.
- The maximum key length is 128 Unicode characters.
- The maximum value length is 256 Unicode characters.
- Tag keys and values are case-sensitive.
- Do not use the `aws:` prefix in your tag names or values, because it is reserved for AWS use. You can't edit or delete tag names or values with this prefix, and they do not count toward your limit of tags per Auto Scaling group.

Tagging Lifecycle

If you have opted to propagate tags to your Amazon EC2 instances, the tags are managed as follows:

- In most cases, when an Auto Scaling group launches instances, it adds tags to the instances during resource creation rather than after the resource is created.
 - The exception is when you use a launch configuration to launch Spot Instances. For this scenario, your Auto Scaling group adds tags while the instances are in the `Pending` lifecycle state. If you have a lifecycle hook, the tags are available when the instance enters the `Pending:Wait` lifecycle state. For more information, see [Auto Scaling Lifecycle \(p. 7\)](#). If you need the Auto Scaling group to add tags to instances as part of the same API call that launches the Spot Instances, consider migrating to launch templates. For more information, see [Launch Templates \(p. 24\)](#).
- The Auto Scaling group automatically adds a tag to the instances with a key of `aws:autoscaling:groupName` and a value of the name of the Auto Scaling group.
- When you attach existing instances, the Auto Scaling group adds the tags to the instances, overwriting any existing tags with the same tag key. In addition, it adds a tag with a key of `aws:autoscaling:groupName` and a value of the name of the Auto Scaling group.
- When you detach an instance from an Auto Scaling group, it removes only the `aws:autoscaling:groupName` tag.
- When you scale in manually or the Auto Scaling group automatically scales in, it removes all tags from the instances that are terminating.

Add or Modify Tags for Your Auto Scaling Group

When you add a tag to your Auto Scaling group, you can specify whether it should be added to instances launched in the Auto Scaling group. If you modify a tag, the updated version of the tag is added to

instances launched in the Auto Scaling group after the change. If you create or modify a tag for an Auto Scaling group, these changes are not made to instances that are already running in the Auto Scaling group.

Contents

- [Add or Modify Tags \(Console\) \(p. 59\)](#)
- [Add or Modify Tags \(AWS CLI\) \(p. 59\)](#)

Add or Modify Tags (Console)

Use the Amazon EC2 console to:

- Add tags to new Auto Scaling groups when you create them
- Add, modify, or delete tags for existing Auto Scaling groups

To tag an Auto Scaling group on creation

When you use the Amazon EC2 console to create an Auto Scaling group, you can specify tag keys and values on the **Configure Tags** page of the Create Auto Scaling Group wizard. To propagate a tag to the instances launched in the Auto Scaling group, make sure that you keep the **Tag New Instances** option for that tag selected. Otherwise, you can deselect it.

To add or modify tags for an existing Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose an existing group from the list.
4. On the **Tags** tab, choose **Add/Edit tags**. The **Add/Edit Auto Scaling Group Tags** page lists any existing tags for the Auto Scaling group.
5. To modify existing tags, edit **Key** and **Value**.
6. To add a new tag, choose **Add tag** and edit **Key** and **Value**. You can keep **Tag New Instances** selected to add the tag to the instances launched in the Auto Scaling group automatically, and deselect it otherwise.
7. When you have finished adding tags, choose **Save**.

Add or Modify Tags (AWS CLI)

The following examples show how to use the AWS CLI to add tags when you create Auto Scaling groups, and to add or modify tags for existing Auto Scaling groups.

To tag an Auto Scaling group on creation

- Use the `create-auto-scaling-group` command to create a new Auto Scaling group and add a tag, for example, `env=prod`, to the Auto Scaling group. The tag is also added to any instances launched in the Auto Scaling group.

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg \
--launch-configuration-name my-launch-config --min-size 1 --max-size 3 \
--vpc-zone-identifier "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782" \
--tags Key=env,Value=prod,PropagateAtLaunch=true
```

To create or modify tags for an existing Auto Scaling group

- Use the [create-or-update-tags](#) command to create or modify a tag. For example, the following command adds the `Name=my-asg` and `cost-center=cc123` tags. The tags are also added to any instances launched in the Auto Scaling group after this change. If a tag with either key already exists, the existing tag is replaced. The Amazon EC2 console associates the display name for each instance with the name that is specified for the `Name` key (case-sensitive).

```
aws autoscaling create-or-update-tags \
  --tags ResourceId=my-asg,ResourceType=auto-scaling-group,Key=Name,Value=my-
asg,PropagateAtLaunch=true \
  ResourceId=my-asg,ResourceType=auto-scaling-group,Key=cost-
center,Value=cc123,PropagateAtLaunch=true
```

To list all tags for an Auto Scaling group

- Use the following [describe-tags](#) command to list the tags for the specified Auto Scaling group.

```
aws autoscaling describe-tags --filters Name=auto-scaling-group,Values=my-asg
```

The following is an example response.

```
{
  "Tags": [
    {
      "ResourceType": "auto-scaling-group",
      "ResourceId": "my-asg",
      "PropagateAtLaunch": true,
      "Value": "prod",
      "Key": "env"
    }
  ]
}
```

- Alternatively, use the following [describe-auto-scaling-groups](#) command to verify that the tag is added to the Auto Scaling group.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

The following is an example response.

```
{
  "AutoScalingGroups": [
    {
      "AutoScalingGroupARN": "arn",
      "HealthCheckGracePeriod": 0,
      "SuspendedProcesses": [],
      "DesiredCapacity": 1,
      "Tags": [
        {
          "ResourceType": "auto-scaling-group",
          "ResourceId": "my-asg",
          "PropagateAtLaunch": true,
          "Value": "prod",
          "Key": "env"
        }
      ]
    }
  ],
}
```

```
        "EnabledMetrics": [],  
        "LoadBalancerNames": [],  
        "AutoScalingGroupName": "my-asg",  
        ...  
    }  
]  
}
```

Delete Tags

You can delete a tag associated with your Auto Scaling group at any time.

Contents

- [Delete Tags \(Console\) \(p. 61\)](#)
- [Delete Tags \(AWS CLI\) \(p. 61\)](#)

Delete Tags (Console)

To delete a tag

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose an existing group from the list.
4. On the **Tags** tab, choose **Add/Edit tags**. The **Add/Edit Auto Scaling Group Tags** page lists any existing tags for the Auto Scaling group.
5. Choose the delete icon next to the tag.
6. Choose **Save**.

Delete Tags (AWS CLI)

Use the **delete-tags** command to delete a tag. For example, the following command deletes a tag with a key of `env`.

```
aws autoscaling delete-tags --tags "ResourceId=my-asg,ResourceType=auto-scaling-group,Key=env"
```

You must specify the tag key, but you don't have to specify the value. If you specify a value and the value is incorrect, the tag is not deleted.

Using a Load Balancer with an Auto Scaling Group

The purpose of automatic scaling is to automatically increase the size of your Auto Scaling group when demand goes up and decrease it when demand goes down. As capacity is increased or decreased, the Amazon EC2 instances being added or removed must be registered or deregistered with a load balancer. This enables your application to automatically distribute incoming web traffic across such a dynamically changing number of instances.

Your load balancer acts as a single point of contact for all incoming web traffic to your Auto Scaling group. When an instance is added to your Auto Scaling group, it needs to register with the load

balancer or no traffic is routed to it. When an instance is removed from your Auto Scaling group, it must deregister from the load balancer or traffic continues to be routed to it.

To use a load balancer with your Auto Scaling group, create the load balancer and then attach it to the group.

You can also use an Elastic Load Balancing health check with your instances to make sure that traffic is routed only to the healthy instances. For more information, see [Adding Elastic Load Balancing Health Checks to an Auto Scaling Group \(p. 65\)](#).

Elastic Load Balancing Types

Elastic Load Balancing provides three types of load balancers that can be used with your Auto Scaling group: Classic Load Balancers, Application Load Balancers, and Network Load Balancers. With Classic Load Balancers, instances are registered with the load balancer. With Application Load Balancers and Network Load Balancers, instances are registered as targets with a target group.

When you plan to use your load balancer with an Auto Scaling group, it's not necessary to register your EC2 instances with the load balancer or target group. When you enable Elastic Load Balancing, instances that are launched by your Auto Scaling group are automatically registered with the load balancer or target group, and instances that are terminated by your Auto Scaling group are automatically deregistered from the load balancer or target group.

Classic Load Balancer

Routes and load balances either at the transport layer (TCP/SSL), or at the application layer (HTTP/HTTPS). A Classic Load Balancer supports either EC2-Classic or a VPC.

Application Load Balancer

Routes and load balances at the application layer (HTTP/HTTPS), and supports path-based routing. An Application Load Balancer can route requests to ports on one or more registered targets, such as EC2 instances, in your virtual private cloud (VPC).

Note

The Application Load Balancer target groups must have a target type of `instance`. For more information, see [Target Type](#) in the *User Guide for Application Load Balancers*.

Network Load Balancer

Routes and load balances at the transport layer (TCP/UDP Layer-4), based on address information extracted from the TCP packet header, not from packet content. Network Load Balancers can handle traffic bursts, retain the source IP of the client, and use a fixed IP for the life of the load balancer.

Note

The Network Load Balancer target groups must have a target type of `instance`. For more information, see [Target Type](#) in the *User Guide for Network Load Balancers*.

To learn more about Elastic Load Balancing, see the following topics:

- [What Is Elastic Load Balancing?](#)
- [What Is a Classic Load Balancer?](#)
- [What Is an Application Load Balancer?](#)
- [What Is a Network Load Balancer?](#)

For information about integrating Amazon EC2 Auto Scaling with Elastic Load Balancing, see the following topics:

Topics

- [Attaching a Load Balancer to Your Auto Scaling Group \(p. 63\)](#)
- [Adding Elastic Load Balancing Health Checks to an Auto Scaling Group \(p. 65\)](#)
- [Expanding Your Scaled and Load-Balanced Application to an Additional Availability Zone \(p. 66\)](#)

Attaching a Load Balancer to Your Auto Scaling Group

Amazon EC2 Auto Scaling integrates with Elastic Load Balancing to enable you to attach one or more load balancers to an existing Auto Scaling group. After you attach the load balancer, it automatically registers the instances in the group and distributes incoming traffic across the instances.

When you attach a load balancer, it enters the `Adding` state while registering the instances in the group. After all instances in the group are registered with the load balancer, it enters the `Added` state. After at least one registered instance passes the health checks, it enters the `InService` state. After the load balancer enters the `InService` state, Amazon EC2 Auto Scaling can terminate and replace any instances that are reported as unhealthy. If no registered instances pass the health checks (for example, due to a misconfigured health check), the load balancer doesn't enter the `InService` state. Amazon EC2 Auto Scaling doesn't terminate and replace the instances.

When you detach a load balancer, it enters the `Removing` state while deregistering the instances in the group. The instances remain running after they are deregistered. If connection draining is enabled, Elastic Load Balancing waits for in-flight requests to complete or for the maximum timeout to expire (whichever comes first) before deregistering the instances. By default, connection draining is enabled for Application Load Balancers but must be enabled for Classic Load Balancers. For more information, see [Connection Draining](#) in the *User Guide for Classic Load Balancers*.

Contents

- [Prerequisites \(p. 63\)](#)
- [Add a Load Balancer \(Console\) \(p. 63\)](#)
- [Add a Load Balancer \(AWS CLI\) \(p. 64\)](#)

Prerequisites

Before you begin, create a load balancer in the same AWS Region as the Auto Scaling group. Elastic Load Balancing supports three types of load balancers: Application Load Balancers, Network Load Balancers, and Classic Load Balancers. You can attach any of these types of load balancers to your Auto Scaling group. For more information, see [Elastic Load Balancing Types \(p. 62\)](#).

To configure the Auto Scaling group to use Elastic Load Balancing health check, see [Adding Elastic Load Balancing Health Checks to an Auto Scaling Group \(p. 65\)](#).

Add a Load Balancer (Console)

Use the following procedure to attach a load balancer to an existing Auto Scaling group. To attach your load balancer to your Auto Scaling group when you create the group, see [Tutorial: Set Up a Scaled and Load-Balanced Application \(p. 18\)](#).

To attach a load balancer to a group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose an existing group from the list.
4. On the **Details** tab, choose **Edit**.

5. Do one of the following:
 - a. [Classic Load Balancers] For **Load Balancers**, choose your load balancer.
 - b. [Application/Network Load Balancers] For **Target Groups**, choose your target group.
6. Choose **Save**.

When you no longer need the load balancer, use the following procedure to detach it from your Auto Scaling group.

To detach a load balancer from a group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose an existing group from the list.
4. On the **Details** tab, choose **Edit**.
5. Do one of the following:
 - a. [Classic Load Balancers] For **Load Balancers**, remove the load balancer.
 - b. [Application/Network Load Balancers] For **Target Groups**, remove the target group.
6. Choose **Save**.

Add a Load Balancer (AWS CLI)

To attach a Classic Load Balancer

Use the following [attach-load-balancers](#) command to attach the specified load balancer to your Auto Scaling group.

```
aws autoscaling attach-load-balancers --auto-scaling-group-name my-asg \  
  --load-balancer-names my-lb
```

To attach a target group for an Application Load Balancer or Network Load Balancer

Use the following [attach-load-balancer-target-groups](#) command to attach the specified target group to your Auto Scaling group.

```
aws autoscaling attach-load-balancer-target-groups --auto-scaling-group-name my-asg \  
  --target-group-arns my-targetgroup-arn
```

To detach a Classic Load Balancer

Use the following [detach-load-balancers](#) command to detach a load balancer from your Auto Scaling group if you no longer need it.

```
aws autoscaling detach-load-balancers --auto-scaling-group-name my-asg \  
  --load-balancer-names my-lb
```

To detach a target group for an Application Load Balancer or Network Load Balancer

Use the following [detach-load-balancer-target-groups](#) command to detach a target group from your Auto Scaling group if you no longer need it.

```
aws autoscaling detach-load-balancer-target-groups --auto-scaling-group-name my-asg \  
  --target-group-arns my-targetgroup-arn
```

```
--target-group-arns my-targetgroup-arn
```

Adding Elastic Load Balancing Health Checks to an Auto Scaling Group

The default health checks for an Auto Scaling group are EC2 status checks only. If an instance fails these status checks, the Auto Scaling group considers the instance unhealthy and replaces it. For more information, see [Health Checks for Auto Scaling Instances](#) (p. 132).

If you attached one or more load balancers or target groups to your Auto Scaling group, the group does not, by default, consider an instance unhealthy and replace it if it fails the load balancer health checks.

However, you can optionally configure the Auto Scaling group to use Elastic Load Balancing health checks. This ensures that the group can determine an instance's health based on additional tests provided by the load balancer. The load balancer periodically sends pings, attempts connections, or sends requests to test the EC2 instances. These tests are called health checks.

To learn more about Elastic Load Balancing health checks, see the following topics:

- [Configure Health Checks for Your Classic Load Balancer](#) in the *User Guide for Classic Load Balancers*
- [Health Checks for Your Target Groups](#) in the *User Guide for Application Load Balancers*
- [Health Checks for Your Target Groups](#) in the *User Guide for Network Load Balancers*

If you configure the Auto Scaling group to use Elastic Load Balancing health checks, it considers the instance unhealthy if it fails either the EC2 status checks or the load balancer health checks. If you attach multiple load balancers to an Auto Scaling group, all of them must report that the instance is healthy in order for it to consider the instance healthy. If one load balancer reports an instance as unhealthy, the Auto Scaling group replaces the instance, even if other load balancers report it as healthy.

The following procedures show how to add Elastic Load Balancing health checks to your Auto Scaling group.

Contents

- [Adding Health Checks \(Console\)](#) (p. 65)
- [Adding Health Checks \(AWS CLI\)](#) (p. 66)

Adding Health Checks (Console)

Use the following procedure to add an ELB health check with a grace period of 300 seconds to an Auto Scaling group with an attached load balancer.

To add health checks

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose an existing group from the list.
4. On the **Details** tab, choose **Edit**.
5. For **Health Check Type**, select **ELB**.
6. For **Health Check Grace Period**, enter 300.
7. Choose **Save**.
8. On the **Instances** tab, the **Health Status** column displays the results of the newly added health checks.

Adding Health Checks (AWS CLI)

Use the following `update-auto-scaling-group` command to create a health check with a grace period of 300 seconds.

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-lb-asg \
--health-check-type ELB --health-check-grace-period 300
```

Expanding Your Scaled and Load-Balanced Application to an Additional Availability Zone

You can take advantage of the safety and reliability of geographic redundancy by spanning your Auto Scaling group across multiple Availability Zones within a Region and then attaching a load balancer to distribute incoming traffic across those zones. Incoming traffic is distributed equally across all Availability Zones enabled for your load balancer.

Note

An Auto Scaling group can contain Amazon EC2 instances from multiple Availability Zones within the same Region. However, an Auto Scaling group can't contain instances from multiple Regions.

When one Availability Zone becomes unhealthy or unavailable, Amazon EC2 Auto Scaling launches new instances in an unaffected zone. When the unhealthy Availability Zone returns to a healthy state, Amazon EC2 Auto Scaling automatically redistributes the application instances evenly across all of the zones for your Auto Scaling group. Amazon EC2 Auto Scaling does this by attempting to launch new instances in the Availability Zone with the fewest instances. If the attempt fails, however, Amazon EC2 Auto Scaling attempts to launch in other Availability Zones until it succeeds.

You can expand the availability of your scaled and load-balanced application by adding an Availability Zone to your Auto Scaling group and then enabling that zone for your load balancer. After you've enabled the new Availability Zone, the load balancer begins to route traffic equally among all the enabled zones.

Contents

- [Add an Availability Zone \(Console\) \(p. 66\)](#)
- [Add an Availability Zone \(AWS CLI\) \(p. 67\)](#)

Add an Availability Zone (Console)

Use the following procedure to expand your Auto Scaling group to an additional subnet (EC2-VPC) or Availability Zone (EC2-Classic).

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose an existing group from the list.
4. On the **Details** tab, choose **Edit**.
5. Do one of the following:
 - [EC2-VPC] In **Subnet(s)**, choose the subnet corresponding to the Availability Zone.
 - [EC2-Classic] In **Availability Zones(s)**, choose the Availability Zone.
6. Choose **Save**.
7. On the navigation pane, choose **Load Balancers**.

8. Choose your load balancer.
9. Do one of the following:
 - [Classic Load Balancer in EC2-Classic] On the **Instances** tab, choose **Edit Availability Zones**. On the **Add and Remove Availability Zones** page, choose the Availability Zone to add.
 - [Classic Load Balancer in a VPC] On the **Instances** tab, choose **Edit Availability Zones**. On the **Add and Remove Subnets** page, for **Available subnets**, choose the add icon (+) for the subnet to add. The subnet is moved under **Selected subnets**.
 - [Application Load Balancer] On the **Description** tab, for **Availability Zones**, choose **Edit**. Choose the add icon (+) for one of the subnets for the Availability Zone to add. The subnet is moved under **Selected subnets**.
10. Choose **Save**.

Add an Availability Zone (AWS CLI)

The commands that you use depend on whether your load balancer is a Classic Load Balancer in a VPC, a Classic Load Balancer in EC2-Classic, or an Application Load Balancer.

For an Auto Scaling group with a Classic Load Balancer in a VPC

1. Add a subnet to the Auto Scaling group using the following **update-auto-scaling-group** command.

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg \
  --vpc-zone-identifier subnet-41767929 subnet-cb663da2 --min-size 2
```

2. Verify that the instances in the new subnet are ready to accept traffic from the load balancer using the following **describe-auto-scaling-groups** command.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

3. Enable the new subnet for your Classic Load Balancer using the following **attach-load-balancer-to-subnets** command.

```
aws elb attach-load-balancer-to-subnets --load-balancer-name my-lb \
  --subnets subnet-41767929
```

For an Auto Scaling group with a Classic Load Balancer in EC2-Classic

1. Add an Availability Zone to the Auto Scaling group using the following **update-auto-scaling-group** command.

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg \
  --availability-zones us-west-2a us-west-2b us-west-2c --min-size 3
```

2. Verify that the instances in the new Availability Zone are ready to accept traffic from the load balancer using the following **describe-auto-scaling-groups** command.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

3. Enable the new Availability Zone for your Classic Load Balancer using the following **enable-availability-zones-for-load-balancer** command.

```
aws elb enable-availability-zones-for-load-balancer --load-balancer-name my-lb \
  --availability-zones us-west-2c
```

For an Auto Scaling group with an Application Load Balancer

1. Add a subnet to the Auto Scaling group using the following `update-auto-scaling-group` command.

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg \
  --vpc-zone-identifier subnet-41767929 subnet-cb663da2 --min-size 2
```

2. Verify that the instances in the new subnet are ready to accept traffic from the load balancer using the following `describe-auto-scaling-groups` command.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

3. Enable the new subnet for your Application Load Balancer using the following `set-subnets` command.

```
aws elbv2 set-subnets --load-balancer-arn my-lb-arn \
  --subnets subnet-41767929 subnet-cb663da2
```

Launching Spot Instances in Your Auto Scaling Group

Spot Instances are a cost-effective choice compared to On-Demand Instances, if you can be flexible about when your applications run and if your applications can be interrupted. This topic describes how to launch only Spot Instances in your Auto Scaling Group by specifying settings in a *launch configuration* or *launch template*, rather than in the Auto Scaling group itself.

Important

You can specify the same settings that are used to launch Spot Instances as part of the settings of an Auto Scaling group. When you specify the settings as part of the Auto Scaling group, you can specify additional options. For example, you can specify whether to launch only Spot Instances, or a combination of both On-Demand Instances and Spot Instances.

For more information, see [Auto Scaling Groups with Multiple Instance Types and Purchase Options](#) (p. 43).

Before launching Spot Instances using Amazon EC2 Auto Scaling, we recommend that you become familiar with launching and managing Spot Instances using Amazon EC2. For more information, see [Spot Instances](#) in the *Amazon EC2 User Guide for Linux Instances*.

When you create a launch configuration or launch template to launch Spot Instances instead of On-Demand Instances, keep the following considerations in mind:

- **Setting your maximum price.** You set the maximum price you are willing to pay as part of the launch configuration or launch template. If the Spot price is within your maximum price, whether your request is fulfilled depends on Spot Instance capacity. You pay only the Spot price for the Spot Instances that you launch. If the price for Spot Instances rises above your maximum price for a running instance in your Auto Scaling group, Amazon EC2 terminates your instance. For more information, see [Pricing and Savings](#) in the *Amazon EC2 User Guide for Linux Instances*.
- **Changing your maximum price.** You must create a launch configuration or launch template version with the new price. With a new launch configuration, you must associate it with your Auto Scaling group. With a launch template, you can configure the Auto Scaling group to use the default template or the latest version of the template. That way, it is automatically associated with the Auto Scaling group. The existing instances continue to run as long as the maximum price specified in the launch configuration or launch template used for those instances is higher than the current Spot price.
- **Maintaining your Spot Instances.** When your Spot Instance is terminated, the Auto Scaling group attempts to launch a replacement instance to maintain the desired capacity for the group. If your

maximum price is higher than the Spot price, then it launches a Spot Instance. Otherwise (or if the request is unsuccessful), it keeps trying.

- **Balancing across Availability Zones.** If you specify multiple Availability Zones, Amazon EC2 Auto Scaling distributes the Spot requests across these Availability Zones. If your maximum price is too low in one Availability Zone for any requests to be fulfilled, Amazon EC2 Auto Scaling checks whether requests were fulfilled in the other zones. If so, Amazon EC2 Auto Scaling cancels the requests that failed and redistributes them across the Availability Zones that have requests fulfilled. If the price in an Availability Zone with no fulfilled requests drops enough that future requests succeed, Amazon EC2 Auto Scaling rebalances across all the Availability Zones. For more information, see [Rebalancing Activities](#) (p. 7).
- **Spot Instance termination.** Amazon EC2 Auto Scaling can terminate or replace Spot Instances in the same way that it can terminate or replace On-Demand Instances. For more information, see [Controlling Which Auto Scaling Instances Terminate During Scale In](#) (p. 110).
- **Spot interruption notices.** You can use Spot Instance interruption notices to monitor the status of your Spot Instances. For example, you can set up a rule in Amazon CloudWatch Events that automatically sends the EC2 Spot two-minute warning to an Amazon SNS topic, an AWS Lambda function, or another target. For more information, see [Spot Instance Interruption Notices](#) in the *Amazon EC2 User Guide for Linux Instances* and the [Amazon CloudWatch Events User Guide](#).

Merging Your Auto Scaling Groups into a Single Multi-Zone Group

To merge separate single-zone Auto Scaling groups into a single group spanning multiple Availability Zones, rezone one of the single-zone groups into a multi-zone group. Then, delete the other groups. This works for groups with or without a load balancer, as long as the new multi-zone group is in one of the same Availability Zones as the original single-zone groups.

The following examples assume that you have two identical groups in two different Availability Zones, `us-west-2a` and `us-west-2c`. These two groups share the following specifications:

- Minimum size = 2
- Maximum size = 5
- Desired capacity = 3

Merge Zones (AWS CLI)

Use the following procedure to merge `my-group-a` and `my-group-c` into a single group that covers both `us-west-2a` and `us-west-2c`.

To merge separate single-zone groups into a single multi-zone group

1. Use the following [update-auto-scaling-group](#) command to add the `us-west-2c` Availability Zone to the supported Availability Zones for `my-group-a`. Increase the maximum size of this group to allow for the instances from both single-zone groups.

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-group-a \  
--availability-zones "us-west-2a" "us-west-2c" \  
--max-size 10 --min-size 4
```

2. Use the following [set-desired-capacity](#) command to increase the size of `my-group-a`.

```
aws autoscaling set-desired-capacity --auto-scaling-group-name my-group-a \  
--desired-capacity 6
```

```
--desired-capacity 6
```

3. (Optional) Use the following [describe-auto-scaling-groups](#) command to verify that my-group-a is at its new size.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-group-a
```

4. Use the following [update-auto-scaling-group](#) command to remove the instances from my-group-c.

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-group-c \  
--min-size 0 --max-size 0
```

5. (Optional) Use the following [describe-auto-scaling-groups](#) command to verify that no instances remain in my-group-c.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-group-c
```

The following is example output.

```
{  
  "AutoScalingGroups": [  
    {  
      "AutoScalingGroupARN": "arn",  
      "HealthCheckGracePeriod": 300,  
      "SuspendedProcesses": [],  
      "DesiredCapacity": 0,  
      "Tags": [],  
      "EnabledMetrics": [],  
      "LoadBalancerNames": [],  
      "AutoScalingGroupName": "my-group-c",  
      "DefaultCooldown": 300,  
      "MinSize": 0,  
      "Instances": [],  
      "MaxSize": 0,  
      "VPCZoneIdentifier": "null",  
      "TerminationPolicies": [  
        "Default"  
      ],  
      "LaunchConfigurationName": "my-launch-config",  
      "CreatedTime": "2015-02-26T18:24:14.449Z",  
      "AvailabilityZones": [  
        "us-west-2c"  
      ],  
      "HealthCheckType": "EC2"  
    }  
  ]  
}
```

6. Use the [delete-auto-scaling-group](#) command to delete my-group-c.

```
aws autoscaling delete-auto-scaling-group --auto-scaling-group-name my-group-c
```

Deleting Your Auto Scaling Infrastructure

To completely delete your scaling infrastructure, complete the following tasks.

Tasks

- [Delete Your Auto Scaling Group](#) (p. 71)
- [\(Optional\) Delete the Launch Configuration](#) (p. 71)
- [\(Optional\) Delete the Launch Template](#) (p. 71)
- [\(Optional\) Delete the Load Balancer](#) (p. 72)
- [\(Optional\) Delete CloudWatch Alarms](#) (p. 73)

Delete Your Auto Scaling Group

When you delete an Auto Scaling group, its desired, minimum, and maximum values are set to 0. As a result, the instances are terminated. Deleting an instance also deletes any associated logs or data, and any volumes on the instance. If do not want to terminate one or more instances, you can detach them before you delete the Auto Scaling group.

To delete your Auto Scaling group (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. On the **Auto Scaling Groups** page, choose your Auto Scaling group and choose **Actions, Delete**.
4. When prompted for confirmation, choose **Yes, Delete**.

To delete your Auto Scaling group (AWS CLI)

Use the following [delete-auto-scaling-group](#) command to delete the Auto Scaling group.

```
aws autoscaling delete-auto-scaling-group --auto-scaling-group-name my-asg
```

(Optional) Delete the Launch Configuration

You can skip this step to keep the launch configuration for future use.

To delete the launch configuration (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Launch Configurations**.
3. On the **Launch Configurations** page, choose your launch configuration and choose **Actions, Delete launch configuration**.
4. When prompted for confirmation, choose **Yes, Delete**.

To delete the launch configuration (AWS CLI)

Use the following [delete-launch-configuration](#) command.

```
aws autoscaling delete-launch-configuration --launch-configuration-name my-launch-config
```

(Optional) Delete the Launch Template

You can delete your launch template or just one version of your launch template. When you delete a launch template, all its versions are deleted.

You can skip this step to keep the launch template for future use.

To delete your launch template (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, choose **Launch Templates**.
3. Select your launch template and then do one of the following:
 - Choose **Actions, Delete template**. When prompted for confirmation, choose **Delete launch template**.
 - Choose **Actions, Delete template version**. Select the version to delete and choose **Delete launch template version**.

To delete the launch template (AWS CLI)

Use the following [delete-launch-template](#) command to delete your template and all its versions.

```
aws ec2 delete-launch-template --launch-template-id lt-068f72b72934aff71
```

Alternatively, you can use the [delete-launch-template-versions](#) command to delete a specific version of a launch template.

```
aws ec2 delete-launch-template-versions --launch-template-id lt-068f72b72934aff71 --versions 1
```

(Optional) Delete the Load Balancer

Skip this step if your Auto Scaling group is not associated with an Elastic Load Balancing load balancer, or if you want to keep the load balancer for future use.

To delete your load balancer (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, choose **Load Balancers**.
3. Choose the load balancer and choose **Actions, Delete**.
4. When prompted for confirmation, choose **Yes, Delete**.

To delete your target group (console)

1. On the navigation pane, choose **Target Groups**.
2. Choose the target group and choose **Actions, Delete**.
3. When prompted for confirmation, choose **Yes**.

To delete the load balancer associated with the Auto Scaling group (AWS CLI)

For Application Load Balancers and Network Load Balancers, use the following [delete-load-balancer](#) and [delete-target-group](#) commands.

```
aws elbv2 delete-load-balancer --load-balancer-arn my-load-balancer-arn
aws elbv2 delete-target-group --target-group-arn my-target-group-arn
```

For Classic Load Balancers, use the following [delete-load-balancer](#) command.

```
aws elb delete-load-balancer --load-balancer-name my-load-balancer
```

(Optional) Delete CloudWatch Alarms

To delete any CloudWatch alarms associated with your Auto Scaling group, complete the following steps.

You can skip this step if your Auto Scaling group is not associated with any CloudWatch alarms, or if you want to keep the alarms for future use.

Note

Deleting an Auto Scaling group automatically deletes the CloudWatch alarms that Amazon EC2 Auto Scaling manages for a target tracking scaling policy.

To delete the CloudWatch alarms (console)

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. On the navigation pane, choose **Alarms**.
3. Choose the alarms and choose **Action, Delete**.
4. When prompted for confirmation, choose **Delete**.

To delete the CloudWatch alarms (AWS CLI)

Use the `delete-alarms` command. You can delete one or more alarms at a time. For example, use the following command to delete the `Step-Scaling-AlarmHigh-AddCapacity` and `Step-Scaling-AlarmLow-RemoveCapacity` alarms.

```
aws cloudwatch delete-alarms --alarm-name Step-Scaling-AlarmHigh-AddCapacity Step-Scaling-AlarmLow-RemoveCapacity
```


Scaling the Size of Your Auto Scaling Group

Scaling is the ability to increase or decrease the compute capacity of your application. Scaling starts with an event, or scaling action, which instructs an Auto Scaling group to either launch or terminate Amazon EC2 instances.

Amazon EC2 Auto Scaling provides a number of ways to adjust scaling to best meet the needs of your applications. As a result, it's important that you have a good understanding of your application. Keep the following considerations in mind:

- What role should Amazon EC2 Auto Scaling play in your application's architecture? It's common to think about automatic scaling primarily as a way to increase and decrease capacity, but it's also useful for maintaining a steady number of servers.
- What cost constraints are important to you? Because Amazon EC2 Auto Scaling uses EC2 instances, you only pay for the resources that you use. Knowing your cost constraints helps you decide when to scale your applications, and by how much.
- What metrics are important to your application? Amazon CloudWatch supports a number of different metrics that you can use with your Auto Scaling group.

Contents

- [Scaling Options \(p. 74\)](#)
- [Maintaining the Number of Instances in Your Auto Scaling Group \(p. 75\)](#)
- [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#)
- [Scheduled Scaling for Amazon EC2 Auto Scaling \(p. 85\)](#)
- [Dynamic Scaling for Amazon EC2 Auto Scaling \(p. 87\)](#)
- [Scaling Cooldowns for Amazon EC2 Auto Scaling \(p. 107\)](#)
- [Controlling Which Auto Scaling Instances Terminate During Scale In \(p. 110\)](#)
- [Amazon EC2 Auto Scaling Lifecycle Hooks \(p. 115\)](#)
- [Temporarily Removing Instances from Your Auto Scaling Group \(p. 123\)](#)
- [Suspending and Resuming Scaling Processes \(p. 127\)](#)

Scaling Options

Amazon EC2 Auto Scaling provides several ways for you to scale your Auto Scaling group.

Maintain current instance levels at all times

You can configure your Auto Scaling group to maintain a specified number of running instances at all times. To maintain the current instance levels, Amazon EC2 Auto Scaling performs a periodic health check on running instances within an Auto Scaling group. When Amazon EC2 Auto Scaling finds an

unhealthy instance, it terminates that instance and launches a new one. For more information, see [Maintaining the Number of Instances in Your Auto Scaling Group \(p. 75\)](#).

Manual scaling

Manual scaling is the most basic way to scale your resources, where you specify only the change in the maximum, minimum, or desired capacity of your Auto Scaling group. Amazon EC2 Auto Scaling manages the process of creating or terminating instances to maintain the updated capacity. For more information, see [Manual Scaling for Amazon EC2 Auto Scaling \(p. 75\)](#).

Scale based on a schedule

Scaling by schedule means that scaling actions are performed automatically as a function of time and date. This is useful when you know exactly when to increase or decrease the number of instances in your group, simply because the need arises on a predictable schedule. For more information, see [Scheduled Scaling for Amazon EC2 Auto Scaling \(p. 85\)](#).

Scale based on demand

A more advanced way to scale your resources, using scaling policies, lets you define parameters that control the scaling process. For example, you have a web application that currently runs on two instances and you want the CPU utilization of the Auto Scaling group to stay at around 50 percent when the load on the application changes. This is useful for scaling in response to changing conditions, when you don't know when those conditions will change. You can set up Amazon EC2 Auto Scaling to respond for you. For more information, see [Dynamic Scaling for Amazon EC2 Auto Scaling \(p. 87\)](#).

Predictive scaling

You can also use Amazon EC2 Auto Scaling in combination with AWS Auto Scaling to scale resources across multiple services. AWS Auto Scaling can help you maintain optimal availability and performance by combining predictive scaling and dynamic scaling (proactive and reactive approaches, respectively) together to scale your Amazon EC2 capacity faster. For more information, see the [AWS Auto Scaling User Guide](#).

Maintaining the Number of Instances in Your Auto Scaling Group

After you have created your Auto Scaling group, the Auto Scaling group starts by launching enough EC2 instances to meet its minimum capacity (or its desired capacity, if specified). If there are no other scaling conditions attached to the Auto Scaling group, the Auto Scaling group maintains this number of running instances even if an instance becomes unhealthy.

To maintain the same number of instances, Amazon EC2 Auto Scaling performs a periodic health check on running instances within an Auto Scaling group. When it finds that an instance is unhealthy, it terminates that instance and launches a new one. If you stop or terminate a running instance, the instance is considered to be unhealthy and replaced. For more information about health check replacements, see [Health Checks for Auto Scaling Instances \(p. 132\)](#).

Manual Scaling for Amazon EC2 Auto Scaling

At any time, you can change the size of an existing Auto Scaling group manually. You can either update the desired capacity of the Auto Scaling group, or update the instances that are attached to the Auto Scaling group.

Changing the Size of Your Auto Scaling Group (Console)

When you change the size of your Auto Scaling group, Amazon EC2 Auto Scaling manages the process of launching or terminating instances to maintain the new group size.

The following example assumes that you've created an Auto Scaling group with a minimum size of 1 and a maximum size of 5. Therefore, the group currently has one running instance.

To change the size of your Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Details** tab, choose **Edit**.
5. For **Desired**, increase the desired capacity by one. For example, if the current value is 1, type 2.

The desired capacity must be less than or equal to the maximum size of the group. If your new value for **Desired** is greater than **Max**, you must update **Max**.

When you are finished, choose **Save**.

Now, verify that your Auto Scaling group has launched one additional instance.

To verify that the size of your Auto Scaling group has changed

1. On the **Activity History** tab, the **Status** column shows the current status of your instance. Use the refresh button until you see the status of your instance change to **Successful**, indicating that your Auto Scaling group has successfully launched a new instance.
2. On the **Instances** tab, the **Lifecycle** column shows the state of your instances. It takes a short time for an instance to launch. After the instance starts, its state changes to **InService**. You can see that your Auto Scaling group has launched 1 new instance, and it is in the **InService** state.

Changing the Size of Your Auto Scaling Group (AWS CLI)

When you change the size of your Auto Scaling group, Amazon EC2 Auto Scaling manages the process of launching or terminating instances to maintain the new group size. The default is not to wait for the default cooldown period to complete, but you can override the default behavior and wait for the cooldown period to complete. For more information, see [Scaling Cooldowns for Amazon EC2 Auto Scaling](#) (p. 107).

The following example assumes that you've created an Auto Scaling group with a minimum size of 1 and a maximum size of 5. Therefore, the group currently has one running instance.

To change the size of your Auto Scaling group

Use the **set-desired-capacity** command to change the size of your Auto Scaling group, as shown in the following example.

```
aws autoscaling set-desired-capacity --auto-scaling-group-name my-asg \
  --desired-capacity 2
```

If you choose to honor the default cooldown period for your Auto Scaling group, you must specify the `--honor-cooldown` option as shown in the following example.

```
aws autoscaling set-desired-capacity --auto-scaling-group-name my-asg \
  --desired-capacity 2 --honor-cooldown
```

To verify the size of your Auto Scaling group

Use the [describe-auto-scaling-groups](#) command to confirm that the size of your Auto Scaling group has changed, as in the following example.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

The following is example output, with details about the group and instances launched.

```
{
  "AutoScalingGroups": [
    {
      "AutoScalingGroupARN": "arn",
      "ServiceLinkedRoleARN": "arn",
      "TargetGroupARNs": [],
      "SuspendedProcesses": [],
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "Tags": [],
      "EnabledMetrics": [],
      "LoadBalancerNames": [],
      "AutoScalingGroupName": "my-asg",
      "DefaultCooldown": 300,
      "MinSize": 1,
      "Instances": [
        {
          "ProtectedFromScaleIn": false,
          "AvailabilityZone": "us-west-2a",
          "LaunchTemplate": {
            "LaunchTemplateName": "my-launch-template",
            "Version": "1",
            "LaunchTemplateId": "lt-050555ad16a3f9c7f"
          },
          "InstanceId": "i-05b4f7d5be44822a6",
          "HealthStatus": "Healthy",
          "LifecycleState": "Pending"
        },
        {
          "ProtectedFromScaleIn": false,
          "AvailabilityZone": "us-west-2a",
          "LaunchTemplate": {
            "LaunchTemplateName": "my-launch-template",
            "Version": "1",
            "LaunchTemplateId": "lt-050555ad16a3f9c7f"
          },
          "InstanceId": "i-0c20ac468fa3049e8",
          "HealthStatus": "Healthy",
          "LifecycleState": "InService"
        }
      ],
      "MaxSize": 5,
      "VPCZoneIdentifier": "subnet-c87f2be0",
      "HealthCheckGracePeriod": 300,
      "TerminationPolicies": [
```

```
        "Default"
      ],
      "CreatedTime": "2019-03-18T23:30:42.611Z",
      "AvailabilityZones": [
        "us-west-2a"
      ],
      "HealthCheckType": "EC2",
      "NewInstancesProtectedFromScaleIn": false,
      "DesiredCapacity": 2
    }
  ]
}
```

Notice that `DesiredCapacity` shows the new value. Your Auto Scaling group has launched an additional instance.

Attach EC2 Instances to Your Auto Scaling Group

Amazon EC2 Auto Scaling provides you with an option to enable automatic scaling for one or more EC2 instances by attaching them to your existing Auto Scaling group. After the instances are attached, they become a part of the Auto Scaling group.

The instance to attach must meet the following criteria:

- The instance is in the `running` state.
- The AMI used to launch the instance must still exist.
- The instance is not a member of another Auto Scaling group.
- The instance is launched into one of the Availability Zones defined in your Auto Scaling group.
- If the Auto Scaling group has an attached load balancer, the instance and the load balancer must both be in EC2-Classic or the same VPC. If the Auto Scaling group has an attached target group, the instance and the load balancer must both be in the same VPC.

When you attach instances, the desired capacity of the group increases by the number of instances being attached. If the number of instances being attached plus the desired capacity exceeds the maximum size of the group, the request fails.

If you attach an instance to an Auto Scaling group that has an attached load balancer, the instance is registered with the load balancer. If you attach an instance to an Auto Scaling group that has an attached target group, the instance is registered with the target group.

The examples use an Auto Scaling group with the following configuration:

- Auto Scaling group name = `my-asg`
- Minimum size = 1
- Maximum size = 5
- Desired capacity = 2
- Availability Zone = `us-west-2a`

Attaching an Instance (Console)

You can attach an existing instance to an existing Auto Scaling group, or to a new Auto Scaling group as you create it.

To attach an instance to a new Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.

2. On the navigation pane, choose **Instances**.
3. Select the instance.
4. Choose **Actions, Instance Settings, Attach to Auto Scaling Group**.
5. On the **Attach to Auto Scaling Group** page, select a new Auto Scaling group, type a name for the group, and then choose **Attach**.

The new Auto Scaling group is created using a new launch configuration with the same name that you specified for the Auto Scaling group. The launch configuration gets its settings (for example, security group and IAM role) from the instance that you attached. The Auto Scaling group gets settings (for example, Availability Zone and subnet) from the instance that you attached, and has a desired capacity and maximum size of 1.

6. (Optional) To edit the settings for the Auto Scaling group, on the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**. Select the new Auto Scaling group, choose **Edit**, change the settings as needed, and then choose **Save**.

To attach an instance to an existing Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. (Optional) On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**. Select the Auto Scaling group and verify that the maximum size of the Auto Scaling group is large enough that you can add another instance. Otherwise, choose **Edit**, increase the maximum size, and then choose **Save**.
3. On the navigation pane, choose **Instances**.
4. Select the instance.
5. Choose **Actions, Instance Settings, Attach to Auto Scaling Group**.
6. On the **Attach to Auto Scaling Group** page, select an existing Auto Scaling group, select the instance, and then choose **Attach**.
7. If the instance doesn't meet the criteria, you get an error message with the details. For example, the instance might not be in the same Availability Zone as the Auto Scaling group. Choose **Close** and try again with an instance that meets the criteria.

Attaching an Instance (AWS CLI)

To attach an instance to an Auto Scaling group

1. Describe a specific Auto Scaling group using the following `describe-auto-scaling-groups` command.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-names my-asg
```

The following example response shows that the desired capacity is 2 and the group has two running instances:

```
{
  "AutoScalingGroups": [
    {
      "AutoScalingGroupARN": "arn",
      "ServiceLinkedRoleARN": "arn",
      "TargetGroupARNs": [],
      "SuspendedProcesses": [],
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      }
    }
  ]
}
```

```
    },
    "Tags": [],
    "EnabledMetrics": [],
    "LoadBalancerNames": [],
    "AutoScalingGroupName": "my-asg",
    "DefaultCooldown": 300,
    "MinSize": 1,
    "Instances": [
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2a",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-05b4f7d5be44822a6",
        "HealthStatus": "Healthy",
        "LifecycleState": "Pending"
      },
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2a",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-0c20ac468fa3049e8",
        "HealthStatus": "Healthy",
        "LifecycleState": "InService"
      }
    ],
    "MaxSize": 5,
    "VPCZoneIdentifier": "subnet-c87f2be0",
    "HealthCheckGracePeriod": 300,
    "TerminationPolicies": [
      "Default"
    ],
    "CreatedTime": "2019-03-18T23:30:42.611Z",
    "AvailabilityZones": [
      "us-west-2a"
    ],
    "HealthCheckType": "EC2",
    "NewInstancesProtectedFromScaleIn": false,
    "DesiredCapacity": 2
  ]
}
```

2. Attach an instance to the Auto Scaling group using the following [attach-instances](#) command.

```
aws autoscaling attach-instances --instance-ids i-0787762faf1c28619 --auto-scaling-group-name my-asg
```

3. To verify that the instance is attached, use the following [describe-auto-scaling-groups](#) command.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-names my-asg
```

The following example response shows that the desired capacity has increased by 1 to 3, and that there is a new instance, i-0787762faf1c28619:

```
{
```

```
"AutoScalingGroups": [
  {
    "AutoScalingGroupARN": "arn",
    "ServiceLinkedRoleARN": "arn",
    "TargetGroupARNs": [],
    "SuspendedProcesses": [],
    "LaunchTemplate": {
      "LaunchTemplateName": "my-launch-template",
      "Version": "1",
      "LaunchTemplateId": "lt-050555ad16a3f9c7f"
    },
    "Tags": [],
    "EnabledMetrics": [],
    "LoadBalancerNames": [],
    "AutoScalingGroupName": "my-asg",
    "DefaultCooldown": 300,
    "MinSize": 1,
    "Instances": [
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2a",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-05b4f7d5be44822a6",
        "HealthStatus": "Healthy",
        "LifecycleState": "Pending"
      },
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2a",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-0c20ac468fa3049e8",
        "HealthStatus": "Healthy",
        "LifecycleState": "InService"
      },
      {
        "ProtectedFromScaleIn": false,
        "AvailabilityZone": "us-west-2a",
        "LaunchTemplate": {
          "LaunchTemplateName": "my-launch-template",
          "Version": "1",
          "LaunchTemplateId": "lt-050555ad16a3f9c7f"
        },
        "InstanceId": "i-0787762faf1c28619",
        "HealthStatus": "Healthy",
        "LifecycleState": "InService"
      }
    ],
    "MaxSize": 5,
    "VPCZoneIdentifier": "subnet-c87f2be0",
    "HealthCheckGracePeriod": 300,
    "TerminationPolicies": [
      "Default"
    ],
    "CreatedTime": "2019-03-18T23:30:42.611Z",
    "AvailabilityZones": [
      "us-west-2a"
    ],
    "HealthCheckType": "EC2",
```



```
        "NewInstancesProtectedFromScaleIn": false,  
        "DesiredCapacity": 3  
    }  
]  
}
```

Detach EC2 Instances from Your Auto Scaling Group

You can remove an instance from an Auto Scaling group. After the instances are detached, you can manage them independently from the rest of the Auto Scaling group. By detaching an instance, you can:

- Move an instance out of one Auto Scaling group and attach it to a different group. For more information, see [Attach EC2 Instances to Your Auto Scaling Group \(p. 78\)](#).
- Test an Auto Scaling group by creating it using existing instances running your application, and then detach these instances from the Auto Scaling group when your tests are complete.

When you detach instances, you have the option of decrementing the desired capacity for the Auto Scaling group by the number of instances you are detaching. If you choose not to decrement the capacity, Amazon EC2 Auto Scaling launches new instances to replace the ones that you detached. If you decrement the capacity but detach multiple instances from the same Availability Zone, Amazon EC2 Auto Scaling can rebalance the Availability Zones unless you suspend the `AZRebalance` process. For more information, see [Scaling Processes \(p. 127\)](#).

If the number of instances that you are detaching decreases the size of the Auto Scaling group below its minimum capacity, you must decrement the minimum capacity for the group before you can detach the instances.

If you detach an instance from an Auto Scaling group that has an attached load balancer, the instance is deregistered from the load balancer. If you detach an instance from an Auto Scaling group that has an attached target group, the instance is deregistered from the target group. If connection draining is enabled for your load balancer, Amazon EC2 Auto Scaling waits for in-flight requests to complete.

The examples use an Auto Scaling group with the following configuration:

- Auto Scaling group name = `my-asg`
- Minimum size = 1
- Maximum size = 5
- Desired capacity = 4
- Availability Zone = `us-west-2a`

Detaching Instances (Console)

Use the following procedure to detach an instance from your Auto Scaling group.

To detach an instance from an existing Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Instances** tab, select the instance and choose **Actions, Detach**.
5. On the **Detach Instance** page, select the check box to launch a replacement instance, or leave it unchecked to decrement the desired capacity. Choose **Detach Instance**.

Detaching Instances (AWS CLI)

Use the following procedure to detach an instance from your Auto Scaling group.

To detach an instance from an existing Auto Scaling group

1. List the current instances using the following [describe-auto-scaling-instances](#) command.

```
aws autoscaling describe-auto-scaling-instances
```

The following example response shows that the group has four running instances:

```
{
  "AutoScalingInstances": [
    {
      "ProtectedFromScaleIn": false,
      "AvailabilityZone": "us-west-2a",
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "InstanceId": "i-05b4f7d5be44822a6",
      "AutoScalingGroupName": "my-asg",
      "HealthStatus": "HEALTHY",
      "LifecycleState": "InService"
    },
    {
      "ProtectedFromScaleIn": false,
      "AvailabilityZone": "us-west-2a",
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "InstanceId": "i-0c20ac468fa3049e8",
      "AutoScalingGroupName": "my-asg",
      "HealthStatus": "HEALTHY",
      "LifecycleState": "InService"
    },
    {
      "ProtectedFromScaleIn": false,
      "AvailabilityZone": "us-west-2a",
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "InstanceId": "i-0787762faf1c28619",
      "AutoScalingGroupName": "my-asg",
      "HealthStatus": "HEALTHY",
      "LifecycleState": "InService"
    },
    {
      "ProtectedFromScaleIn": false,
      "AvailabilityZone": "us-west-2a",
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "InstanceId": "i-0f280a4c58d319a8a",
      "AutoScalingGroupName": "my-asg",
    }
  ]
}
```

```
        "HealthStatus": "HEALTHY",  
        "LifecycleState": "InService"  
    }  
]  
}
```

2. Detach an instance and decrement the desired capacity using the following [detach-instances](#) command.

```
aws autoscaling detach-instances --instance-ids i-05b4f7d5be44822a6 \  
    --auto-scaling-group-name my-asg --should-decrement-desired-capacity
```

3. Verify that the instance is detached using the following [describe-auto-scaling-instances](#) command.

```
aws autoscaling describe-auto-scaling-instances
```

The following example response shows that there are now three running instances:

```
{  
  "AutoScalingInstances": [  
    {  
      "ProtectedFromScaleIn": false,  
      "AvailabilityZone": "us-west-2a",  
      "LaunchTemplate": {  
        "LaunchTemplateName": "my-launch-template",  
        "Version": "1",  
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"  
      },  
      "InstanceId": "i-0c20ac468fa3049e8",  
      "AutoScalingGroupName": "my-asg",  
      "HealthStatus": "HEALTHY",  
      "LifecycleState": "InService"  
    },  
    {  
      "ProtectedFromScaleIn": false,  
      "AvailabilityZone": "us-west-2a",  
      "LaunchTemplate": {  
        "LaunchTemplateName": "my-launch-template",  
        "Version": "1",  
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"  
      },  
      "InstanceId": "i-0787762faf1c28619",  
      "AutoScalingGroupName": "my-asg",  
      "HealthStatus": "HEALTHY",  
      "LifecycleState": "InService"  
    },  
    {  
      "ProtectedFromScaleIn": false,  
      "AvailabilityZone": "us-west-2a",  
      "LaunchTemplate": {  
        "LaunchTemplateName": "my-launch-template",  
        "Version": "1",  
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"  
      },  
      "InstanceId": "i-0f280a4c58d319a8a",  
      "AutoScalingGroupName": "my-asg",  
      "HealthStatus": "HEALTHY",  
      "LifecycleState": "InService"  
    }  
  ]  
}
```

Scheduled Scaling for Amazon EC2 Auto Scaling

Scaling based on a schedule allows you to set your own scaling schedule for predictable load changes. For example, every week the traffic to your web application starts to increase on Wednesday, remains high on Thursday, and starts to decrease on Friday. You can plan your scaling actions based on the predictable traffic patterns of your web application. Scaling actions are performed automatically as a function of time and date.

Note

For scaling based on predictable load changes, you can also use the predictive scaling feature of AWS Auto Scaling. For more information, see the [AWS Auto Scaling User Guide](#).

To configure your Auto Scaling group to scale based on a schedule, you create a scheduled action. The scheduled action tells Amazon EC2 Auto Scaling to perform a scaling action at specified times. To create a scheduled scaling action, you specify the start time when the scaling action should take effect, and the new minimum, maximum, and desired sizes for the scaling action. At the specified time, Amazon EC2 Auto Scaling updates the group with the values for minimum, maximum, and desired size specified by the scaling action.

You can create scheduled actions for scaling one time only or for scaling on a recurring schedule.

Considerations

When you create a scheduled action, keep the following in mind.

- The order of execution for scheduled actions is guaranteed within the same group, but not for scheduled actions across groups.
- A scheduled action generally executes within seconds. However, the action may be delayed for up to two minutes from the scheduled start time. Because actions within an Auto Scaling group are executed in the order that they are specified, scheduled actions with scheduled start times close to each other can take longer to execute.
- You can create a maximum of 125 scheduled actions per Auto Scaling group.
- A scheduled action must have a unique time value. If you attempt to schedule an activity at a time when another scaling activity is already scheduled, the call is rejected with an error message noting the conflict.
- A scheduled action does not persist in your account once it has reached its end time.
- You can temporarily disable scheduled scaling without deleting your scheduled actions. For more information, see [Suspending and Resuming Scaling Processes \(p. 127\)](#).
- Cooldown periods are not supported.
- You can also schedule scaling actions for resources beyond Amazon EC2. For more information, see [Scheduled Scaling](#) in the *Application Auto Scaling User Guide*.

Create and Manage Scheduled Actions (Console)

You can create scheduled actions that scale one time only or that scale on a recurring schedule using the console. Complete the following procedure to create a scheduled action to scale your Auto Scaling group.

To create a scheduled action for an Auto Scaling group

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.

4. On the **Scheduled Actions** tab, choose **Create Scheduled Action**.
5. On the **Create Scheduled Action** page, do the following:
 - Specify the size of the group using at least one of the following values: **Min**, **Max**, or **Desired Capacity**.
 - Choose an option for **Recurrence**. If you choose **Once**, the action is performed at the specified time. If you select **Cron**, type a cron expression that specifies when to perform the action, in UTC. If you select an option that begins with **Every**, the cron expression is created for you.
 - If you chose **Once** for **Recurrence**, specify the time for the action in **Start Time**.
 - If you specified a recurring schedule, you can specify values for **Start Time** and **End Time**. If you specify a start time, the earliest time the action is performed is at this time. If you specify an end time, the action is not performed after this time.
6. Choose **Create**.

Update a Scheduled Action

If your requirements change, you can update a scheduled action.

To update a scheduled action

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Scheduled Actions** tab, select the scheduled action.
5. Choose **Actions**, **Edit**.
6. On the **Edit Scheduled Action** page, do the following:
 - Update the size of the group as needed using **Min**, **Max**, or **Desired Capacity**.
 - Update the specified recurrence as needed.
 - Update the start and end time as needed.
 - Choose **Save**.

Delete a Scheduled Action

When you no longer need a scheduled action, you can delete it.

To delete a scheduled action

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Scheduled Actions** tab, select the scheduled action.
5. Choose **Actions**, **Delete**.
6. When prompted for confirmation, choose **Yes**, **Delete**.

Create and Manage Scheduled Actions (AWS CLI)

You can create and update scheduled actions that scale one time only or that scale on a recurring schedule using the [put-scheduled-update-group-action](#) command.

To scale one time only

You can specify a one-time schedule to automatically scale your Auto Scaling group at a certain date and time, in UTC.

- To decrease the number of running instances in your Auto Scaling group at a specific time, use the following command. At the date and time specified for `--start-time`, if the group currently has more than 1 instance, the group scales in to 1 instance.

```
aws autoscaling put-scheduled-update-group-action --scheduled-action-name my-one-time-action \
--auto-scaling-group-name my-asg --start-time "2019-05-13T08:00:00Z" --desired-
capacity 1
```

- To increase the number of running instances in your Auto Scaling group at a specific time, use the following command. At the date and time specified for `--start-time`, if the group currently has fewer than 3 instances, the group scales out to 3 instances.

```
aws autoscaling put-scheduled-update-group-action --scheduled-action-name my-one-time-action \
--auto-scaling-group-name my-asg --start-time "2019-05-12T08:00:00Z" --desired-
capacity 3
```

To scale on a recurring schedule

You can specify a recurrence schedule, in UTC, using the Unix cron syntax format. This format consists of five fields separated by white spaces: [Minute] [Hour] [Day_of_Month] [Month_of_Year] [Day_of_Week]. For more information about this format, see [Crontab](#).

Use the following `put-scheduled-update-group-action` command to create a scheduled action that runs at 00:30 hours on the first of January, June, and December each year.

```
aws autoscaling put-scheduled-update-group-action --scheduled-action-name my-recurring-action \
--auto-scaling-group-name my-asg --recurrence "30 0 1 1,6,12 *" --desired-capacity 3
```

Delete a Scheduled Action

To delete a scheduled action

Use the following `delete-scheduled-action` command.

```
aws autoscaling delete-scheduled-action --scheduled-action-name my-recurring-action
```

Dynamic Scaling for Amazon EC2 Auto Scaling

When you configure dynamic scaling, you must define how to scale in response to changing demand. For example, you have a web application that currently runs on two instances and you want the CPU utilization of the Auto Scaling group to stay at around 50 percent when the load on the application changes. This gives you extra capacity to handle traffic spikes without maintaining an excessive amount of idle resources. You can configure your Auto Scaling group to scale automatically to meet this need. The policy type determines how the scaling action is performed.

Contents

- [Scaling Policy Types \(p. 88\)](#)
- [Multiple Scaling Policies \(p. 88\)](#)
- [Target Tracking Scaling Policies for Amazon EC2 Auto Scaling \(p. 89\)](#)
- [Simple and Step Scaling Policies for Amazon EC2 Auto Scaling \(p. 93\)](#)
- [Add a Scaling Policy to an Existing Auto Scaling Group \(p. 101\)](#)
- [Scaling Based on Amazon SQS \(p. 102\)](#)
- [Deleting a Scaling Policy \(p. 106\)](#)

Scaling Policy Types

Amazon EC2 Auto Scaling supports the following types of scaling policies:

- **Target tracking scaling**—Increase or decrease the current capacity of the group based on a target value for a specific metric. This is similar to the way that your thermostat maintains the temperature of your home – you select a temperature and the thermostat does the rest.
- **Step scaling**—Increase or decrease the current capacity of the group based on a set of scaling adjustments, known as *step adjustments*, that vary based on the size of the alarm breach.
- **Simple scaling**—Increase or decrease the current capacity of the group based on a single scaling adjustment.

If you are scaling based on a utilization metric that increases or decreases proportionally to the number of instances in an Auto Scaling group, we recommend that you use target tracking scaling policies. Otherwise, we recommend that you use step scaling policies.

Multiple Scaling Policies

In most cases, a target tracking scaling policy is sufficient to configure your Auto Scaling group to scale out or scale in automatically. A target tracking scaling policy allows you to select a desired outcome and have the Auto Scaling group add and remove instances as needed to achieve that outcome.

For an advanced scaling configuration, your Auto Scaling group can have more than one scaling policy. For example, you can define one or more target tracking scaling policies, one or more step scaling policies, or both. This provides greater flexibility to cover multiple scenarios.

To illustrate how multiple policies work together, consider an application that uses an Auto Scaling group and an Amazon SQS queue to send requests to a single EC2 instance. To help ensure that the application performs at optimum levels, there are two policies that control when the Auto Scaling group should scale out. One is a target tracking policy that uses a custom metric to add and remove capacity based on the number of SQS messages in the queue. The other is a step policy that uses the Amazon CloudWatch `CPUUtilization` metric to add capacity when the instance exceeds 90 percent utilization for a specified length of time.

When there are multiple policies in force at the same time, there's a chance that each policy could instruct the Auto Scaling group to scale out (or in) at the same time. For example, it's possible that the EC2 instance could trigger the CloudWatch alarm for the `CPUUtilization` metric at the same time that the SQS queue triggers the alarm for the custom metric.

When these situations occur, Amazon EC2 Auto Scaling chooses the policy that provides the largest capacity for both scale out and scale in. Suppose, for example, that the policy for CPU utilization launches one instance, while the policy for the SQS queue launches two instances. If the scale-out

criteria for both policies are met at the same time, Amazon EC2 Auto Scaling gives precedence to the SQS queue policy. This results in the Auto Scaling group launching two instances.

The approach of giving precedence to the policy that provides the largest capacity applies even when the policies use different criteria for scaling in. For example, if one policy terminates three instances, another policy decreases the number of instances by 25 percent, and the group has eight instances at the time of scale in, Amazon EC2 Auto Scaling gives precedence to the policy that provides the largest number of instances for the group. This results in the Auto Scaling group terminating two instances (25% of 8 = 2). The intention is to prevent Amazon EC2 Auto Scaling from removing too many instances.

Target Tracking Scaling Policies for Amazon EC2 Auto Scaling

With target tracking scaling policies, you select a scaling metric and set a target value. Amazon EC2 Auto Scaling creates and manages the CloudWatch alarms that trigger the scaling policy and calculates the scaling adjustment based on the metric and the target value. The scaling policy adds or removes capacity as required to keep the metric at, or close to, the specified target value. In addition to keeping the metric close to the target value, a target tracking scaling policy also adjusts to the changes in the metric due to a changing load pattern.

For example, you can use target tracking scaling to:

- Configure a target tracking scaling policy to keep the average aggregate CPU utilization of your Auto Scaling group at 40 percent.
- Configure a target tracking scaling policy to keep the request count per target of your Elastic Load Balancing target group at 1000 for your Auto Scaling group.

We recommend that you scale on Amazon EC2 instance metrics with a 1-minute frequency because that ensures a faster response to utilization changes. Scaling on metrics with a 5-minute frequency can result in slower response times and scaling on stale metric data. By default, Amazon EC2 instances are enabled for basic monitoring, which means metric data for instances is available at 5-minute intervals. You can enable detailed monitoring to get metric data for instances at 1-minute frequency. For more information, see [Configure Monitoring for Auto Scaling Instances \(p. 137\)](#).

Considerations

Keep the following considerations in mind:

- A target tracking scaling policy assumes that it should scale out your Auto Scaling group when the specified metric is above the target value. You cannot use a target tracking scaling policy to scale out your Auto Scaling group when the specified metric is below the target value.
- You may see gaps between the target value and the actual metric data points. This is because we act conservatively by rounding up or down when determining how many instances to add or remove. This prevents us from adding an insufficient number of instances or removing too many instances. However, for smaller Auto Scaling groups with fewer instances, the utilization of the group may seem far from the target value. For example, you set a target value of 50 percent for CPU utilization and your Auto Scaling group then exceeds the target. We might determine that adding 1.5 instances will decrease the CPU utilization to close to 50 percent. Because it is not possible to add 1.5 instances, we round up and add two instances. This might decrease the CPU utilization to a value below 50 percent, but it ensures that your application has enough resources to support it. Similarly, if we determine that removing 1.5 instances increases your CPU utilization to above 50 percent, we remove just one instance.
- For larger Auto Scaling groups with more instances, the utilization is spread over a larger number of instances, in which case adding or removing instances causes less of a gap between the target value and the actual metric data points.

- To ensure application availability, the Auto Scaling group scales out proportionally to the metric as fast as it can, but scales in more gradually.
- An Auto Scaling group can have multiple scaling policies in force at the same time. For more information, see [Multiple Scaling Policies](#) (p. 88).
- You can have multiple target tracking scaling policies for an Auto Scaling group, provided that each of them uses a different metric. The intention of Amazon EC2 Auto Scaling is to always prioritize availability, so its behavior differs depending on whether the target tracking policies are ready for scale out or scale in. It will scale out the Auto Scaling group if any of the target tracking policies are ready for scale out, but will scale in only if all of the target tracking policies (with the scale-in portion enabled) are ready to scale in.
- You can disable the scale-in portion of a target tracking scaling policy. This feature provides you with the flexibility to scale in your Auto Scaling group using a different method. For example, you can use a different scaling policy type for scale in while using a target tracking scaling policy for scale out.
- Do not edit or delete the CloudWatch alarms that are configured for the target tracking scaling policy. CloudWatch alarms that are associated with your target tracking scaling policies are deleted automatically when you delete the scaling policies.
- You can also use target tracking scaling for resources beyond Amazon EC2. For more information, see [Target Tracking Scaling Policies](#) in the *Application Auto Scaling User Guide*.

Choosing Metrics

You can use the Amazon EC2 console to apply a target tracking scaling policy based on a predefined metric. Alternatively, you can use the Amazon EC2 Auto Scaling CLI or API to apply a scaling policy based on a predefined or customized metric. The following predefined metrics are available:

- `ASGAverageCPUUtilization`—Average CPU utilization of the Auto Scaling group.
- `ASGAverageNetworkIn`—Average number of bytes received on all network interfaces by the Auto Scaling group.
- `ASGAverageNetworkOut`—Average number of bytes sent out on all network interfaces by the Auto Scaling group.
- `ALBRequestCountPerTarget`—Number of requests completed per target in an Application Load Balancer target group.

You can choose other available Amazon CloudWatch metrics by specifying a customized metric.

Keep the following in mind when choosing a metric:

- Not all metrics work for target tracking. This can be important when you are specifying a customized metric. The metric must be a valid utilization metric and describe how busy an instance is. The metric value must increase or decrease proportionally to the number of instances in the Auto Scaling group. That's so the metric data can be used to proportionally scale out or in the number of instances. For example, the CPU utilization of an Auto Scaling group (that is, the Amazon EC2 metric `CPUUtilization` with the metric dimension `AutoScalingGroupName`) works, if the load on the Auto Scaling group is distributed across the instances.
- The following metrics do not work for target tracking:
 - The number of requests received by the load balancer fronting the Auto Scaling group (that is, the Elastic Load Balancing metric `RequestCount`). The number of requests received by the load balancer doesn't change based on the utilization of the Auto Scaling group.
 - Load balancer request latency (that is, the Elastic Load Balancing metric `Latency`). Request latency can increase based on increasing utilization, but doesn't necessarily change proportionally.
 - The CloudWatch SQS queue metric `ApproximateNumberOfMessagesVisible`. The number of messages in a queue may not change proportionally to the size of the Auto Scaling group that

processes messages from the queue. However, a customized metric that measures the number of messages in the queue per EC2 instance in the Auto Scaling group can work. For more information, see [Scaling Based on Amazon SQS \(p. 102\)](#).

- A target tracking scaling policy does not scale in your Auto Scaling group when the specified metric has insufficient data unless you use the `ALBRequestCountPerTarget` metric. This works because the `ALBRequestCountPerTarget` metric emits zeros for periods with no associated data, and the target tracking policy requires metric data to interpret a low utilization trend. To have your Auto Scaling group scale in to 0 instances when no requests are routed to the target group, the group's minimum capacity must be set to 0.

Create an Auto Scaling Group with a Target Tracking Scaling Policy (Console)

To create an Auto Scaling group with a target tracking scaling policy

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Choose **Create Auto Scaling group**.
4. On the **Create Auto Scaling Group** page, do one of the following:
 - Select **Create an Auto Scaling group from an existing launch configuration**, select an existing launch configuration, and then choose **Next Step**.
 - If you don't have a launch configuration that you'd like to use, choose **Create a new launch configuration** and follow the directions. For more information, see [Creating a Launch Configuration \(p. 32\)](#).
5. On the **Configure Auto Scaling group details** page, do the following:
 - a. For **Group name**, type a name for your Auto Scaling group.
 - b. For **Group size**, type the desired capacity for your Auto Scaling group.
 - c. If the launch configuration specifies instances that require a VPC, such as T2 instances, you must select a VPC from **Network**. Otherwise, if your AWS account supports EC2-Classic and the instances don't require a VPC, you can select either **Launch into EC2-Classic** or a VPC.
 - d. If you selected a VPC in the previous step, select one or more subnets from **Subnet**. If you selected EC2-Classic in the previous step, select one or more Availability Zones from **Availability Zone(s)**.
 - e. Choose **Next: Configure scaling policies**.
6. On the **Configure scaling policies** page, do the following:
 - a. Select **Use scaling policies to adjust the capacity of this group**.
 - b. Specify the minimum and maximum size for your Auto Scaling group using the row that begins with **Scale between**. For example, if your group is already at its maximum size, specify a new maximum in order to scale out.
 - c. For **Scale Group Size**, specify a scaling policy. You can optionally specify a name for the policy, and then choose a value for **Metric type**.
 - d. Specify a **Target value** for the metric.
 - e. Specify an instance warm-up value for **Instances need**, which allows you to control the time until a newly launched instance can contribute to the CloudWatch metrics.
 - f. Check the **Disable scale-in** option to create only a scale-out policy. This allows you to create a separate scale-in policy of a different type if wanted.
 - g. Choose **Review**.
 - h. On the **Review** page, choose **Create Auto Scaling group**.

Instance Warmup

You can specify the number of seconds that it takes for a newly launched instance to warm up. Until its specified warm-up time has expired, an instance is not counted toward the aggregated metrics of the Auto Scaling group.

While scaling out, we do not consider instances that are warming up as part of the current capacity of the group. This ensures that we don't add more instances than you need.

While scaling in, we consider instances that are terminating as part of the current capacity of the group. Therefore, we don't remove more instances from the Auto Scaling group than necessary.

A scale-in activity can't start while a scale-out activity is in progress.

Create a Target Tracking Scaling Policy (AWS CLI)

Use the AWS CLI as follows to configure target tracking scaling policies for your Auto Scaling group.

Tasks

- [Step 1: Create an Auto Scaling Group \(p. 92\)](#)
- [Step 2: Create a Target Tracking Scaling Policy \(p. 92\)](#)

Step 1: Create an Auto Scaling Group

Use the `create-auto-scaling-group` command to create an Auto Scaling group named `my-asg` using the launch configuration `my-launch-config`. If you don't have a launch configuration that you'd like to use, you can create one. For more information, see [create-launch-configuration](#).

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg \
  --launch-configuration-name my-launch-config \
  --vpc-zone-identifier "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782" \
  --max-size 5 --min-size 1
```

Step 2: Create a Target Tracking Scaling Policy

You can create a target tracking scaling policy that tells the Auto Scaling group to increase and decrease the number of running EC2 instances in the group dynamically when the load on the application changes.

Example: target tracking configuration file

The following is an example of a target tracking configuration file, which you should save as `config.json`:

```
{
  "TargetValue": 40.0,
  "PredefinedMetricSpecification":
    {
      "PredefinedMetricType": "ASGAverageCPUUtilization"
    }
}
```

Alternatively, you can customize the metric used for scaling by creating a customized metric specification and adding values for each parameter from CloudWatch. The following is an example target tracking configuration that keeps the average utilization of the specified metric at 40 percent.

```
{
  "TargetValue":40.0,
  "CustomizedMetricSpecification":{
    "MetricName":"MyUtilizationMetric",
    "Namespace":"MyNamespace",
    "Dimensions":[
      {
        "Name":"MyOptionalMetricDimensionName",
        "Value":"MyOptionalMetricDimensionValue"
      }
    ],
    "Statistic":"Average",
    "Unit":"Percent"
  }
}
```

Example: cpu40-target-tracking-scaling-policy

Use the [put-scaling-policy](#) command, along with the `config.json` file that you created previously, to create a scaling policy named `cpu40-target-tracking-scaling-policy` that keeps the average CPU utilization of the Auto Scaling group at 40 percent:

```
aws autoscaling put-scaling-policy --policy-name cpu40-target-tracking-scaling-policy \
  --auto-scaling-group-name my-asg --policy-type TargetTrackingScaling \
  --target-tracking-configuration file://config.json
```

Simple and Step Scaling Policies for Amazon EC2 Auto Scaling

With simple and step scaling policies, you choose scaling metrics and threshold values for the CloudWatch alarms that trigger the scaling process as well as define how your Auto Scaling group should be scaled when a threshold is in breach for a specified number of evaluation periods.

We recommend that you use step scaling policies instead of simple scaling policies, even if you have a single scaling adjustment. Amazon EC2 Auto Scaling originally supported only simple scaling policies. If you created your scaling policy before target tracking and step policies were introduced, your policy is treated as a simple scaling policy. For more information, see [Simple Scaling Policies and Cooldowns \(p. 94\)](#).

If your scaling metric is a utilization metric that increases or decreases proportionally to the number of instances in the Auto Scaling group, we recommend that you use a target tracking scaling policy. For more information, see [Target Tracking Scaling Policies for Amazon EC2 Auto Scaling \(p. 89\)](#). You still have the option to use target tracking scaling with step scaling for a more advanced scaling policy configuration. For example, if you want, you can configure a more aggressive response when utilization reaches a certain level.

Step Scaling Policies

Step scaling policies increase or decrease the current capacity of your Auto Scaling group based on a set of scaling adjustments, known as *step adjustments*. The adjustments vary based on the size of the alarm breach.

After a scaling activity is started, the policy continues to respond to additional alarms, even while a scaling activity or health check replacement is in progress. Therefore, all alarms that are breached are evaluated by Amazon EC2 Auto Scaling as it receives the alarm messages.

If you are creating a policy to scale out, you can specify the estimated warm-up time that it takes for a newly launched instance to be ready to contribute to the aggregated metrics. For more information, see [Instance Warmup](#) (p. 96).

You can also use step scaling for resources beyond Amazon EC2. For more information, see [Step Scaling Policies](#) in the *Application Auto Scaling User Guide*.

Simple Scaling Policies and Cooldowns

In most cases, step scaling policies are a better choice than simple scaling policies. With simple scaling, after a scaling activity is started, the policy must wait for the scaling activity or health check replacement to complete and the cooldown period to expire before responding to additional alarms. Cooldown periods help to prevent the initiation of additional scaling activities before the effects of previous activities are visible. For more information, see [Scaling Cooldowns for Amazon EC2 Auto Scaling](#) (p. 107).

Amazon EC2 Auto Scaling does not support cooldown periods for step scaling policies. Therefore, you can't specify a cooldown period for these policies and the default cooldown period for the group doesn't apply.

Scaling Adjustment Types

When a step scaling or simple scaling policy is executed, it changes the current capacity of your Auto Scaling group using the scaling adjustment specified in the policy. A scaling adjustment can't change the capacity of the group above the maximum group size or below the minimum group size.

Amazon EC2 Auto Scaling supports the following adjustment types for step scaling and simple scaling:

- **ChangeInCapacity**—Increase or decrease the current capacity of the group by the specified number of instances. A positive value increases the capacity and a negative adjustment value decreases the capacity.

Example: If the current capacity of the group is 3 instances and the adjustment is 5, then when this policy is performed, there are 5 instances added to the group for a total of 8 instances.

- **ExactCapacity**—Change the current capacity of the group to the specified number of instances. Specify a positive value with this adjustment type.

Example: If the current capacity of the group is 3 instances and the adjustment is 5, then when this policy is performed, the capacity is set to 5 instances.

- **PercentChangeInCapacity**—Increment or decrement the current capacity of the group by the specified percentage. A positive value increases the capacity and a negative value decreases the capacity. If the resulting value is not an integer, it is rounded as follows:
 - Values greater than 1 are rounded down. For example, 12.7 is rounded to 12.
 - Values between 0 and 1 are rounded to 1. For example, .67 is rounded to 1.
 - Values between 0 and -1 are rounded to -1. For example, -.58 is rounded to -1.
 - Values less than -1 are rounded up. For example, -6.67 is rounded to -6.

Example: If the current capacity is 10 instances and the adjustment is 10 percent, then when this policy is performed, 1 instance is added to the group for a total of 11 instances.

With **PercentChangeInCapacity**, you can also specify the minimum number of instances to scale (using the **MinAdjustmentMagnitude** parameter or **Add instances in increments of at least** in the console). For example, suppose that you create a policy that adds 25 percent and you specify a minimum increment of 2 instances. If you have an Auto Scaling group with 4 instances and the scaling policy is executed, 25 percent of 4 is 1 instance. However, because you specified a minimum increment of 2, there are 2 instances added.

Step Adjustments

When you create a step scaling policy, you add one or more step adjustments that enable you to scale based on the size of the alarm breach. Each step adjustment specifies the following:

- A lower bound for the metric value
- An upper bound for the metric value
- The amount by which to scale, based on the scaling adjustment type

There are a few rules for the step adjustments for your policy:

- The ranges of your step adjustments can't overlap or have a gap.
- Only one step adjustment can have a null lower bound (negative infinity). If one step adjustment has a negative lower bound, then there must be a step adjustment with a null lower bound.
- Only one step adjustment can have a null upper bound (positive infinity). If one step adjustment has a positive upper bound, then there must be a step adjustment with a null upper bound.
- The upper and lower bound can't be null in the same step adjustment.
- If the metric value is above the breach threshold, the lower bound is inclusive and the upper bound is exclusive. If the metric value is below the breach threshold, the lower bound is exclusive and the upper bound is inclusive.

CloudWatch aggregates metric data points based on the statistic for the metric associated with your CloudWatch alarm. When the alarm is breached, the appropriate scaling policy is triggered. Amazon EC2 Auto Scaling applies the aggregation type to the most recent metric data points from CloudWatch (as opposed to the raw metric data). It compares this aggregated metric value against the upper and lower bounds defined by the step adjustments to determine which step adjustment to perform.

If you are using the AWS Management Console, you specify the upper and lower bounds as absolute values. If you are using the API or the CLI, you specify the upper and lower bounds relative to the breach threshold. For example, suppose that you have an alarm with a breach threshold of 50 and a scaling adjustment type of `PercentChangeInCapacity`. You also have scale-out and scale-in policies with the following step adjustments:

Scale out policy			
Lower bound	Upper bound	Adjustment	Metric value
0	10	0	$50 \leq \text{value} < 60$
10	20	10	$60 \leq \text{value} < 70$
20	null	30	$70 \leq \text{value} < +\text{infinity}$
Scale in policy			
Lower bound	Upper bound	Adjustment	Metric value
-10	0	0	$40 < \text{value} \leq 50$
-20	-10	-10	$30 < \text{value} \leq 40$
null	-20	-30	$-\text{infinity} < \text{value} \leq 30$

Your group has both a current capacity and a desired capacity of 10 instances. The group maintains its current and desired capacity while the aggregated metric value is greater than 40 and less than 60.

If the metric value gets to 60, the desired capacity of the group increases by 1 instance, to 11 instances, based on the second step adjustment of the scale-out policy (add 10 percent of 10 instances). After the new instance is running and its specified warm-up time has expired, the current capacity of the group increases to 11 instances. If the metric value rises to 70 even after this increase in capacity, the desired capacity of the group increases by another 3 instances, to 14 instances, based on the third step adjustment of the scale-out policy (add 30 percent of 11 instances, 3.3 instances, rounded down to 3 instances).

If the metric value gets to 40, the desired capacity of the group decreases by 1 instance, to 13 instances, based on the second step adjustment of the scale-in policy (remove 10 percent of 14 instances, 1.4 instances, rounded down to 1 instance). If the metric value falls to 30 even after this decrease in capacity, the desired capacity of the group decreases by another 3 instances, to 10 instances, based on the third step adjustment of the scale-in policy (remove 30 percent of 13 instances, 3.9 instances, rounded down to 3 instances).

Instance Warmup

With step scaling policies, you can specify the number of seconds that it takes for a newly launched instance to warm up. Until its specified warm-up time has expired, an instance is not counted toward the aggregated metrics of the Auto Scaling group. While scaling out, AWS also does not consider instances that are warming up as part of the current capacity of the group. Therefore, multiple alarm breaches that fall in the range of the same step adjustment result in a single scaling activity. This ensures that we don't add more instances than you need.

Using the example in the previous section, suppose that the metric gets to 60, and then it gets to 62 while the new instance is still warming up. The current capacity is still 10 instances, so 1 instance is added (10 percent of 10 instances). However, the desired capacity of the group is already 11 instances, so AWS does not increase the desired capacity further. If the metric gets to 70 while the new instance is still warming up, we should add 3 instances (30 percent of 10 instances). However, the desired capacity of the group is already 11, so we add only 2 instances, for a new desired capacity of 13 instances.

While scaling in, AWS considers instances that are terminating as part of the current capacity of the group. Therefore, we don't remove more instances from the Auto Scaling group than necessary.

A scale-in activity can't start while a scale-out activity is in progress.

Create an Auto Scaling Group with Step Scaling Policies (Console)

You can create a scaling policy that uses CloudWatch alarms to determine when your Auto Scaling group should scale out or scale in. Each CloudWatch alarm watches a single metric and sends messages to Amazon EC2 Auto Scaling when the metric breaches a threshold that you specify in your policy. You can use alarms to monitor any of the metrics that are sent to CloudWatch from the services in AWS that you're using. Or, you can create and monitor your own metrics.

When you create a CloudWatch alarm, you can specify an Amazon SNS topic to send an email notification to when the alarm changes state. For more information, see [Create Amazon CloudWatch Alarms](#) (p. 139).

Use the console to create an Auto Scaling group with two scaling policies: a scale-out policy that increases the capacity of the group by 30 percent, and a scale-in policy that decreases the capacity of the group to two instances.

To create an Auto Scaling group with step scaling policies

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.

3. Choose **Create Auto Scaling group**.
4. On the **Create Auto Scaling Group** page, do one of the following:
 - Select **Create an Auto Scaling group from an existing launch configuration**, select an existing launch configuration, and then choose **Next Step**.
 - If you don't have a launch configuration that you'd like to use, choose **Create a new launch configuration** and follow the directions. For more information, see [Creating a Launch Configuration \(p. 32\)](#).
5. On the **Configure Auto Scaling group details** page, do the following:
 - a. For **Group name**, type a name for your Auto Scaling group.
 - b. For **Group size**, type the desired capacity for your Auto Scaling group.
 - c. If the launch configuration specifies instances that require a VPC, such as T2 instances, you must select a VPC from **Network**. Otherwise, if your AWS account supports EC2-Classic and the instances don't require a VPC, you can select either **Launch into EC2-Classic** or a VPC.
 - d. If you selected a VPC in the previous step, select one or more subnets from **Subnet**. If you selected EC2-Classic in the previous step, select one or more Availability Zones from **Availability Zone(s)**.
 - e. Choose **Next: Configure scaling policies**.
6. On the **Configure scaling policies** page, do the following:
 - a. Select **Use scaling policies to adjust the capacity of this group**.
 - b. Specify the minimum and maximum size for your Auto Scaling group using the row that begins with **Scale between**. For example, if your group is already at its maximum size, you need to specify a new maximum in order to scale out.

Scale between and instances. These will be the minimum and maximum size of your group.
 - c. Specify your scale-out policy under **Increase Group Size**. You can optionally specify a name for the policy, then choose **Add new alarm**.
 - d. On the **Create Alarm** page, choose **create topic**. For **Send a notification to**, type a name for the SNS topic. For **With these recipients**, type one or more email addresses to receive notification. You can replace the default name for your alarm with a custom name. Next, specify the metric and the criteria for the policy. For example, you can leave the default settings for **Whenever** (Average of CPU Utilization). For **Is**, choose **>=** and type 80 percent. For **For at least**, type 1 consecutive period of 5 Minutes. Choose **Create Alarm**.

The screenshot shows the 'Create Alarm' interface. On the left, there are input fields for notification settings: 'Send a notification to' (AddCapacityNotification), 'With these recipients' (mymail@example.com), 'Whenever' (Average of CPU Utilization), 'Is' (>= 80 Percent), 'For at least' (1 consecutive period(s) of 5 Minutes), and 'Name of alarm' (AddCapacityAlarm). On the right, there is a line graph titled 'CPU Utilization Percent' showing a blue line representing 'my-test-asg' over a time period from 12/10 02:00 to 12/10 06:00. The 'Create Alarm' button is highlighted in blue at the bottom right.

- e. For **Take the action**, choose **Add**, type 30 in the next field, and then choose **percent of group**. By default, the lower bound for this step adjustment is the alarm threshold and the upper bound is null (positive infinity).

To add another step adjustment, choose **Add step**. To set a minimum number of instances to scale, update the number field in **Add instances in increments of at least 1 instance(s)**.

(Optional) We recommend that you use the default to create scaling policies with steps. To create simple scaling policies, choose **Create a simple scaling policy**. For more information, see [Simple and Step Scaling Policies for Amazon EC2 Auto Scaling \(p. 93\)](#).

Increase Group Size

Name:

Execute policy when: [AddCapacityAlarm](#) [Edit](#) [Remove](#)
breaches the alarm threshold: CPUUtilization >= 80 for 300 seconds
for the metric dimensions AutoScalingGroupName = my-asg

Take the action: percent of group <= CPUUtilization < +infinity

[Add step](#) ⓘ

Add instances in increments of at least instance(s)

Instances need: seconds to warm up after each step

[Create a simple scaling policy](#) ⓘ

- f. Specify an instance warm-up value for **Instances need**, which allows you to control the amount of time until a newly launched instance can contribute to the CloudWatch metrics.
- g. Specify your scale-in policy under **Decrease Group Size**. You can optionally specify a name for the policy, then choose **Add new alarm**.
- h. On the **Create Alarm** page, you can select the same notification that you created for the scale-out policy or create a new one for the scale-in policy. You can replace the default name for your alarm with a custom name. Keep the default settings for **Whenever** (Average of CPU Utilization). For **Is**, choose <= and type 40 percent. For **For at least**, type 1 consecutive period of 5 Minutes. Choose **Create Alarm**.
- i. For **Take the action**, choose Remove, type 2 in the next field, and then choose instances. By default, the upper bound for this step adjustment is the alarm threshold and the lower bound is null (negative infinity). To add another step adjustment, choose **Add step**.

(Optional) We recommend that you use the default to create scaling policies with steps. To create simple scaling policies, choose **Create a simple scaling policy**. For more information, see [Scaling Policy Types \(p. 88\)](#).

Decrease Group Size

Name:

Execute policy when: [DecreaseCapacityAlarm](#) [Edit](#) [Remove](#)
breaches the alarm threshold: CPUUtilization <= 40 for 300 seconds
for the metric dimensions AutoScalingGroupName = my-asg

Take the action: instances >= CPUUtilization > -infinity

[Add step](#) ⓘ

[Create a simple scaling policy](#) ⓘ

- j. Choose **Review**.
 - k. On the **Review** page, choose **Create Auto Scaling group**.
7. Use the following steps to verify the scaling policies for your Auto Scaling group.
- a. The **Auto Scaling Group creation status** page confirms that your Auto Scaling group was successfully created. Choose **View your Auto Scaling Groups**.
 - b. On the **Auto Scaling Groups** page, select the Auto Scaling group that you just created.
 - c. On the **Activity History** tab, the **Status** column shows whether your Auto Scaling group has successfully launched instances.

- d. On the **Instances** tab, the **Lifecycle** column contains the state of your instances. It takes a short time for an instance to launch. After the instance starts, its lifecycle state changes to **InService**.

The **Health Status** column shows the result of the EC2 instance health check on your instance.
- e. On the **Scaling Policies** tab, you can see the policies that you created for the Auto Scaling group.

Configure Step Scaling Policies (AWS CLI)

Use the AWS CLI as follows to configure step scaling policies for your Auto Scaling group.

Tasks

- [Step 1: Create an Auto Scaling Group \(p. 99\)](#)
- [Step 2: Create Scaling Policies \(p. 99\)](#)
- [Step 3: Create CloudWatch Alarms \(p. 101\)](#)

Step 1: Create an Auto Scaling Group

Use the following [create-auto-scaling-group](#) command to create an Auto Scaling group named `my-asg` using the launch configuration `my-launch-config`. If you don't have a launch configuration that you'd like to use, you can create one. For more information, see [create-launch-configuration](#).

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg \  
  --launch-configuration-name my-launch-config \  
  --vpc-zone-identifier "subnet-5ea0c127,subnet-6194ea3b,subnet-c934b782" \  
  --max-size 5 --min-size 1
```

Step 2: Create Scaling Policies

You can create step or simple scaling policies that tell the Auto Scaling group what to do when the load on the application changes.

Example: my-step-scaleout-policy

Use the following [put-scaling-policy](#) command to create a step scaling policy named `my-step-scaleout-policy` with an adjustment type of `PercentChangeInCapacity` that increases the capacity of the group based on the following step adjustments (assuming a CloudWatch alarm threshold of 60%):

- Increase the instance count by 10% when the value of the metric is greater than or equal to 70% but less than 80%
- Increase the instance count by 20% when the value of the metric is greater than or equal to 80% but less than 90%
- Increase the instance count by 30% when the value of the metric is greater than or equal to 90%

```
aws autoscaling put-scaling-policy \  
  --auto-scaling-group-name my-asg \  
  --policy-name my-step-scaleout-policy \  
  --policy-type StepScaling \  
  --adjustment-type PercentChangeInCapacity \  
  --metric-aggregation-type Average \  
  --step-adjustments  
  MetricIntervalLowerBound=10.0,MetricIntervalUpperBound=20.0,ScalingAdjustment=10 \  
  \
```

```
MetricIntervalLowerBound=20.0,MetricIntervalUpperBound=30.0,ScalingAdjustment=20 \
    MetricIntervalLowerBound=30.0,ScalingAdjustment=30 \
--min-adjustment-magnitude 1
```

The output includes the ARN for the policy. Store this ARN in a safe place. You need it to create CloudWatch alarms.

```
{
  "PolicyARN": "arn:aws:autoscaling:region:123456789012:scalingPolicy:4ee9e543-86b5-4121-
b53b-aa4c23b5bbcc:autoScalingGroupName/my-asg:policyName/my-step-scalein-policy"
}
```

Example: my-step-scalein-policy

Use the following [put-scaling-policy](#) command to create a step scaling policy named `my-step-scalein-policy` with an adjustment type of `ChangeInCapacity` that decreases the capacity of the group by 2 instances:

```
aws autoscaling put-scaling-policy \
  --auto-scaling-group-name my-asg \
  --policy-name my-step-scalein-policy \
  --policy-type StepScaling \
  --adjustment-type ChangeInCapacity \
  --step-adjustments MetricIntervalUpperBound=0.0,ScalingAdjustment=-2
```

The output includes the ARN that serves as a unique name for the policy. Later, you can use either the ARN or a combination of the policy name and group name to specify the policy. Store this ARN in a safe place. You need it to create CloudWatch alarms.

```
{
  "PolicyARN": "arn:aws:autoscaling:region:123456789012:scalingPolicy:ac542982-
cbeb-4294-891c-a5a941dfa787:autoScalingGroupName/my-asg:policyName/my-step-scaleout-policy"
}
```

Examples of Simple Scaling Policies

Alternatively, you can create simple scaling policies by using the following CLI commands instead of the preceding CLI commands. The output of these commands includes an ARN for each policy, which you need to create the CloudWatch alarms. Keep in mind that a cooldown period will be in place due to the use of simple scaling policies.

Example: my-simple-scaleout-policy

Use the following [put-scaling-policy](#) command to create a simple scaling policy named `my-simple-scaleout-policy` with an adjustment type of `PercentChangeInCapacity` that increases the capacity of the group by 30 percent:

```
aws autoscaling put-scaling-policy --policy-name my-simple-scaleout-policy \
  --auto-scaling-group-name my-asg --scaling-adjustment 30 \
  --adjustment-type PercentChangeInCapacity
```

Example: my-simple-scalein-policy

Use the following [put-scaling-policy](#) command to create a simple scaling policy named `my-simple-scalein-policy` with an adjustment type of `ChangeInCapacity` that decreases the capacity of the group by two instances:

```
aws autoscaling put-scaling-policy --policy-name my-simple-scalein-policy \  
  --auto-scaling-group-name my-asg --scaling-adjustment -2 \  
  --adjustment-type ChangeInCapacity
```

Step 3: Create CloudWatch Alarms

In step 2, you created scaling policies that provided instructions to the Auto Scaling group about how to scale in and scale out when the conditions that you specify change. In this step, you create alarms by identifying the metrics to watch, defining the conditions for scaling, and then associating the alarms with the scaling policies.

Example: AddCapacity

Use the following CloudWatch **put-metric-alarm** command to create an alarm that increases the size of the Auto Scaling group based on an average CPU threshold value of 60% for at least two consecutive evaluation periods of two minutes. To use your own custom metric, specify its name in `--metric-name` and its namespace in `--namespace`.

```
aws cloudwatch put-metric-alarm --alarm-name Step-Scaling-AlarmHigh-AddCapacity \  
  --metric-name CPUUtilization --namespace AWS/EC2 --statistic Average \  
  --period 120 --evaluation-periods 2 --threshold 60 \  
  --comparison-operator GreaterThanOrEqualToThreshold \  
  --dimensions "Name=AutoScalingGroupName,Value=my-asg" \  
  --alarm-actions PolicyARN
```

Example: RemoveCapacity

Use the following CloudWatch **put-metric-alarm** command to create an alarm that decreases the size of the Auto Scaling group based on average CPU threshold value of 40% for at least two consecutive evaluation periods of two minutes. To use your own custom metric, specify its name in `--metric-name` and its namespace in `--namespace`.

```
aws cloudwatch put-metric-alarm --alarm-name Step-Scaling-AlarmLow-RemoveCapacity \  
  --metric-name CPUUtilization --namespace AWS/EC2 --statistic Average \  
  --period 120 --evaluation-periods 2 --threshold 40 \  
  --comparison-operator LessThanOrEqualToThreshold \  
  --dimensions "Name=AutoScalingGroupName,Value=my-asg" \  
  --alarm-actions PolicyARN
```

Add a Scaling Policy to an Existing Auto Scaling Group

Use the console to add a scaling policy to an existing Auto Scaling group. A scaling policy is used to increase and decrease the number of running EC2 instances in the group dynamically to meet changing conditions. For more information, see [Dynamic Scaling for Amazon EC2 Auto Scaling](#) (p. 87).

To update an Auto Scaling group with scaling based on metrics

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. On the **Scaling Policies** tab, choose **Add policy**.

The page displays all the available resources.

5. If you are adding a target tracking scaling policy, use the following sub-steps. If you are using a simple or step scaling policy, skip to the next step.

- a. For **Name**, type a name for the policy.
 - b. Choose a **Metric type** and specify a **Target value** for the metric.
 - c. Specify an instance warm-up value for **Instances need**, which allows you to control the amount of time until a newly launched instance can contribute to the CloudWatch metrics.
 - d. Check the **Disable scale-in** option if you only want a scale-out policy created. This allows you to create a separate scale-in policy if wanted.
 - e. Choose **Create**.
6. If you are using a step scaling policy, choose **Create a scaling policy with steps**, and then do the following:

Note

We recommend that you use the option to create scaling policies with steps. To use a simple scaling policy, choose **Create a simple scaling policy** instead. For more information about simple and step scaling, see [Simple and Step Scaling Policies for Amazon EC2 Auto Scaling](#) (p. 93).

- a. For **Name**, type a name for the policy, and then choose **Create new alarm**.
- b. On the **Create Alarm** page, choose **create topic**. For **Send a notification to**, type a name for the SNS topic. For **With these recipients**, type one or more email addresses to receive notification. You can replace the default name for your alarm with a custom name. Next, specify the metric and the criteria for the alarm, using **Whenever**, **Is**, and **For at least**. Choose **Create Alarm**.
- c. Specify the scaling activity for the policy using **Take the action**. By default, the lower bound for this step adjustment is the alarm threshold and the upper bound is null (positive infinity). To add another step adjustment, choose **Add step**.
- d. Specify an instance warm-up value for **Instances need**, which allows you to control the amount of time until a newly launched instance can contribute to the CloudWatch metrics.
- e. Choose **Create**.

Scaling Based on Amazon SQS

This section shows you how to scale your Auto Scaling group in response to changing demand from an Amazon Simple Queue Service (Amazon SQS) queue. Amazon SQS offers a secure, durable, and available hosted queue that lets you integrate and decouple distributed software systems and components. If you're not familiar with Amazon SQS, see the [Amazon Simple Queue Service Developer Guide](#) for more information.

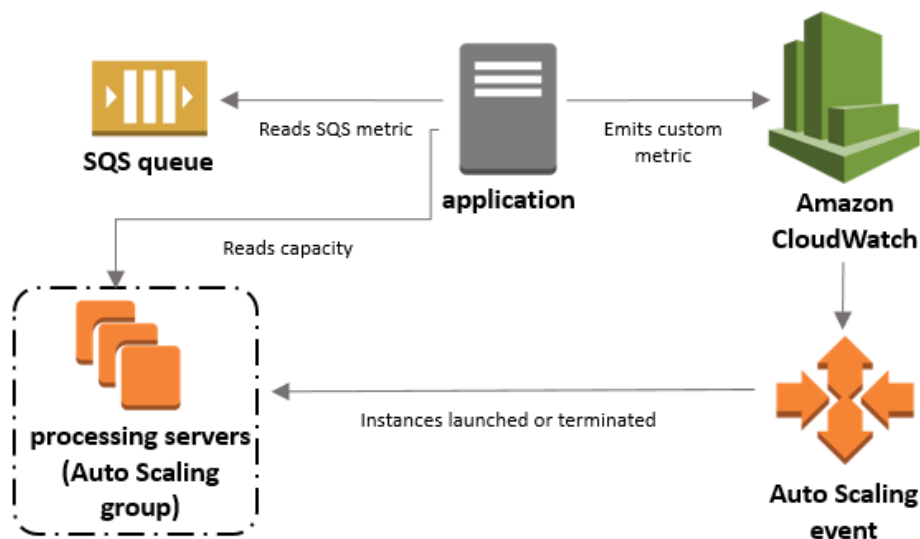
There are several scenarios where you might think about scaling in response to activity in an Amazon SQS queue. Suppose that you have a web app that lets users upload images and use them online. Each image requires resizing and encoding before it can be published. The app runs on EC2 instances in an Auto Scaling group that is configured to handle your typical upload rates. Unhealthy instances are terminated and replaced to maintain current instance levels at all times. The app places the raw bitmap data of the images in an Amazon SQS queue for processing. It processes the images and then publishes the processed images where they can be viewed by users.

This architecture works well if the number of image uploads doesn't vary over time. What happens if your upload levels change? If your uploads increase and decrease on a predictable schedule, you can specify the time and date to perform scaling activities. For more information, see [Scheduled Scaling for Amazon EC2 Auto Scaling](#) (p. 85). A more dynamic way to scale your Auto Scaling group, scaling by policy, lets you define parameters that control the scaling process. For example, you can create a policy that calls for enlarging your fleet of EC2 instances whenever the average number of uploads reaches a certain level. This is useful for scaling in response to changing conditions, when you don't know when those conditions will change.

There are three main parts to this configuration:

- An Auto Scaling group to manage EC2 instances for the purposes of processing messages from an SQS queue.
- A custom metric to send to Amazon CloudWatch that measures the number of messages in the queue per EC2 instance in the Auto Scaling group.
- A target tracking policy that configures your Auto Scaling group to scale based on the custom metric and a set target value. CloudWatch alarms invoke the scaling policy.

The following diagram illustrates the architecture of this configuration.



Choosing an Effective Metric and Target Value

The number of messages in your Amazon SQS queue does not solely define the number of instances needed. In fact, the number of instances in the fleet can be driven by multiple factors, including how long it takes to process a message and the acceptable amount of latency (queue delay).

The solution is to use a *backlog per instance* metric with the target value being the *acceptable backlog per instance* to maintain. You can calculate these numbers as follows:

- **Backlog per instance:** To determine your backlog per instance, start with the Amazon SQS metric `ApproximateNumberOfMessages` to determine the length of the SQS queue (number of messages available for retrieval from the queue). Divide that number by the fleet's running capacity, which for an Auto Scaling group is the number of instances in the `InService` state, to get the backlog per instance.
- **Acceptable backlog per instance:** To determine your target value, first calculate what your application can accept in terms of latency. Then, take the acceptable latency value and divide it by the average time that an EC2 instance takes to process a message.

To illustrate with an example, the current `ApproximateNumberOfMessages` is 1500 and the fleet's running capacity is 10. If the average processing time is 0.1 seconds for each message and the longest acceptable latency is 10 seconds, then the acceptable backlog per instance is $10 / 0.1$, which equals 100. This means that 100 is the target value for your target tracking policy. Because the backlog per instance is currently at 150 ($1500 / 10$), your fleet scales out by five instances to maintain proportion to the target value.

The following examples create the custom metric and target tracking scaling policy that configures your Auto Scaling group to scale based on these calculations.

Configure Scaling with Amazon SQS

The following section shows you how to set up automatic scaling for an SQS queue using the AWS CLI. The procedures assume that you already have a queue (standard or FIFO), an Auto Scaling group, and EC2 instances running the application that uses the queue.

Tasks

- [Step 1: Create a CloudWatch Custom Metric](#) (p. 104)
- [Step 2: Create a Target Tracking Scaling Policy](#) (p. 105)
- [Step 3: Test Your Scaling Policy](#) (p. 105)

Step 1: Create a CloudWatch Custom Metric

Create the custom calculated metric by first reading metrics from your AWS account. Then, calculate the backlog per instance metric recommended in an earlier section. Lastly, publish this number to CloudWatch at a 1-minute granularity.

Wherever possible, you should scale on EC2 instance metrics with a 1-minute frequency (also known as detailed monitoring) because that ensures a faster response to utilization changes. Scaling on metrics with a 5-minute frequency can result in slower response times and scaling on stale metric data. By default, EC2 instances are enabled for basic monitoring, which means metric data for instances is available at 5-minute intervals. You can enable detailed monitoring to get metric data for instances at a 1-minute frequency. For more information, see [Monitoring Your Auto Scaling Groups and Instances Using Amazon CloudWatch](#) (p. 135).

To create a CloudWatch custom metric

1. Use the SQS [get-queue-attributes](#) command to get the number of messages waiting in the queue (ApproximateNumberOfMessages):

```
aws sqs get-queue-attributes --queue-url https://sqs.region.amazonaws.com/123456789/MyQueue \
--attribute-names ApproximateNumberOfMessages
```

2. Use the [describe-auto-scaling-groups](#) command to get the running capacity of the group, which is the number of instances in the InService lifecycle state. This command returns the instances of an Auto Scaling group along with their lifecycle state.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-names my-asg
```

3. Calculate the backlog per instance by dividing the ApproximateNumberOfMessages metric by the fleet's running capacity metric.
4. Publish the results at a 1-minute granularity as a CloudWatch custom metric using the AWS CLI or an API. A custom metric is defined using a metric name and namespace of your choosing. Namespaces for custom metrics cannot start with "AWS/". For more information about publishing custom metrics, see the [Publish Custom Metrics](#) topic in the *Amazon CloudWatch User Guide*.

Here is an example CLI [put-metric-data](#) command.

```
aws cloudwatch put-metric-data --metric-name MyBacklogPerInstance --
namespace MyNamespace \
--unit None --value 20 --
dimensions MyOptionalMetricDimensionName=MyOptionalMetricDimensionValue
```

After your application is emitting the desired metric, the data is sent to CloudWatch. The metric is visible in the CloudWatch console. You can access it by logging into the AWS Management Console and navigating to the CloudWatch page. Then, view the metric by navigating to the metrics page or searching for it using the search box. For help viewing metrics, see [View Available Metrics](#) in the *Amazon CloudWatch User Guide*.

Step 2: Create a Target Tracking Scaling Policy

Next, create a target tracking scaling policy that tells the Auto Scaling group to increase and decrease the number of running EC2 instances in the group dynamically when the load on the application changes. You can use a target tracking scaling policy because the scaling metric is a utilization metric that increases and decreases proportionally to the capacity of the group.

To create a target tracking scaling policy

1. Use the following command to specify a target value for your scaling policy in a `config.json` file in your home directory.

For the `TargetValue`, calculate the acceptable backlog per instance metric and enter it here. To calculate this number, decide on a normal latency value and divide it by the average time that it takes to process a message.

```
$ cat ~/config.json
{
  "TargetValue": 100,
  "CustomizedMetricSpecification": {
    "MetricName": "MyBacklogPerInstance",
    "Namespace": "MyNamespace",
    "Dimensions": [
      {
        "Name": "MyOptionalMetricDimensionName",
        "Value": "MyOptionalMetricDimensionValue"
      }
    ],
    "Statistic": "Average",
    "Unit": "None"
  }
}
```

2. Use the `put-scaling-policy` command, along with the `config.json` file that you created in the previous step, to create your scaling policy:

```
aws autoscaling put-scaling-policy --policy-name my-scaling-policy --auto-scaling-
group-name my-asg \
  --policy-type TargetTrackingScaling --target-tracking-configuration file://~/
config.json
```

This creates two alarms: one for scaling out and one for scaling in. It also returns the Amazon Resource Name (ARN) of the policy that is registered with CloudWatch, which CloudWatch uses to invoke scaling whenever the metric is in breach.

Step 3: Test Your Scaling Policy

After your setup is complete, verify your scaling policy is working. You can test it by increasing the number of messages in your SQS queue and then verifying that your Auto Scaling group has launched an additional EC2 instance, and also by decreasing the number of messages in your SQS queue and then verifying that the Auto Scaling group has terminated an EC2 instance.

To test the scale-out function

1. Follow the steps in [Tutorial: Sending a Message to an Amazon SQS Queue](#) to add messages to your queue. Make sure that you have increased the number of messages in the queue so that the backlog per instance metric exceeds the target value.

It can take a few minutes for your changes to trigger the CloudWatch alarm.

2. Use the `describe-auto-scaling-groups` command to verify that the group has launched an instance:

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

To test the scale-in function

1. Follow the steps in [Tutorial: Sending a Message to an Amazon SQS Queue](#) to remove messages from the queue. Make sure that you have decreased the number of messages in the queue so that the backlog per instance metric is below the target value.

It can take a few minutes for your changes to trigger the CloudWatch alarm.

2. Use the `describe-auto-scaling-groups` command to verify that the group has terminated an instance:

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

Deleting a Scaling Policy

After you no longer need a scaling policy, you can delete it. You might also need to delete the CloudWatch alarms. Deleting a target tracking scaling policy also deletes any associated CloudWatch alarms. Deleting a step scaling policy or a simple scaling policy deletes the underlying alarm action, but does not delete the CloudWatch alarm associated with the scaling policy, even if it no longer has an associated action.

To delete a scaling policy (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. On the **Scaling Policies** tab, choose **Actions, Delete**.
5. When prompted for confirmation, choose **Yes, Delete**.
6. (Optional) If you deleted a step scaling policy or a simple scaling policy, do the following to delete the CloudWatch alarm that was associated with the policy. You can skip this step to keep the alarm for future use.
 - a. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
 - b. On the navigation pane, choose **Alarms**.
 - c. Choose the alarm (for example, `awsec2-my-auto-scaling-group-CPU-Utilization`) and choose **Action, Delete**.
 - d. When prompted for confirmation, choose **Delete**.

To delete your scaling policy (AWS CLI)

Use the following `delete-policy` command.

```
aws autoscaling delete-policy --auto-scaling-group-name my-asg \  
--policy-name my-step-scalein-policy
```

To delete your CloudWatch alarm (AWS CLI)

For step scaling policies and simple scaling policies, use the [delete-alarms](#) command to delete the CloudWatch alarm that was associated with the policy. You can skip this step to keep the alarm for future use. You can delete one or more alarms at a time. For example, use the following command to delete the Step-Scaling-AlarmHigh-AddCapacity and Step-Scaling-AlarmLow-RemoveCapacity alarms.

```
aws cloudwatch delete-alarms --alarm-name Step-Scaling-AlarmHigh-AddCapacity Step-Scaling-AlarmLow-RemoveCapacity
```

Scaling Cooldowns for Amazon EC2 Auto Scaling

The cooldown period helps to ensure that your Auto Scaling group doesn't launch or terminate additional instances before the previous scaling activity takes effect. You can configure the length of time based on your instance warmup period or other application needs.

Amazon EC2 Auto Scaling supports cooldown periods when using simple scaling policies, but not when using other scaling policies. After the Auto Scaling group dynamically scales using a simple scaling policy, it waits for the cooldown period to complete before resuming scaling activities. You can use the default cooldown period associated with your Auto Scaling group, or you can override the default by specifying a cooldown period for your policy. For more information about simple scaling, see [Simple and Step Scaling Policies for Amazon EC2 Auto Scaling](#) (p. 93).

When you manually scale your Auto Scaling group, the default is not to wait for the cooldown period, but you can override the default and honor the cooldown period. If an instance becomes unhealthy, the Auto Scaling group does not wait for the cooldown period to complete before replacing the unhealthy instance. For more information about manual scaling, see [Manual Scaling for Amazon EC2 Auto Scaling](#) (p. 75).


Amazon EC2 Auto Scaling supports both default cooldown periods and scaling-specific cooldown periods.

Contents

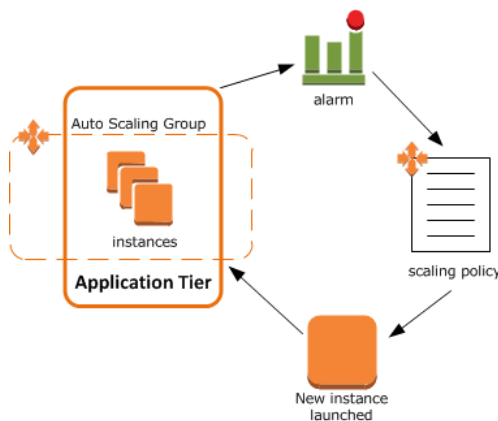
- [Example: Cooldowns](#) (p. 107)
- [Default Cooldowns](#) (p. 108)
- [Scaling-Specific Cooldowns](#) (p. 109)
- [Cooldowns and Multiple Instances](#) (p. 109)
- [Cooldowns and Lifecycle Hooks](#) (p. 109)

Example: Cooldowns

Consider the following scenario: you have a web application running in AWS. This web application consists of three basic tiers: web, application, and database. To make sure that the application always has the resources to meet traffic demands, you create two Auto Scaling groups: one for your web tier and one for your application tier.



To help ensure that the Auto Scaling group for the application tier has the appropriate number of EC2 instances, [create a CloudWatch alarm \(p. 139\)](#) to scale out whenever the **CPUUtilization** metric for the instances exceeds 90 percent. When the alarm occurs, the Auto Scaling group launches and configures another instance.




These instances use a configuration script to install and configure software before the instance is put into service. As a result, it takes around two or three minutes from the time the instance launches until it's fully in service. The actual time depends on several factors, such as the size of the instance and whether there are startup scripts to complete.

Now a spike in traffic occurs, causing the CloudWatch alarm to fire. When it does, the Auto Scaling group launches an instance to help with the increase in demand. However, there's a problem: the instance takes a couple of minutes to launch. During that time, the CloudWatch alarm could continue to fire, causing the Auto Scaling group to launch another instance each time the alarm fires.

However, with a cooldown period in place, the Auto Scaling group launches an instance and then suspends scaling activities due to simple scaling policies or manual scaling until the specified time elapses. (The default is 300 seconds.) This gives newly launched instances time to start handling application traffic. After the cooldown period expires, any suspended scaling actions resume. If the CloudWatch alarm fires again, the Auto Scaling group launches another instance, and the cooldown period takes effect again. If, however, the additional instance was enough to bring the CPU utilization back down, then the group remains at its current size.

Default Cooldowns

The default cooldown period is applied when you create your Auto Scaling group. Its default value is 300 seconds. This cooldown period automatically applies to any scaling activities for simple scaling policies, and you can optionally request to have it apply to your manual scaling activities.



You can configure the default cooldown period when you create the Auto Scaling group, using the AWS Management Console, the [create-auto-scaling-group](#) command (AWS CLI), or the [CreateAutoScalingGroup](#) API operation.

You can change the default cooldown period at any time, using the AWS Management Console, the [update-auto-scaling-group](#) command (AWS CLI), or the [UpdateAutoScalingGroup](#) API operation.

Scaling-Specific Cooldowns

In addition to specifying the default cooldown period for your Auto Scaling group, you can create cooldowns that apply to a specific simple scaling policy or manual scaling. A scaling-specific cooldown period overrides the default cooldown period.

One common use for scaling-specific cooldowns is with a scale-in policy—a policy that terminates instances based on a specific criteria or metric. Because this policy terminates instances, Amazon EC2 Auto Scaling needs less time to determine whether to terminate additional instances. The default cooldown period of 300 seconds is too long, in which case a scaling-specific cooldown period of 180 seconds for your scale-in policy can reduce costs by allowing the group to scale in faster.

You can create a scaling-specific cooldown period using the AWS Management Console, the [put-scaling-policy](#) command (AWS CLI), or the [PutScalingPolicy](#) API operation.

Cooldowns and Multiple Instances

The preceding sections have provided examples that show how cooldown periods affect Auto Scaling groups when a single instance launches or terminates. However, it is common for Auto Scaling groups to launch more than one instance at a time. For example, you might choose to have the Auto Scaling group launch three instances when a specific metric threshold is met.

With multiple instances, the cooldown period (either the default cooldown or the scaling-specific cooldown) takes effect starting when the last instance launches.

Cooldowns and Lifecycle Hooks

You can add lifecycle hooks to your Auto Scaling groups. These hooks enable you to control how instances launch and terminate within an Auto Scaling group. You can perform actions on the instance before it is put into service or before it is terminated.

Lifecycle hooks can affect the impact of any cooldown periods configured for the Auto Scaling group, manual scaling, or a simple scaling policy. The cooldown period does not begin until after the instance moves out of the wait state.

Controlling Which Auto Scaling Instances Terminate During Scale In

With each Auto Scaling group, you control when it adds instances (referred to as *scaling out*) or removes instances (referred to as *scaling in*) from your network architecture. You can scale the size of your group manually by attaching and detaching instances, or you can automate the process through the use of a scaling policy.

To specify which instances to terminate first during scale in, configure a termination policy for the Auto Scaling group.

Optionally, you can use instance protection to prevent specific instances from being terminated during automatic scale in. For instances in an Auto Scaling group, use Amazon EC2 Auto Scaling features to protect an instance when a scale-in event occurs. If you want to protect your instance from being accidentally terminated, use Amazon EC2 termination protection.

Note

For [Auto Scaling Groups with Multiple Instance Types and Purchase Options \(p. 43\)](#), the scale-in action respects the On-Demand base amount and the percentages set for On-Demand and Spot Instances. When the group scales in, the Auto Scaling group first identifies which of these two types of instances it will terminate. Then, it chooses the Availability Zone with the most instances of that type, and applies the default or customized termination policy.

Contents

- [Default Termination Policy \(p. 110\)](#)
- [Customizing the Termination Policy \(p. 112\)](#)
- [Instance Protection \(p. 113\)](#)

Default Termination Policy

This section describes the default termination policy used by an Auto Scaling group when a scale-in event occurs. The default termination policy is designed to help ensure that your instances span Availability Zones evenly for high availability. The default policy is kept generic and flexible to cover a range of scenarios.

With the default termination policy, the behavior of the Auto Scaling group is as follows:

1. Determine which Availability Zone(s) have the most instances, and at least one instance that is not protected from scale in.

If there are multiple unprotected instances to choose from in the Availability Zone(s) with the most instances, an instance is selected for termination based on the following criteria (applied in the order shown).

2. [For [Auto Scaling Groups with Multiple Instance Types and Purchase Options \(p. 43\)](#) only]

Determine which instance to terminate so as to align the remaining instances to the allocation strategy for the On-Demand or Spot Instance that is terminating and your current selection of instance types. If there is one such instance, terminate it. Otherwise, apply the next condition.

3. [For Auto Scaling groups that use a launch template]

Determine whether any of the instances use the oldest launch template. If there is one such instance, terminate it. (Note that there is one exception: if the group originally used a launch configuration. Amazon EC2 Auto Scaling terminates instances that use a launch configuration before instances with the oldest launch template.)

4. [For Auto Scaling groups that use a launch configuration]

Determine whether any of the instances use the oldest launch configuration. If there is one such instance, terminate it.

5. After applying all of the criteria in 2 through 4, if there are multiple unprotected instances to terminate, determine which instances are closest to the next billing hour. If there is one such instance, terminate it. (Terminating the instance closest to the next billing hour helps you maximize the use of your instances that have an hourly charge.)


Note

With On-Demand Instances, you pay for compute capacity by per hour or per second depending on which instances you run. If your Auto Scaling group uses Amazon Linux or Ubuntu, your EC2 usage is billed in one-second increments. For more information, see [Amazon EC2 Pricing](#).


6. If there are multiple unprotected instances closest to the next billing hour, choose one of these instances at random.

Example

Consider an Auto Scaling group that uses a launch configuration. It has one instance type, two Availability Zones, a desired capacity of two instances, and scaling policies that increase and decrease the number of instances by one when certain thresholds are met. The two instances in this group are distributed as follows.



When the threshold for the scale-out policy is met, the policy takes effect and the Auto Scaling group launches a new instance. The Auto Scaling group now has three instances, distributed as follows.



When the threshold for the scale-in policy is met, the policy takes effect and the Auto Scaling group terminates one of the instances. If you did not assign a specific termination policy to the group, it uses the default termination policy. It selects the Availability Zone with two instances, and terminates the instance launched from the oldest launch configuration. If the instances were launched from the same launch configuration, the Auto Scaling group selects the instance that is closest to the next billing hour and terminates it.

For more information about high availability with Amazon EC2 Auto Scaling, see [Distributing Instances Across Availability Zones \(p. 6\)](#).

Customizing the Termination Policy

The default termination policy assigned to an Auto Scaling group is typically sufficient for most situations. However, you have the option of replacing the default policy with a customized one.

When you customize the termination policy, if one Availability Zone has more instances than the other Availability Zones that are used by the group, your termination policy is applied to the instances from the imbalanced Availability Zone. If the Availability Zones used by the group are balanced, the termination policy is applied across all of the Availability Zones for the group.

Amazon EC2 Auto Scaling supports the following custom termination policies:

- **OldestInstance**. Terminate the oldest instance in the group. This option is useful when you're upgrading the instances in the Auto Scaling group to a new EC2 instance type. You can gradually replace instances of the old type with instances of the new type.
- **NewestInstance**. Terminate the newest instance in the group. This policy is useful when you're testing a new launch configuration but don't want to keep it in production.
- **OldestLaunchConfiguration**. Terminate instances that have the oldest launch configuration. This policy is useful when you're updating a group and phasing out the instances from a previous configuration.
- **ClosestToNextInstanceHour**. Terminate instances that are closest to the next billing hour. This policy helps you maximize the use of your instances that have an hourly charge.
- **Default**. Terminate instances according to the default termination policy. This policy is useful when you have more than one scaling policy for the group.
- **OldestLaunchTemplate**. Terminate instances that have the oldest launch template. With this policy, instances that use the non-current launch template are terminated first, followed by instances that use the oldest version of the current launch template. This policy is useful when you're updating a group and phasing out the instances from a previous configuration.
- **AllocationStrategy**. Terminate instances in the Auto Scaling group to align the remaining instances to the allocation strategy for the type of instance that is terminating (either a Spot Instance or an On-Demand Instance). This policy is useful when your preferred instance types have changed. If the Spot allocation strategy is `lowest-price`, you can gradually rebalance the distribution of Spot Instances across your N lowest priced Spot pools. If the Spot allocation strategy is `capacity-optimized`, you can gradually rebalance the distribution of Spot Instances across Spot pools where there is more available Spot capacity. You can also gradually replace On-Demand Instances of a lower priority type with On-Demand Instances of a higher priority type.

To customize a termination policy (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. For **Actions**, choose **Edit**.
5. On the **Details** tab, locate **Termination Policies**. Choose one or more termination policies. If you choose multiple policies, list them in the order in which they should apply. If you use the **Default** policy, make it the last one in the list.
6. Choose **Save**.

To customize a termination policy (AWS CLI)

Use one of the following commands:

- [create-auto-scaling-group](#)
- [update-auto-scaling-group](#)

You can use these policies individually, or combine them into a list of policies. For example, use the following command to update an Auto Scaling group to use the `OldestLaunchConfiguration` policy first and then use the `ClosestToNextInstanceHour` policy:

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg --termination-policies "OldestLaunchConfiguration" "ClosestToNextInstanceHour"
```

If you use the `Default` termination policy, make it the last one in the list of termination policies. For example, `--termination-policies "OldestLaunchConfiguration" "Default"`.

Instance Protection

To control whether an Auto Scaling group can terminate a particular instance when scaling in, use instance protection. You can enable the instance protection setting on an Auto Scaling group or on an individual Auto Scaling instance. When the Auto Scaling group launches an instance, it inherits the instance protection setting of the Auto Scaling group. You can change the instance protection setting for an Auto Scaling group or an Auto Scaling instance at any time.

Instance protection starts when the instance state is `InService`. If you detach an instance that is protected from termination, its instance protection setting is lost. When you attach the instance to the group again, it inherits the current instance protection setting of the group.

If all instances in an Auto Scaling group are protected from termination during scale in, and a scale-in event occurs, its desired capacity is decremented. However, the Auto Scaling group can't terminate the required number of instances until their instance protection settings are disabled.

Instance protection does not protect Auto Scaling instances from the following:

- Manual termination through the Amazon EC2 console, the `terminate-instances` command, or the `TerminateInstances` action. To protect Auto Scaling instances from manual termination, enable Amazon EC2 termination protection. For more information, see [Enabling Termination Protection](#) in the *Amazon EC2 User Guide for Linux Instances*.
- Health check replacement if the instance fails health checks. For more information, see [Health Checks for Auto Scaling Instances \(p. 132\)](#). To prevent Amazon EC2 Auto Scaling from terminating unhealthy instances, suspend the `ReplaceUnhealthy` process. For more information, see [Suspending and Resuming Scaling Processes \(p. 127\)](#).
- Spot Instance interruptions. A Spot instance is terminated when capacity is no longer available or the Spot price exceeds your maximum price.

Tasks

- [Enable Instance Protection for a Group \(p. 113\)](#)
- [Modify the Instance Protection Setting for a Group \(p. 114\)](#)
- [Modify the Instance Protection Setting for an Instance \(p. 115\)](#)

Enable Instance Protection for a Group

You can enable instance protection when you create an Auto Scaling group. By default, instance protection is disabled.

To enable instance protection (console)

When you create the Auto Scaling group, on the **Configure Auto Scaling group details** page, under **Advanced Details**, select the **Protect From Scale In** option from **Instance Protection**.

▼ **Advanced Details**

Load Balancing ⓘ	You currently don't have any load balancers Learn about Elastic Load Balancing
Health Check Grace Period ⓘ	<input type="text" value="300"/> seconds
Monitoring ⓘ	Amazon EC2 Detailed Monitoring metrics, which are provided at 1 minute frequency, are not enabled for the launch configuration my-lc. Instances launched from it will use Basic Monitoring metrics, provided at 5 minute frequency. Learn more
Instance Protection ⓘ	<div><input type="text" value="Protect From Scale In"/></div>

To enable instance protection (AWS CLI)

Use the following [create-auto-scaling-group](#) command to enable instance protection:

```
aws autoscaling create-auto-scaling-group --auto-scaling-group-name my-asg --new-instances-protected-from-scale-in ...
```

Modify the Instance Protection Setting for a Group

You can enable or disable the instance protection setting for an Auto Scaling group.

To change the instance protection setting for a group (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. On the **Details** tab, choose **Edit**.
5. For **Instance Protection**, select **Protect From Scale In**.

Instance Protection ⓘ

6. Choose **Save**.

To change the instance protection setting for a group (AWS CLI)

Use the following [update-auto-scaling-group](#) command to enable instance protection for the specified Auto Scaling group:

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg --new-instances-protected-from-scale-in
```

Use the following command to disable instance protection for the specified group:

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-asg --no-new-  
instances-protected-from-scale-in
```

Modify the Instance Protection Setting for an Instance

By default, an instance gets its instance protection setting from its Auto Scaling group. However, you can enable or disable instance protection for an instance at any time.

To change the instance protection setting for an instance (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. On the **Instances** tab, select the instance.
5. To enable instance protection, choose **Actions, Instance Protection, Set Scale In Protection**. When prompted, choose **Set Scale In Protection**.
6. To disable instance protection, choose **Actions, Instance Protection, Remove Scale In Protection**. When prompted, choose **Remove Scale In Protection**.

To change the instance protection setting for an instance (AWS CLI)

Use the following [set-instance-protection](#) command to enable instance protection for the specified instance:

```
aws autoscaling set-instance-protection --instance-ids i-5f2e8a0d --auto-scaling-group-  
name my-asg --protected-from-scale-in
```

Use the following command to disable instance protection for the specified instance:

```
aws autoscaling set-instance-protection --instance-ids i-5f2e8a0d --auto-scaling-group-  
name my-asg --no-protected-from-scale-in
```

Amazon EC2 Auto Scaling Lifecycle Hooks

Lifecycle hooks enable you to perform custom actions by *pausing* instances as an Auto Scaling group launches or terminates them. When an instance is paused, it remains in a wait state until either you complete the lifecycle action using the `complete-lifecycle-action` CLI command or `CompleteLifecycleAction` API action, or the timeout period ends (one hour by default).

For example, your newly launched instance completes its startup sequence and a lifecycle hook pauses the instance. While the instance is in a wait state, you can install or configure software on it, making sure that your instance is fully ready before it starts receiving traffic. For another example of the use of lifecycle hooks, when a scale-in event occurs, the terminating instance is first deregistered from the load balancer (if the Auto Scaling group is being used with Elastic Load Balancing). Then, a lifecycle hook pauses the instance before it is terminated. While the instance is in the wait state, you can, for example, connect to the instance and download logs or other data before the instance is fully terminated.

Each Auto Scaling group can have multiple lifecycle hooks. However, there is a limit on the number of hooks per Auto Scaling group. For more information, see [Amazon EC2 Auto Scaling Limits \(p. 9\)](#).

Contents

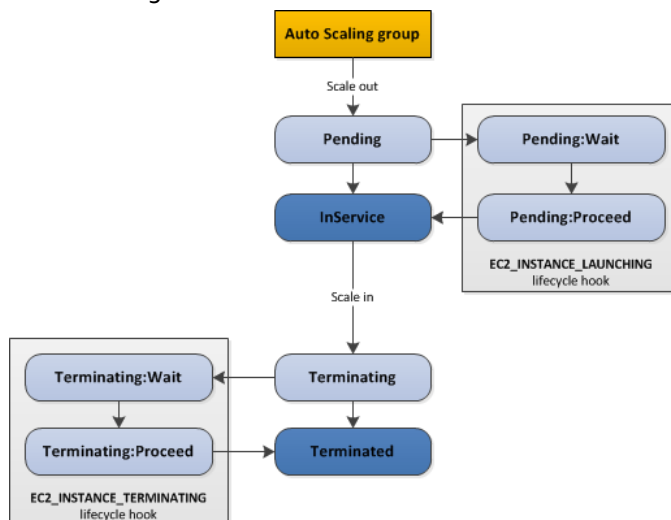
- [How Lifecycle Hooks Work \(p. 116\)](#)
- [Considerations When Using Lifecycle Hooks \(p. 117\)](#)
- [Prepare for Notifications \(p. 118\)](#)
- [Add Lifecycle Hooks \(p. 121\)](#)
- [Complete the Lifecycle Hook \(p. 122\)](#)
- [Test the Notification \(p. 122\)](#)

How Lifecycle Hooks Work

After you add lifecycle hooks to your Auto Scaling group, they work as follows:

1. The Auto Scaling group responds to scale-out events by launching instances and scale-in events by terminating instances.
2. The lifecycle hook puts the instance into a wait state (`Pending:Wait` or `Terminating:Wait`). The instance is paused until you continue or the timeout period ends.
3. You can perform a custom action using one or more of the following options:
 - Define a CloudWatch Events target to invoke a Lambda function when a lifecycle action occurs. The Lambda function is invoked when Amazon EC2 Auto Scaling submits an event for a lifecycle action to CloudWatch Events. The event contains information about the instance that is launching or terminating, and a token that you can use to control the lifecycle action.
 - Define a notification target for the lifecycle hook. Amazon EC2 Auto Scaling sends a message to the notification target. The message contains information about the instance that is launching or terminating, and a token that you can use to control the lifecycle action.
 - Create a script that runs on the instance as the instance starts. The script can control the lifecycle action using the ID of the instance on which it runs.
4. By default, the instance remains in a wait state for one hour, and then the Auto Scaling group continues the launch or terminate process (`Pending:Proceed` or `Terminating:Proceed`). If you need more time, you can restart the timeout period by recording a heartbeat. If you finish before the timeout period ends, you can complete the lifecycle action, which continues the launch or termination process.

The following illustration shows the transitions between instance states in this process:



For more information about the complete lifecycle of instances in an Auto Scaling group, see [Auto Scaling Lifecycle](#) (p. 7).

Considerations When Using Lifecycle Hooks

Adding lifecycle hooks to your Auto Scaling group gives you greater control over how instances launch and terminate. The following are things to consider when adding a lifecycle hook to your Auto Scaling group, to help ensure that the group continues to perform as expected.

Considerations

- [Keeping Instances in a Wait State](#) (p. 117)
- [Cooldowns and Custom Actions](#) (p. 117)
- [Health Check Grace Period](#) (p. 117)
- [Lifecycle Action Result](#) (p. 118)
- [Spot Instances](#) (p. 118)

Keeping Instances in a Wait State

Instances can remain in a wait state for a finite period of time. The default is one hour (3600 seconds). You can adjust this time in the following ways:

- Set the heartbeat timeout for the lifecycle hook when you create the lifecycle hook. With the **put-lifecycle-hook** command, use the `--heartbeat-timeout` parameter. With the `PutLifecycleHook` operation, use the `HeartbeatTimeout` parameter.
- Continue to the next state if you finish before the timeout period ends, using the **complete-lifecycle-action** command or the `CompleteLifecycleAction` operation.
- Postpone the end of the timeout period by recording a heartbeat, using the **record-lifecycle-action-heartbeat** command or the `RecordLifecycleActionHeartbeat` operation. This extends the timeout period by the timeout value specified when you created the lifecycle hook. For example, if the timeout value is one hour, and you call this command after 30 minutes, the instance remains in a wait state for an additional hour, or a total of 90 minutes.

The maximum amount of time that you can keep an instance in a wait state is 48 hours or 100 times the heartbeat timeout, whichever is smaller.

Cooldowns and Custom Actions

When an Auto Scaling group launches or terminates an instance due to a simple scaling policy, a **cooldown** (p. 107) takes effect. The cooldown period helps ensure that the Auto Scaling group does not launch or terminate more instances than needed.

Consider an Auto Scaling group with a lifecycle hook that supports a custom action at instance launch. When the application experiences an increase in demand, the group launches instances to add capacity. Because there is a lifecycle hook, the instance is put into the `Pending:Wait` state, which means that it is not available to handle traffic yet. When the instance enters the wait state, scaling actions due to simple scaling policies are suspended. When the instance enters the `InService` state, the cooldown period starts. When the cooldown period expires, any suspended scaling actions resume.

Health Check Grace Period

If you add a lifecycle hook, the health check grace period does not start until the lifecycle hook actions complete and the instance enters the `InService` state.

Lifecycle Action Result

At the conclusion of a lifecycle hook, the result is either `ABANDON` or `CONTINUE`.

If the instance is launching, `CONTINUE` indicates that your actions were successful, and that the instance can be put into service. Otherwise, `ABANDON` indicates that your custom actions were unsuccessful, and that the instance can be terminated.

If the instance is terminating, both `ABANDON` and `CONTINUE` allow the instance to terminate. However, `ABANDON` stops any remaining actions, such as other lifecycle hooks, while `CONTINUE` allows any other lifecycle hooks to complete.

Spot Instances

You can use lifecycle hooks with Spot Instances. However, a lifecycle hook does not prevent an instance from terminating in the event that capacity is no longer available. In addition, when a Spot Instance terminates, you must still complete the lifecycle action (using the **complete-lifecycle-action** command or the `CompleteLifecycleAction` operation).

Prepare for Notifications

You can configure notifications when the instance enters a wait state, which enables you to perform a custom action. You can use Amazon CloudWatch Events, Amazon SNS, or Amazon SQS to receive the notifications. Choose whichever option you prefer.

Important

AWS resources for notifications must always be created in the same AWS Region where you create your lifecycle hook. For example, if you configure notifications using Amazon SNS, the Amazon SNS topic must reside in the same region as your lifecycle hook.

Alternatively, if you have a script that configures your instances when they launch, you do not need to receive notification when the lifecycle action occurs. If you are not doing so already, update your script to retrieve the instance ID of the instance from the instance metadata. For more information, see [Retrieving Instance Metadata](#).

Notification Options

- [Route Notifications to Lambda Using CloudWatch Events \(p. 118\)](#)
- [Receive Notification Using Amazon SNS \(p. 119\)](#)
- [Receive Notification Using Amazon SQS \(p. 120\)](#)

Route Notifications to Lambda Using CloudWatch Events

You can use CloudWatch Events to set up a target to invoke a Lambda function when a lifecycle action occurs.

To set up notifications using CloudWatch Events

1. Create a Lambda function using the steps in [Create a Lambda Function \(p. 144\)](#) and note its Amazon Resource Name (ARN). For example, `arn:aws:lambda:region:123456789012:function:my-function`.
2. Create a CloudWatch Events rule that matches the lifecycle action using the following **put-rule** command.

```
aws events put-rule --name my-rule --event-pattern file://pattern.json --state ENABLED
```

The `pattern.json` for an instance launch lifecycle action is:

```
{
  "source": [ "aws.autoscaling" ],
  "detail-type": [ "EC2 Instance-launch Lifecycle Action" ]
}
```

The `pattern.json` for an instance terminate lifecycle action is:

```
{
  "source": [ "aws.autoscaling" ],
  "detail-type": [ "EC2 Instance-terminate Lifecycle Action" ]
}
```

3. Grant the rule permission to invoke your Lambda function using the following **add-permission** command. This command trusts the CloudWatch Events service principal (`events.amazonaws.com`) and scopes permissions to the specified rule.

```
aws lambda add-permission --function-name LogScheduledEvent --statement-id my-scheduled-event \
  --action 'lambda:InvokeFunction' --principal events.amazonaws.com --source-arn
arn:aws:events:region:123456789012:rule/my-scheduled-rule
```

4. Create a target that invokes your Lambda function when the lifecycle action occurs, using the following **put-targets** command.

```
aws events put-targets --rule my-rule --targets
Id=1,Arn=arn:aws:lambda:region:123456789012:function:my-function
```

5. After you have followed these instructions, continue on to [Add Lifecycle Hooks \(p. 121\)](#) as a next step.

When the Auto Scaling group responds to a scale-out or scale-in event, it puts the instance in a wait state. While the instance is in a wait state, the Lambda function is invoked. For more information about the event data, see [Auto Scaling Events \(p. 140\)](#).

Receive Notification Using Amazon SNS

You can use Amazon SNS to set up a notification target to receive notifications when a lifecycle action occurs.

To set up notifications using Amazon SNS

1. Create an Amazon SNS topic using the following **create-topic** command. For more information, see [Create a Topic](#) in the *Amazon Simple Notification Service Developer Guide*.

```
aws sns create-topic --name my-sns-topic
```

Note the ARN of the target (for example, `arn:aws:sns:region:123456789012:my-sns-topic`).

2. Create a service role (or *assume* role) for Amazon EC2 Auto Scaling to which you can grant permission to access your notification target.

To create an IAM role and allow Amazon EC2 Auto Scaling to assume it

- a. Open the IAM console at <https://console.aws.amazon.com/iam/>.

- b. In the navigation pane, choose **Roles**, **Create new role**.
 - c. Under **Select type of trusted entity**, choose **AWS service**.
 - d. Under **Choose the service that will use this role**, choose **EC2 Auto Scaling** from the list.
 - e. Under **Select your use case**, choose **EC2 Auto Scaling Notification Access**, and then choose **Next:Permissions**.
 - f. Choose **Next:Tags**, (optional) add metadata to the role by attaching tags as key-value pairs, and then choose **Next:Review**.
 - g. On the **Review** page, type a name for the role (e.g. my-notification-role) and choose **Create role**.
 - h. On the **Roles** page, choose the role you just created to open the **Summary** page. Make a note of the **Role ARN**. For example, `arn:aws:iam::123456789012:role/my-notification-role`. You will specify the role ARN when you create the lifecycle hook in the next procedure.
3. After you have followed these instructions, continue on to [Add Lifecycle Hooks \(p. 121\)](#) as a next step.

When the Auto Scaling group responds to a scale-out or scale-in event, it puts the instance in a wait state. While the instance is in a wait state, a message is published to the notification target. The message includes the following event data:

- `LifecycleActionToken` — The lifecycle action token.
- `AccountId` — The AWS account ID.
- `AutoScalingGroupName` — The name of the Auto Scaling group.
- `LifecycleHookName` — The name of the lifecycle hook.
- `EC2InstanceId` — The ID of the EC2 instance.
- `LifecycleTransition` — The lifecycle hook type.

For example:

```
Service: AWS Auto Scaling
Time: 2019-04-30T20:42:11.305Z
RequestId: 18b2ec17-3e9b-4c15-8024-ff2e8ce8786a
LifecycleActionToken: 71514b9d-6a40-4b26-8523-05e7ee35fa40
AccountId: 123456789012
AutoScalingGroupName: my-asg
LifecycleHookName: my-hook
EC2InstanceId: i-0598c7d356eba48d7
LifecycleTransition: autoscaling:EC2_INSTANCE_LAUNCHING
NotificationMetadata: null
```

Receive Notification Using Amazon SQS

You can use Amazon SQS to set up a notification target to receive notifications when a lifecycle action occurs.

Important

FIFO queues are not compatible with lifecycle hooks.

To set up notifications using Amazon SQS

1. Create the target using Amazon SQS. For more information, see [Getting Started with Amazon SQS](#) in the *Amazon Simple Queue Service Developer Guide*. Note the ARN of the target (for example, `arn:aws:sqs:region:123456789012:my-sqs-queue`).

2. Create a service role (or *assume* role) for Amazon EC2 Auto Scaling to which you can grant permission to access your notification target.

To create an IAM role and allow Amazon EC2 Auto Scaling to assume it

- a. Open the IAM console at <https://console.aws.amazon.com/iam/>.
 - b. In the navigation pane, choose **Roles, Create new role**.
 - c. Under **Select type of trusted entity**, choose **AWS service**.
 - d. Under **Choose the service that will use this role**, choose **EC2 Auto Scaling** from the list.
 - e. Under **Select your use case**, choose **EC2 Auto Scaling Notification Access**, and then choose **Next:Permissions**.
 - f. Choose **Next:Tags**, (optional) add metadata to the role by attaching tags as key–value pairs, and then choose **Next:Review**.
 - g. On the **Review** page, type a name for the role (e.g. my-notification-role) and choose **Create role**.
 - h. On the **Roles** page, choose the role you just created to open the **Summary** page. Make a note of the **Role ARN**. For example, `arn:aws:iam::123456789012:role/my-notification-role`. You will specify the role ARN when you create the lifecycle hook in the next procedure.
3. After you have followed these instructions, continue on to [Add Lifecycle Hooks \(p. 121\)](#) as a next step.

When the Auto Scaling group responds to a scale-out or scale-in event, it puts the instance in a wait state. While the instance is in a wait state, a message is published to the notification target.

Add Lifecycle Hooks

You can create and update lifecycle hooks using the [put-lifecycle-hook](#) command.

To perform an action on scale out, use the following command.

```
aws autoscaling put-lifecycle-hook --lifecycle-hook-name my-hook --auto-scaling-group-name my-asg \
  --lifecycle-transition autoscaling:EC2_INSTANCE_LAUNCHING
```

To perform an action on scale in, use the following command instead.

```
aws autoscaling put-lifecycle-hook --lifecycle-hook-name my-hook --auto-scaling-group-name my-asg \
  --lifecycle-transition autoscaling:EC2_INSTANCE_TERMINATING
```

To receive notifications using Amazon SNS or Amazon SQS, you must specify additional options. If you followed the previous procedure, your notification target and IAM role are already created. For more information, see [Prepare for Notifications \(p. 118\)](#).

For example, add the following options to specify an SNS topic as the notification target.

```
--notification-target-arn arn:aws:sns:region:123456789012:my-sns-topic --role-arn arn:aws:iam::123456789012:role/my-notification-role
```

The topic receives a test notification with the following key-value pair.

```
"Event": "autoscaling:TEST_NOTIFICATION"
```


Complete the Lifecycle Hook

When an Auto Scaling group responds to a scale-out or scale-in event, it puts the instance in a wait state and sends any notifications. It continues the launch or terminate processes after you complete the lifecycle hook.

To complete a lifecycle hook

1. While the instance is in a wait state, you can perform a custom action. For information about how to create a notification, see [Prepare for Notifications \(p. 118\)](#).
2. If you need more time to complete the custom action, use the [record-lifecycle-action-heartbeat](#) command to restart the timeout period and keep the instance in a wait state. You can specify the lifecycle action token you received in the previous step, as shown in the following command.

```
aws autoscaling record-lifecycle-action-heartbeat --lifecycle-hook-name my-launch-hook \
  --auto-scaling-group-name my-asg --lifecycle-action-
token bcd2f1b8-9a78-44d3-8a7a-4dd07d7cf635
```

Alternatively, you can specify the ID of the instance you retrieved in the previous step, as shown in the following command.

```
aws autoscaling record-lifecycle-action-heartbeat --lifecycle-hook-name my-launch-hook \
  --auto-scaling-group-name my-asg --instance-id i-1a2b3c4d
```

3. If you finish the custom action before the timeout period ends, use the [complete-lifecycle-action](#) command so that the Auto Scaling group can continue launching or terminating the instance. You can specify the lifecycle action token, as shown in the following command.

```
aws autoscaling complete-lifecycle-action --lifecycle-action-result CONTINUE \
  --lifecycle-hook-name my-launch-hook --auto-scaling-group-name my-asg \
  --lifecycle-action-token bcd2f1b8-9a78-44d3-8a7a-4dd07d7cf635
```

Alternatively, you can specify the ID of the instance, as shown in the following command.

```
aws autoscaling complete-lifecycle-action --lifecycle-action-result CONTINUE \
  --instance-id i-1a2b3c4d --lifecycle-hook-name my-launch-hook \
  --auto-scaling-group-name my-asg
```

Test the Notification

To generate a notification for a launch event, update the Auto Scaling group by increasing the desired capacity of the Auto Scaling group by 1. You receive a notification within a few minutes after instance launch.

To change the desired capacity (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Details** tab, choose **Edit**.
5. For **Desired**, increase the current value by 1. If this value exceeds **Max**, you must also increase the value of **Max** by 1.

6. Choose **Save**.
7. After a few minutes, you'll receive notification for the event. If you do not need the additional instance that you launched for this test, you can decrease **Desired** by 1. After a few minutes, you'll receive notification for the event.

Temporarily Removing Instances from Your Auto Scaling Group

You can put an instance that is in the `InService` state into the `Standby` state, update or troubleshoot the instance, and then return the instance to service. Instances that are on standby are still part of the Auto Scaling group, but they do not actively handle application traffic.

Important

You are billed for instances that are in a standby state.

For example, you can change the launch configuration for an Auto Scaling group at any time, and any subsequent instances that the Auto Scaling group launches use this configuration. However, the Auto Scaling group does not update the instances that are currently in service. You can terminate these instances and let the Auto Scaling group replace them. Or, you can put the instances on standby, update the software, and then put the instances back in service.

Contents

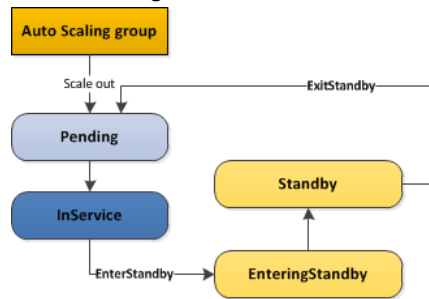
- [How the Standby State Works \(p. 123\)](#)
- [Health Status of an Instance in a Standby State \(p. 124\)](#)
- [Temporarily Remove an Instance \(Console\) \(p. 124\)](#)
- [Temporarily Remove an Instance \(AWS CLI\) \(p. 124\)](#)

How the Standby State Works

The standby state works as follows to help you temporarily remove an instance from your Auto Scaling group:

1. You put the instance into the standby state. The instance remains in this state until you exit the standby state.
2. If there is a load balancer or target group attached to your Auto Scaling group, the instance is deregistered from the load balancer or target group.
3. By default, the value that you specified as your desired capacity is decremented when you put an instance on standby. This prevents the launch of an additional instance while you have this instance on standby. Alternatively, you can specify that your desired capacity is not decremented. If you specify this option, the Auto Scaling group launches an instance to replace the one on standby. The intention is to help you maintain capacity for your application while one or more instances are on standby.
4. You can update or troubleshoot the instance.
5. You return the instance to service by exiting the standby state.
6. After you put an instance that was on standby back in service, the desired capacity is incremented. If you did not decrement the capacity when you put the instance on standby, the Auto Scaling group detects that you have more instances than you need. It applies the termination policy in effect to reduce the size of the group. For more information, see [Controlling Which Auto Scaling Instances Terminate During Scale In \(p. 110\)](#).
7. If there is a load balancer or target group attached to your Auto Scaling group, the instance is registered with the load balancer or target group.

The following illustration shows the transitions between instance states in this process:



For more information about the complete lifecycle of instances in an Auto Scaling group, see [Auto Scaling Lifecycle](#) (p. 7).

Health Status of an Instance in a Standby State

Amazon EC2 Auto Scaling does not perform health checks on instances that are in a standby state. While the instance is in a standby state, its health status reflects the status that it had before you put it on standby. Amazon EC2 Auto Scaling does not perform a health check on the instance until you put it back in service.

For example, if you put a healthy instance on standby and then terminate it, Amazon EC2 Auto Scaling continues to report the instance as healthy. If you return the terminated instance to service, Amazon EC2 Auto Scaling performs a health check on the instance, determines that it is terminating and unhealthy, and launches a replacement instance.

Temporarily Remove an Instance (Console)

The following procedure demonstrates the general process for updating an instance that is currently in service.

To temporarily remove an instance

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. On the **Instances** tab, select the instance.
5. Choose **Actions, Set to Standby**.
6. On the **Set to Standby** page, select the check box to launch a replacement instance. Leave it unchecked to decrement the desired capacity. Choose **Set to Standby**.
7. You can update or troubleshoot your instance as needed. When you have finished, continue with the next step to return the instance to service.
8. Select the instance, choose **Actions, Set to InService**. On the **Set to InService** page, choose **Set to InService**.

Temporarily Remove an Instance (AWS CLI)

The following procedure demonstrates the general process for updating an instance that is currently in service.

To temporarily remove an instance

1. Use the following [describe-auto-scaling-instances](#) command to identify the instance to update:

```
aws autoscaling describe-auto-scaling-instances
```

The following is an example response:

```
{
  "AutoScalingInstances": [
    {
      "ProtectedFromScaleIn": false,
      "AvailabilityZone": "us-west-2a",
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "InstanceId": "i-05b4f7d5be44822a6",
      "AutoScalingGroupName": "my-asg",
      "HealthStatus": "HEALTHY",
      "LifecycleState": "InService"
    },
    ...
  ]
}
```

2. Move the instance into a Standby state using the following **enter-standby** command. The `--should-decrement-desired-capacity` option decreases the desired capacity so that the Auto Scaling group does not launch a replacement instance.

```
aws autoscaling enter-standby --instance-ids i-05b4f7d5be44822a6 \
  --auto-scaling-group-name my-asg --should-decrement-desired-capacity
```

The following is an example response:

```
{
  "Activities": [
    {
      "Description": "Moving EC2 instance to Standby: i-05b4f7d5be44822a6",
      "AutoScalingGroupName": "my-asg",
      "ActivityId": "3b1839fe-24b0-40d9-80ae-bcd883c2be32",
      "Details": "{\"Availability Zone\":\"us-west-2a\"}",
      "StartTime": "2014-12-15T21:31:26.150Z",
      "Progress": 50,
      "Cause": "At 2014-12-15T21:31:26Z instance i-05b4f7d5be44822a6 was moved to
standby
      in response to a user request, shrinking the capacity from 4 to 3.",
      "StatusCode": "InProgress"
    }
  ]
}
```

3. (Optional) Verify that the instance is in Standby using the following **describe-auto-scaling-instances** command.

```
aws autoscaling describe-auto-scaling-instances --instance-ids i-05b4f7d5be44822a6
```

The following is an example response. Notice that the status of the instance is now Standby.

```
{
  "AutoScalingInstances": [
```

```
{
  "ProtectedFromScaleIn": false,
  "AvailabilityZone": "us-west-2a",
  "LaunchTemplate": {
    "LaunchTemplateName": "my-launch-template",
    "Version": "1",
    "LaunchTemplateId": "lt-050555ad16a3f9c7f"
  },
  "InstanceId": "i-05b4f7d5be44822a6",
  "AutoScalingGroupName": "my-asg",
  "HealthStatus": "HEALTHY",
  "LifecycleState": "Standby"
},
...
]
```

4. You can update or troubleshoot your instance as needed. When you have finished, continue with the next step to return the instance to service.
5. Put the instance back in service using the following **exit-standby** command.

```
aws autoscaling exit-standby --instance-ids i-05b4f7d5be44822a6 --auto-scaling-group-name my-asg
```

The following is an example response:

```
{
  "Activities": [
    {
      "Description": "Moving EC2 instance out of Standby: i-05b4f7d5be44822a6",
      "AutoScalingGroupName": "my-asg",
      "ActivityId": "db12b166-cdcc-4c54-8aac-08c5935f8389",
      "Details": "{\n\"Availability Zone\":\n\"us-west-2a\n\"}",
      "StartTime": "2014-12-15T21:46:14.678Z",
      "Progress": 30,
      "Cause": "At 2014-12-15T21:46:14Z instance i-05b4f7d5be44822a6 was moved out of standby in response to a user request, increasing the capacity from 3 to 4.",
      "StatusCode": "PreInService"
    }
  ]
}
```

6. (Optional) Verify that the instance is back in service using the following **describe-auto-scaling-instances** command.

```
aws autoscaling describe-auto-scaling-instances --instance-ids i-05b4f7d5be44822a6
```

The following is an example response. Notice that the status of the instance is **InService**.

```
{
  "AutoScalingInstances": [
    {
      "ProtectedFromScaleIn": false,
      "AvailabilityZone": "us-west-2a",
      "LaunchTemplate": {
        "LaunchTemplateName": "my-launch-template",
        "Version": "1",
        "LaunchTemplateId": "lt-050555ad16a3f9c7f"
      },
      "InstanceId": "i-05b4f7d5be44822a6",

```

```
        "AutoScalingGroupName": "my-asg",  
        "HealthStatus": "HEALTHY",  
        "LifecycleState": "InService"  
    },  
    ...  
]  
}
```

Suspending and Resuming Scaling Processes

This topic explains how to suspend and then resume one or more of the scaling processes for your Auto Scaling group. It also describes the issues to consider when choosing to use the suspend-resume feature of Amazon EC2 Auto Scaling.

Important

Use the standby feature instead of the suspend-resume feature if you need to troubleshoot or reboot an instance. For more information, see [Temporarily Removing Instances from Your Auto Scaling Group \(p. 123\)](#). Use the instance protection feature to prevent specific instances from being terminated during automatic scale in. For more information, see [Instance Protection \(p. 113\)](#).

In addition to suspensions that you initiate, Amazon EC2 Auto Scaling can also suspend processes for Auto Scaling groups that repeatedly fail to launch instances. This is known as an *administrative suspension*. An administrative suspension most commonly applies to Auto Scaling groups that have been trying to launch instances for over 24 hours but have not succeeded in launching any instances. You can resume processes that were suspended by Amazon EC2 Auto Scaling for administrative reasons.

Contents

- [Scaling Processes \(p. 127\)](#)
- [Choosing to Suspend \(p. 128\)](#)
- [Suspend and Resume Scaling Processes \(Console\) \(p. 130\)](#)
- [Suspend and Resume Scaling Processes \(AWS CLI\) \(p. 130\)](#)

Scaling Processes

For Amazon EC2 Auto Scaling, there are two primary process types: `Launch` and `Terminate`. The `Launch` process adds a new Amazon EC2 instance to an Auto Scaling group, increasing its capacity, and the `Terminate` process removes an Amazon EC2 instance from the group, decreasing its capacity.

The other process types for Amazon EC2 Auto Scaling relate to specific scaling features:

- `AddToLoadBalancer`—Adds instances to the attached load balancer or target group when they are launched.
- `AlarmNotification`—Accepts notifications from CloudWatch alarms that are associated with the group's scaling policies.
- `AZRebalance`—Balances the number of EC2 instances in the group evenly across all of the specified Availability Zones when the group becomes unbalanced, for example, a previously unavailable Availability Zone returns to a healthy state. For more information, see [Rebalancing Activities \(p. 7\)](#).
- `HealthCheck`—Checks the health of the instances and marks an instance as unhealthy if Amazon EC2 or Elastic Load Balancing tells Amazon EC2 Auto Scaling that the instance is unhealthy. This process can override the health status of an instance that you set manually. For more information, see [Health Checks for Auto Scaling Instances \(p. 132\)](#).

- `ReplaceUnhealthy`—Terminates instances that are marked as unhealthy and then creates new instances to replace them.
- `ScheduledActions`—Performs the scheduled scaling actions that you create or that are created by the predictive scaling feature of AWS Auto Scaling.

Choosing to Suspend

Each process type can be suspended and resumed independently. This section provides some guidance and behavior to take into account before deciding to suspend a scaling process. Keep in mind that suspending individual processes may interfere with other processes. Depending on the reason for suspending a process, you might need to suspend multiple processes together.

The following descriptions explain what happens when individual process types are suspended.

Warning

If you suspend either the `Launch` or `Terminate` process types, it can prevent other process types from functioning properly.

`Terminate`

- Your Auto Scaling group does not scale in for alarms or scheduled actions that occur while the process is suspended. In addition, the following processes are disrupted:
 - `AZRebalance` is still active but does not function properly. It can launch new instances without terminating the old ones. This could cause your Auto Scaling group to grow up to 10 percent larger than its maximum size, because this is allowed temporarily during rebalancing activities. Your Auto Scaling group could remain above its maximum size until you resume the `Terminate` process. When `Terminate` resumes, `AZRebalance` gradually rebalances the Auto Scaling group if the group is no longer balanced between Availability Zones or if different Availability Zones are specified.
 - `ReplaceUnhealthy` is inactive but not `HealthCheck`. When `Terminate` resumes, the `ReplaceUnhealthy` process immediately starts running. If any instances were marked as unhealthy while `Terminate` was suspended, they will be immediately replaced.

`Launch`

- Your Auto Scaling group does not scale out for alarms or scheduled actions that occur while the process is suspended. `AZRebalance` stops rebalancing the group. `ReplaceUnhealthy` continues to terminate unhealthy instances, but does not launch replacements. When you resume `Launch`, rebalancing activities and health check replacements are handled in the following way:
 - `AZRebalance` gradually rebalances the Auto Scaling group if the group is no longer balanced between Availability Zones or if different Availability Zones are specified.
 - `ReplaceUnhealthy` immediately replaces any instances that it terminated during the time that `Launch` was suspended.

`AddToLoadBalancer`

- Amazon EC2 Auto Scaling launches the instances but does not add them to the load balancer or target group. When you resume the `AddToLoadBalancer` process, it resumes adding instances to the load balancer or target group when they are launched. However, it does not add the instances that were launched while this process was suspended. You must register those instances manually.

`AlarmNotification`

- Amazon EC2 Auto Scaling does not execute scaling policies when a CloudWatch alarm threshold is in breach. Suspending `AlarmNotification` allows you to temporarily stop scaling events triggered

by the group's scaling policies without deleting the scaling policies or their associated CloudWatch alarms. When you resume `AlarmNotification`, Amazon EC2 Auto Scaling considers policies with alarm thresholds that are currently in breach.

`AZRebalance`

- Your Auto Scaling group does not attempt to redistribute instances after certain events. However, if a scale-out or scale-in event occurs, the scaling process still tries to balance the Availability Zones. For example, during scale out, it launches the instance in the Availability Zone with the fewest instances. If the group becomes unbalanced while `AZRebalance` is suspended and you resume it, Amazon EC2 Auto Scaling attempts to rebalance the group. It first calls `Launch` and then `Terminate`.

`HealthCheck`

- Amazon EC2 Auto Scaling stops marking instances unhealthy as a result of EC2 and Elastic Load Balancing health checks. Your custom health checks continue to function properly, however. After you suspend `HealthCheck`, if you need to, you can manually set the health state of instances in your group and have `ReplaceUnhealthy` replace them.

`ReplaceUnhealthy`

- Amazon EC2 Auto Scaling stops replacing instances that are marked as unhealthy. Instances that fail EC2 or Elastic Load Balancing health checks will still be marked as unhealthy. As soon as you resume the `ReplaceUnhealthy` process, Amazon EC2 Auto Scaling replaces instances that were marked unhealthy while this process was suspended. The `ReplaceUnhealthy` process calls both of the primary process types—first `Terminate` and then `Launch`.

`ScheduledActions`

- Amazon EC2 Auto Scaling does not execute scaling actions that are scheduled to run during the suspension period. When you resume `ScheduledActions`, Amazon EC2 Auto Scaling only considers scheduled actions whose execution time has not yet passed.

Suspending Both Launch and Terminate

When you suspend the `Launch` and `Terminate` process types together, the following happens:

- Your Auto Scaling group cannot initiate scaling activities or maintain its desired capacity.
- If the group becomes unbalanced between Availability Zones, Amazon EC2 Auto Scaling does not attempt to redistribute instances evenly between the Availability Zones that are specified for your Auto Scaling group.
- Your Auto Scaling group cannot replace instances that are marked unhealthy.

When you resume the `Launch` and `Terminate` process types, Amazon EC2 Auto Scaling replaces instances that were marked unhealthy while the processes were suspended and may attempt to rebalance the group. Scaling activities will also resume.

Additional Considerations

There are some outside operations that may be affected while `Launch` and `Terminate` are suspended.

- **Spot Instance Interruptions**—If `Terminate` is suspended and your Auto Scaling group has Spot Instances, they can still terminate in the event that Spot capacity is no longer available. While `Launch`

is suspended, Amazon EC2 Auto Scaling cannot launch replacement instances from another Spot Instance pool or from the same Spot Instance pool when it is available again.

- **Attaching and Detaching Instances**—When `Launch` and `Terminate` are suspended, you can detach instances that are attached to your Auto Scaling group, but you can't attach new instances to the group. To attach instances, you must first resume `Launch`.

Note

If detaching an instance will immediately be followed by manually terminating it, you can call the `terminate-instance-in-auto-scaling-group` CLI command instead. This terminates the specified instance and optionally adjusts the group's desired capacity. In addition, if the Auto Scaling group is being used with lifecycle hooks, the custom actions that you specified for instance termination will run before the instance is fully terminated.

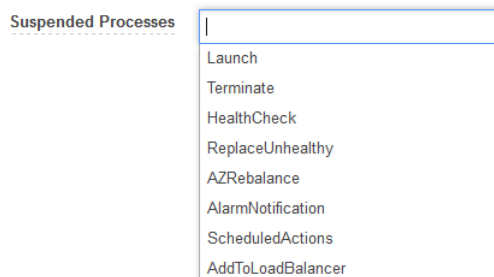
- **Standby Instances**—While `Launch` is suspended, you cannot return an instance in the `Standby` state to service. To return the instance to service, you must first resume `Launch`.

Suspend and Resume Scaling Processes (Console)


You can suspend and resume individual processes or all processes.

To suspend and resume processes

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select the Auto Scaling group.
4. On the **Details** tab, choose **Edit**.
5. For **Suspended Processes**, select the process to suspend.



To resume a suspended process, remove it from **Suspended Processes**.



6. Choose **Save**.

Suspend and Resume Scaling Processes (AWS CLI)

You can suspend and resume individual processes or all processes.

To suspend a process

Use the `suspend-processes` command with the `--scaling-processes` option as follows:

```
aws autoscaling suspend-processes --auto-scaling-group-name my-asg --scaling-processes AlarmNotification
```

To suspend all processes

Use the **suspend-processes** command as follows (omitting the `--scaling-processes` option):

```
aws autoscaling suspend-processes --auto-scaling-group-name my-asg
```

To resume a suspended process

Use the **resume-processes** command as follows:

```
aws autoscaling resume-processes --auto-scaling-group-name my-asg --scaling-  
processes AlarmNotification
```

To resume all suspended processes

Use the **resume-processes** command as follows (omitting the `--scaling-processes` option):

```
aws autoscaling resume-processes --auto-scaling-group-name my-asg
```

Monitoring Your Auto Scaling Instances and Groups

You can use the following features to monitor your Auto Scaling instances and groups.

Health checks

Amazon EC2 Auto Scaling periodically performs health checks on the instances in your Auto Scaling group and identifies any instances that are unhealthy. You can configure Auto Scaling to determine the health status of an instance using Amazon EC2 status checks, Elastic Load Balancing health checks, or custom health checks. For more information, see [Health Checks for Auto Scaling Instances \(p. 132\)](#).

CloudWatch metrics

Amazon EC2 Auto Scaling publishes data points to Amazon CloudWatch about your Auto Scaling groups. CloudWatch enables you to retrieve statistics about those data points as an ordered set of time series data, known as *metrics*. You can use these metrics to verify that your system is performing as expected. For more information, see [Monitoring Your Auto Scaling Groups and Instances Using Amazon CloudWatch \(p. 135\)](#).

CloudWatch Events

Amazon EC2 Auto Scaling can submit events to Amazon CloudWatch Events when your Auto Scaling groups launch or terminate instances, or when a lifecycle action occurs. This enables you to invoke a Lambda function when the event occurs. For more information, see [Getting CloudWatch Events When Your Auto Scaling Group Scales \(p. 140\)](#).

SNS notifications

Amazon EC2 Auto Scaling can send Amazon SNS notifications when your Auto Scaling groups launch or terminate instances. For more information, see [Getting Amazon SNS Notifications When Your Auto Scaling Group Scales \(p. 145\)](#).

CloudTrail logs

AWS CloudTrail enables you to track the calls made to the Amazon EC2 Auto Scaling API by or on behalf of your AWS account. CloudTrail stores the information in log files in the Amazon S3 bucket that you specify. You can use these log files to monitor activity of your Auto Scaling groups. Logs include which requests were made, the source IP addresses where the requests came from, who made the request, when the request was made, and so on. For more information, see [Logging Amazon EC2 Auto Scaling API Calls with AWS CloudTrail \(p. 150\)](#).

Health Checks for Auto Scaling Instances

The health status of an Auto Scaling instance is either healthy or unhealthy. All instances in your Auto Scaling group start in the healthy state. Instances are assumed to be healthy unless Amazon EC2 Auto Scaling receives notification that they are unhealthy. This notification can come from one or more of the following sources: Amazon EC2, Elastic Load Balancing, or a custom health check.

After Amazon EC2 Auto Scaling marks an instance as unhealthy, it is scheduled for replacement. If you do not want instances to be replaced, you can suspend the health check process for any individual Auto Scaling group.

Instance Health Status

Amazon EC2 Auto Scaling can determine the health status of an instance using one or more of the following:

- Status checks provided by Amazon EC2 to identify hardware and software issues that may impair an instance. This includes both instance status checks and system status checks. For more information, see [Types of Status Checks](#) in the *Amazon EC2 User Guide for Linux Instances*.
- Health checks provided by Elastic Load Balancing.
- Your custom health checks.

The EC2 status checks are the default health checks for Amazon EC2 Auto Scaling and do not require any special configuration. You can customize the default health checks conducted by your Auto Scaling group by specifying additional checks, however. For more information, see [Adding Elastic Load Balancing Health Checks to an Auto Scaling Group](#) (p. 65) and [Custom Health Checks](#) (p. 134).

Determining Instance Health

After an instance is fully configured and passes the initial health checks, it is considered healthy by Amazon EC2 Auto Scaling and enters the `InService` state. Amazon EC2 Auto Scaling periodically performs health checks on the instances in your Auto Scaling group and identifies any instances that are unhealthy.

Amazon EC2 Auto Scaling health checks use the results of the Amazon EC2 status checks to determine the health status of an instance. If the instance is in any state other than `running` or if the system status is `impaired`, Amazon EC2 Auto Scaling considers the instance to be unhealthy and launches a replacement instance. This includes when the instance has any of the following states:

- `stopping`
- `stopped`
- `terminating`
- `terminated`

If you attached a load balancer or target group to your Auto Scaling group, you can configure the group to mark an instance as unhealthy when Elastic Load Balancing reports it as `OutOfService`. If connection draining is enabled for your load balancer, Amazon EC2 Auto Scaling waits for in-flight requests to complete or the maximum timeout to expire, whichever comes first, before terminating instances due to a scaling event or health check replacement. If you configure your Auto Scaling group to use the Elastic Load Balancing health checks, Amazon EC2 Auto Scaling determines the health status of the instances by checking both the EC2 status checks and the Elastic Load Balancing health checks. For more information, see [Adding Elastic Load Balancing Health Checks to an Auto Scaling Group](#) (p. 65).

If you have custom health checks, you can send the information from your health checks to Amazon EC2 Auto Scaling so that Amazon EC2 Auto Scaling can use this information. For example, if you determine that an instance is not functioning as expected, you can set the health status of the instance to `Unhealthy`. The next time that Amazon EC2 Auto Scaling performs a health check on the instance, it will determine that the instance is unhealthy and then launch a replacement instance.

Health Check Grace Period

Frequently, an Auto Scaling instance that has just come into service needs to warm up before it can pass the health check. Amazon EC2 Auto Scaling waits until the health check grace period ends before checking the health status of the instance. Amazon EC2 status checks and Elastic Load Balancing health

checks can complete before the health check grace period expires. However, Amazon EC2 Auto Scaling does not act on them until the health check grace period expires. To provide ample warm-up time for your instances, ensure that the health check grace period covers the expected startup time for your application. If you add a lifecycle hook, the grace period does not start until the lifecycle hook actions are completed and the instance enters the `InService` state.

Replacing Unhealthy Instances

After an instance has been marked unhealthy because of an Amazon EC2 or Elastic Load Balancing health check, it is almost immediately scheduled for replacement. It never automatically recovers its health. You can intervene manually by calling the `SetInstanceHealth` action (or the **set-instance-health** command) to set the instance's health status back to healthy. If the instance is already terminating, you get an error. Because the interval between marking an instance unhealthy and its actual termination is so small, attempting to set an instance's health status back to healthy with the `SetInstanceHealth` action (or, **set-instance-health** command) is probably useful only for a suspended group. For more information, see [Suspending and Resuming Scaling Processes](#) (p. 127).

Amazon EC2 Auto Scaling creates a new scaling activity for terminating the unhealthy instance and then terminates it. Later, another scaling activity launches a new instance to replace the terminated instance.

When your instance is terminated, any associated Elastic IP addresses are disassociated and are not automatically associated with the new instance. You must associate these Elastic IP addresses with the new instance manually. Similarly, when your instance is terminated, its attached EBS volumes are detached. You must attach these EBS volumes to the new instance manually. For more information, see [Disassociating an Elastic IP Address and Reassociating with a Different Instance](#) and [Attaching an Amazon EBS Volume to an Instance](#) in the *Amazon EC2 User Guide for Linux Instances*.

Custom Health Checks

If you have your own health check system, you can send the instance's health information directly from your system to Amazon EC2 Auto Scaling.

Use the following **set-instance-health** command to set the health state of the specified instance to Unhealthy.

```
aws autoscaling set-instance-health --instance-id i-123abc45d --health-status Unhealthy
```

Use the following **describe-auto-scaling-groups** command to verify that the instance state is Unhealthy.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-names my-asg
```

The following is an example response that shows that the health status of the instance is Unhealthy and that the instance is terminating.

```
{
  "AutoScalingGroups": [
    {
      ...
      "Instances": [
        {
          "InstanceId": "i-123abc45d",
          "AvailabilityZone": "us-west-2a",
          "HealthStatus": "Unhealthy",
          "LifecycleState": "Terminating",
          "LaunchConfigurationName": "my-lc"
        }
      ]
    }
  ]
}
```

```
    },  
    ...  
  ]  
}  
]
```

Monitoring Your Auto Scaling Groups and Instances Using Amazon CloudWatch

Amazon CloudWatch enables you to retrieve statistics as an ordered set of time-series data, known as metrics. You can use these metrics to verify that your system is performing as expected.

Amazon EC2 sends metrics to CloudWatch that describe your Auto Scaling instances. These metrics are available for any EC2 instance, not just those in an Auto Scaling group. For more information, see [Instance Metrics](#) in the *Amazon EC2 User Guide for Linux Instances*.

Auto Scaling groups can send metrics to CloudWatch that describe the group itself. You must enable these metrics.

Contents

- [Auto Scaling Group Metrics \(p. 135\)](#)
- [Dimensions for Auto Scaling Group Metrics \(p. 136\)](#)
- [Enable Auto Scaling Group Metrics \(p. 136\)](#)
- [Configure Monitoring for Auto Scaling Instances \(p. 137\)](#)
- [View CloudWatch Metrics \(p. 138\)](#)
- [Create Amazon CloudWatch Alarms \(p. 139\)](#)

Auto Scaling Group Metrics

The `AWS/AutoScaling` namespace includes the following metrics.

Metric	Description
<code>GroupMinSize</code>	The minimum size of the Auto Scaling group.
<code>GroupMaxSize</code>	The maximum size of the Auto Scaling group.
<code>GroupDesiredCapacity</code>	The number of instances that the Auto Scaling group attempts to maintain.
<code>GroupInServiceInstances</code>	The number of instances that are running as part of the Auto Scaling group. This metric does not include instances that are pending or terminating.
<code>GroupPendingInstances</code>	The number of instances that are pending. A pending instance is not yet in service. This metric does not include instances that are in service or terminating.
<code>GroupStandbyInstances</code>	The number of instances that are in a <code>Standby</code> state. Instances in this state are still running but are not actively in service.

Metric	Description
GroupTerminatingInstances	The number of instances that are in the process of terminating. This metric does not include instances that are in service or pending.
GroupTotalInstances	The total number of instances in the Auto Scaling group. This metric identifies the number of instances that are in service, pending, and terminating.

Dimensions for Auto Scaling Group Metrics


To filter the metrics for your Auto Scaling group by group name, use the `AutoScalingGroupName` dimension.

Enable Auto Scaling Group Metrics

When you enable Auto Scaling group metrics, Auto Scaling sends sampled data to CloudWatch every minute.


To enable group metrics (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Monitoring** tab, for **Auto Scaling Metrics**, choose **Enable Group Metrics Collection**. If you don't see this option, select **Auto Scaling** for **Display**.



To disable group metrics (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Monitoring** tab, for **Auto Scaling Metrics**, choose **Disable Group Metrics Collection**. If you don't see this option, select **Auto Scaling** for **Display**.



To enable group metrics (AWS CLI)

Enable one or more group metrics using the [enable-metrics-collection](#) command. For example, the following command enables the GroupDesiredCapacity metric.

```
aws autoscaling enable-metrics-collection --auto-scaling-group-name my-asg \  
--metrics GroupDesiredCapacity --granularity "1Minute"
```

If you omit the `--metrics` option, all metrics are enabled.

```
aws autoscaling enable-metrics-collection --auto-scaling-group-name my-asg \  
--granularity "1Minute"
```

To disable group metrics (AWS CLI)

Use the [disable-metrics-collection](#) command. For example, the following command disables all Auto Scaling group metrics.

```
aws autoscaling disable-metrics-collection --auto-scaling-group-name my-asg
```

Configure Monitoring for Auto Scaling Instances

You configure monitoring for EC2 instances using a launch configuration or template. Monitoring is enabled whenever an instance is launched, either basic monitoring (5-minute granularity) or detailed monitoring (1-minute granularity). For detailed monitoring, additional charges apply. For more information, see [Amazon CloudWatch Pricing](#).

Note

By default, basic monitoring is enabled when you create a launch template or when you use the AWS Management Console to create a launch configuration. Detailed monitoring is enabled by default when you create a launch configuration using the AWS CLI or an SDK.

To change the type of monitoring enabled on new EC2 instances, update the launch template or update the Auto Scaling group to use a new launch configuration. Existing instances continue to use the previously enabled monitoring type. To update all instances, terminate them so that they are replaced by your Auto Scaling group or update instances individually using [monitor-instances](#) and [unmonitor-instances](#).

If you have CloudWatch alarms associated with your Auto Scaling group, use the [put-metric-alarm](#) command to update each alarm. Make each period match the monitoring type (300 seconds for basic monitoring and 60 seconds for detailed monitoring). If you change from detailed monitoring to basic monitoring but do not update your alarms to match the five-minute period, they continue to check for statistics every minute. They might find no data available for as many as four out of every five periods.

To configure CloudWatch monitoring (console)

When you create the launch configuration using the AWS Management Console, on the **Configure Details** page, select **Enable CloudWatch detailed monitoring**. Otherwise, basic monitoring is enabled. For more information, see [Creating a Launch Configuration](#) (p. 32).

To enable detailed monitoring for a launch template using the AWS Management Console, in the **Advanced Details** section, for **Monitoring**, choose **Enable**. Otherwise, basic monitoring is enabled. For more information, see [Creating a Launch Template for an Auto Scaling Group](#) (p. 24).

To configure CloudWatch monitoring (AWS CLI)

For launch configurations, use the [create-launch-configuration](#) command with the `--instance-monitoring` option. Set this option to `true` to enable detailed monitoring or `false` to enable basic monitoring.


```
--instance-monitoring Enabled=true
```

For launch templates, use the [create-launch-template](#) command and pass a JSON file that contains the information for creating the launch template. Set the monitoring attribute to "Monitoring": {"Enabled": true} to enable detailed monitoring or "Monitoring": {"Enabled": false} to enable basic monitoring.

View CloudWatch Metrics

You can view the CloudWatch metrics for your Auto Scaling groups and instances using the Amazon EC2 console. These metrics are displayed as monitoring graphs.

Alternatively, you can view these metrics using the CloudWatch console.

To view metrics using the Amazon EC2 console

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. Choose the **Monitoring** tab.
5. (Optional) To filter the results by time, select a time range from **Showing data for**.
6. To view the metrics for your groups, for **Display**, choose **Auto Scaling**. To get a larger view of a single metric, select its graph. The following metrics are available for groups:
 - Minimum Group Size — `GroupMinSize`
 - Maximum Group Size — `GroupMaxSize`
 - Desired Capacity — `GroupDesiredCapacity`
 - In Service Instances — `GroupInServiceInstances`
 - Pending Instances — `GroupPendingInstances`
 - Standby Instances — `GroupStandbyInstances`
 - Terminating Instances — `GroupTerminatingInstances`
 - Total Instances — `GroupTotalInstances`
7. To view metrics for your instances, for **Display**, choose **EC2**. To get a larger view of a single metric, select its graph. The following metrics are available for instances:
 - CPU Utilization — `CPUUtilization`
 - Disk Reads — `DiskReadBytes`
 - Disk Read Operations — `DiskReadOps`
 - Disk Writes — `DiskWriteBytes`
 - Disk Write Operations — `DiskWriteOps`
 - Network In — `NetworkIn`
 - Network Out — `NetworkOut`
 - Status Check Failed (Any) — `StatusCheckFailed`
 - Status Check Failed (Instance) — `StatusCheckFailed_Instance`
 - Status Check Failed (System) — `StatusCheckFailed_System`

To view metrics using the CloudWatch console

For more information, see [Aggregating Statistics by Auto Scaling Group](#).

To view CloudWatch metrics (AWS CLI)

To view all metrics for all your Auto Scaling groups, use the following [list-metrics](#) command.

```
aws cloudwatch list-metrics --namespace "AWS/AutoScaling"
```

To view the metrics for a single Auto Scaling group, specify the `AutoScalingGroupName` dimension as follows.

```
aws cloudwatch list-metrics --namespace "AWS/AutoScaling" --dimensions  
Name=AutoScalingGroupName,Value=my-asg
```

To view a single metric for all your Auto Scaling groups, specify the name of the metric as follows.

```
aws cloudwatch list-metrics --namespace "AWS/AutoScaling" --metric-name  
GroupDesiredCapacity
```

Create Amazon CloudWatch Alarms

One purpose for monitoring metrics is to verify that your application is performing as expected. If a metric goes beyond what you consider an acceptable threshold, you can have a CloudWatch alarm trigger an action. You can specify any of the alarm actions that are supported by CloudWatch.

You configure an alarm by identifying the metric to monitor. For example, you can configure an alarm to watch over the average CPU usage of the EC2 instances in your Auto Scaling group. The action can be a notification that is sent to you when the average CPU usage of the group breaches the threshold that you specified for the consecutive periods you specified. For example, if the metric stays at or above 70 percent for 4 consecutive periods of 1 minute each.

For more information, see [Using Amazon CloudWatch Alarms](#) in the *Amazon CloudWatch User Guide*.

To create a CloudWatch alarm based on average CPU utilization

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. In the navigation pane, choose **Alarms**, **Create Alarm**.
3. Choose **Select Metric** and then **EC2**.
4. Choose a metric category to filter the results. To see the Auto Scaling group metrics, choose **By Auto Scaling Group**. If you only want the metrics for individual instances, choose **Per-Instance Metrics**.
5. Select a metric as follows:
 - a. Select the row that contains the Auto Scaling group or instance that you want to create an alarm on and the **CPUUtilization** metric.
 - b. Choose the **Graphed metrics** tab.
 - c. Under **Statistic**, choose **Average**.
 - d. Under **Period**, choose the evaluation period for the alarm, for example, 1 minute. When evaluating the alarm, each period is aggregated into one data point.

Note

A shorter period creates a more sensitive alarm.
 - e. Choose **Select metric**.
6. Under **Conditions**, define the alarm by defining the threshold condition. For example, you can define a threshold to trigger the alarm whenever the value of the metric is greater than or equal to 70 percent.

7. Under **Additional configuration**, for **Datapoints to alarm**, specify how many datapoints (evaluation periods) must be in the ALARM state to trigger the alarm, for example, 4 out of 4. This creates an alarm that goes to ALARM state if that many consecutive periods are breaching.
8. For **Missing data treatment**, choose one of the options. For a metric that continually reports data, such as CPUUtilization, you might want to choose **Treat missing data as bad (breaching threshold)**, as missing datapoints may indicate that something is wrong. For more information, see [Configuring How CloudWatch Alarms Treat Missing Data](#) in the *Amazon CloudWatch User Guide*.
9. Choose **Next**.
10. Under **Configure actions**, define the action to take.
11. Choose **Next**.
12. Under **Add a description**, enter a name and description for the alarm and choose **Next**.
13. Choose **Create Alarm**.

Getting CloudWatch Events When Your Auto Scaling Group Scales

Amazon CloudWatch Events lets you automate AWS services and respond to system events such as application availability issues or resource changes. Events from AWS services are delivered to CloudWatch Events nearly in real time. You can write simple rules to indicate which events are of interest to you and what automated actions to take when an event matches a rule.

CloudWatch Events lets you set a variety of targets—such as a Lambda function or an Amazon SNS topic—which receive events in JSON format. For more information, see the [Amazon CloudWatch Events User Guide](#).

It is useful to know when Amazon EC2 Auto Scaling is launching or terminating the EC2 instances in your Auto Scaling group. You can configure Amazon EC2 Auto Scaling to send events to CloudWatch Events whenever your Auto Scaling group scales.

Note

You can also receive a two-minute warning when Spot Instances are about to be reclaimed by Amazon EC2. For an example of the event for Spot Instance interruption, see [Spot Instance Interruption Notices](#) in the *Amazon EC2 User Guide for Linux Instances*.

Contents

- [Auto Scaling Events](#) (p. 140)
- [Create a Lambda Function](#) (p. 144)
- [Route Events to Your Lambda Function](#) (p. 144)

Auto Scaling Events

Amazon EC2 Auto Scaling supports sending events to CloudWatch Events when the following events occur:

- [EC2 Instance-launch Lifecycle Action](#) (p. 141)
- [EC2 Instance Launch Successful](#) (p. 141)
- [EC2 Instance Launch Unsuccessful](#) (p. 142)
- [EC2 Instance-terminate Lifecycle Action](#) (p. 142)
- [EC2 Instance Terminate Successful](#) (p. 143)
- [EC2 Instance Terminate Unsuccessful](#) (p. 143)

EC2 Instance-launch Lifecycle Action

Amazon EC2 Auto Scaling moved an instance to a Pending:Wait state due to a lifecycle hook.

Event Data

The following is example data for this event.

```
{
  "version": "0",
  "id": "12345678-1234-1234-1234-123456789012",
  "detail-type": "EC2 Instance-launch Lifecycle Action",
  "source": "aws.autoscaling",
  "account": "123456789012",
  "time": "yyyy-mm-ddThh:mm:ssZ",
  "region": "us-west-2",
  "resources": [
    "auto-scaling-group-arn"
  ],
  "detail": {
    "LifecycleActionToken": "87654321-4321-4321-4321-210987654321",
    "AutoScalingGroupName": "my-asg",
    "LifecycleHookName": "my-lifecycle-hook",
    "EC2InstanceId": "i-1234567890abcdef0",
    "LifecycleTransition": "autoscaling:EC2_INSTANCE_LAUNCHING",
    "NotificationMetadata": "additional-info"
  }
}
```

EC2 Instance Launch Successful

Amazon EC2 Auto Scaling successfully launched an instance.

Event Data

The following is example data for this event.

```
{
  "version": "0",
  "id": "12345678-1234-1234-1234-123456789012",
  "detail-type": "EC2 Instance Launch Successful",
  "source": "aws.autoscaling",
  "account": "123456789012",
  "time": "yyyy-mm-ddThh:mm:ssZ",
  "region": "us-west-2",
  "resources": [
    "auto-scaling-group-arn",
    "instance-arn"
  ],
  "detail": {
    "StatusCode": "InProgress",
    "Description": "Launching a new EC2 instance: i-12345678",
    "AutoScalingGroupName": "my-auto-scaling-group",
    "ActivityId": "87654321-4321-4321-4321-210987654321",
    "Details": {
      "Availability Zone": "us-west-2b",
      "Subnet ID": "subnet-12345678"
    },
    "RequestId": "12345678-1234-1234-1234-123456789012",
    "StatusMessage": "",
    "EndTime": "yyyy-mm-ddThh:mm:ssZ",
    "EC2InstanceId": "i-1234567890abcdef0",
    "StartTime": "yyyy-mm-ddThh:mm:ssZ",
  }
}
```

```
    "Cause": "description-text"
  }
}
```

EC2 Instance Launch Unsuccessful

Amazon EC2 Auto Scaling failed to launch an instance.

Event Data

The following is example data for this event.

```
{
  "version": "0",
  "id": "12345678-1234-1234-1234-123456789012",
  "detail-type": "EC2 Instance Launch Unsuccessful",
  "source": "aws.autoscaling",
  "account": "123456789012",
  "time": "yyyy-mm-ddThh:mm:ssZ",
  "region": "us-west-2",
  "resources": [
    "auto-scaling-group-arn",
    "instance-arn"
  ],
  "detail": {
    "StatusCode": "Failed",
    "AutoScalingGroupName": "my-auto-scaling-group",
    "ActivityId": "87654321-4321-4321-4321-210987654321",
    "Details": {
      "Availability Zone": "us-west-2b",
      "Subnet ID": "subnet-12345678"
    },
    "RequestId": "12345678-1234-1234-1234-123456789012",
    "StatusMessage": "message-text",
    "EndTime": "yyyy-mm-ddThh:mm:ssZ",
    "EC2InstanceId": "i-1234567890abcdef0",
    "StartTime": "yyyy-mm-ddThh:mm:ssZ",
    "Cause": "description-text"
  }
}
```

EC2 Instance-terminate Lifecycle Action

Amazon EC2 Auto Scaling moved an instance to a `Terminating:Wait` state due to a lifecycle hook.

Event Data

The following is example data for this event.

```
{
  "version": "0",
  "id": "12345678-1234-1234-1234-123456789012",
  "detail-type": "EC2 Instance-terminate Lifecycle Action",
  "source": "aws.autoscaling",
  "account": "123456789012",
  "time": "yyyy-mm-ddThh:mm:ssZ",
  "region": "us-west-2",
  "resources": [
    "auto-scaling-group-arn"
  ],
  "detail": {
    "LifecycleActionToken": "87654321-4321-4321-4321-210987654321",

```

```
{
  "AutoScalingGroupName": "my-asg",
  "LifecycleHookName": "my-lifecycle-hook",
  "EC2InstanceId": "i-1234567890abcdef0",
  "LifecycleTransition": "autoscaling:EC2_INSTANCE_TERMINATING",
  "NotificationMetadata": "additional-info"
}
```

EC2 Instance Terminate Successful

Amazon EC2 Auto Scaling successfully terminated an instance.

Event Data

The following is example data for this event.

```
{
  "version": "0",
  "id": "12345678-1234-1234-1234-123456789012",
  "detail-type": "EC2 Instance Terminate Successful",
  "source": "aws.autoscaling",
  "account": "123456789012",
  "time": "yyyy-mm-ddThh:mm:ssZ",
  "region": "us-west-2",
  "resources": [
    "auto-scaling-group-arn",
    "instance-arn"
  ],
  "detail": {
    "StatusCode": "InProgress",
    "Description": "Terminating EC2 instance: i-12345678",
    "AutoScalingGroupName": "my-auto-scaling-group",
    "ActivityId": "87654321-4321-4321-4321-210987654321",
    "Details": {
      "Availability Zone": "us-west-2b",
      "Subnet ID": "subnet-12345678"
    },
    "RequestId": "12345678-1234-1234-1234-123456789012",
    "StatusMessage": "",
    "EndTime": "yyyy-mm-ddThh:mm:ssZ",
    "EC2InstanceId": "i-1234567890abcdef0",
    "StartTime": "yyyy-mm-ddThh:mm:ssZ",
    "Cause": "description-text"
  }
}
```

EC2 Instance Terminate Unsuccessful

Amazon EC2 Auto Scaling failed to terminate an instance.

Event Data

The following is example data for this event.

```
{
  "version": "0",
  "id": "12345678-1234-1234-1234-123456789012",
  "detail-type": "EC2 Instance Terminate Unsuccessful",
  "source": "aws.autoscaling",
  "account": "123456789012",
  "time": "yyyy-mm-ddThh:mm:ssZ",
  "region": "us-west-2",

```

```
"resources": [
  "auto-scaling-group-arn",
  "instance-arn"
],
"detail": {
  "StatusCode": "Failed",
  "AutoScalingGroupName": "my-auto-scaling-group",
  "ActivityId": "87654321-4321-4321-4321-210987654321",
  "Details": {
    "Availability Zone": "us-west-2b",
    "Subnet ID": "subnet-12345678"
  },
  "RequestId": "12345678-1234-1234-1234-123456789012",
  "StatusMessage": "message-text",
  "EndTime": "yyyymm-ddThh:mm:ssZ",
  "EC2InstanceId": "i-1234567890abcdef0",
  "StartTime": "yyyymm-ddThh:mm:ssZ",
  "Cause": "description-text"
}
}
```

Create a Lambda Function

Use the following procedure to create a Lambda function to handle an Auto Scaling event.

To create a Lambda function

1. Open the AWS Lambda console at <https://console.aws.amazon.com/lambda/>.
2. If you are new to Lambda, you see a welcome page; choose **Get Started Now**; otherwise, choose **Create a Lambda function**.
3. On the **Select blueprint** page, type hello-world for **Filter**, and then select the **hello-world** blueprint.
4. On the **Configure triggers** page, choose **Next**.
5. On the **Configure function** page, do the following:
 - a. Type a name and description for the Lambda function.
 - b. Edit the code for the Lambda function. For example, the following code simply logs the event:

```
console.log('Loading function');

exports.handler = function(event, context) {
  console.log("AutoScalingEvent()");
  console.log("Event data:\n" + JSON.stringify(event, null, 4));
  context.succeed("...");
};
```

- c. For **Role**, choose **Choose an existing role** if you have an existing role that you'd like to use, and then choose your role from **Existing role**. Alternatively, to create a new role, choose one of the other options for **Role** and then follow the directions.
 - d. (Optional) For **Advanced settings**, make any changes that you need.
 - e. Choose **Next**.
6. On the **Review** page, choose **Create function**.

Route Events to Your Lambda Function

Use the following procedure to route Auto Scaling events to your Lambda function.

To route events to your Lambda function

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. On the navigation pane, choose **Events**.
3. Choose **Create rule**.
4. For **Event selector**, choose **Auto Scaling** as the event source. By default, the rule applies to all Auto Scaling events for all of your Auto Scaling groups. Alternatively, you can select specific events or a specific Auto Scaling group.
5. For **Targets**, choose **Add target**. Choose **Lambda function** as the target type, and then select your Lambda function.
6. Choose **Configure details**.
7. For **Rule definition**, type a name and description for your rule.
8. Choose **Create rule**.

To test your rule, change the size of your Auto Scaling group. If you used the example code for your Lambda function, it logs the event to CloudWatch Logs.

To test your rule

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, choose **Auto Scaling Groups**, and then select your Auto Scaling group.
3. On the **Details** tab, choose **Edit**.
4. Change the value of **Desired**, and then choose **Save**.
5. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
6. On the navigation pane, choose **Logs**.
7. Select the log group for your Lambda function (for example, `/aws/lambda/my-function`).
8. Select a log stream to view the event data. The data is displayed, similar to the following:

```
Event Data
▼ 2016-02-22T17:48:20.778Z ealfjqinxq6pwo9d Loading function
▼ START RequestId: 7560439b-d98c-11e5-932d-f52757e7aee0 Version: $LATEST
▼ 2016-02-22T17:48:20.813Z 7560439b-d98c-11e5-932d-f52757e7aee0 AutoScalingEvent()
▼ 2016-02-22T17:48:20.814Z 7560439b-d98c-11e5-932d-f52757e7aee0 Event data:
{
  "version": "0",
  "id": "df9b0c8c-89c8-4748-92cb-ac68a9029ada",
  "detail-type": "EC2 Instance Launch Successful",
  "source": "aws.autoscaling",
```

Getting Amazon SNS Notifications When Your Auto Scaling Group Scales

It is useful to know when Amazon EC2 Auto Scaling is launching or terminating the EC2 instances in your Auto Scaling group. Amazon SNS coordinates and manages the delivery or sending of notifications to subscribing clients or endpoints. You can configure Amazon EC2 Auto Scaling to send an SNS notification whenever your Auto Scaling group scales.

Amazon SNS can deliver notifications as HTTP or HTTPS POST, email (SMTP, either plaintext or in JSON format), or as a message posted to an Amazon SQS queue. For more information, see [What Is Amazon SNS](#) in the *Amazon Simple Notification Service Developer Guide*.

For example, if you configure your Auto Scaling group to use the `autoscaling:EC2_INSTANCE_TERMINATE` notification type, and your Auto Scaling group terminates an instance,

it sends an email notification. This email contains the details of the terminated instance, such as the instance ID and the reason that the instance was terminated.

Tip

If you prefer, you can use Amazon CloudWatch Events to configure a target to invoke a Lambda function when your Auto Scaling group scales or when a lifecycle action occurs. For more information, see [Getting CloudWatch Events When Your Auto Scaling Group Scales \(p. 140\)](#).

Contents

- [SNS Notifications \(p. 146\)](#)
- [Configure Amazon SNS \(p. 147\)](#)
- [Configure Your Auto Scaling Group to Send Notifications \(p. 147\)](#)
- [Test the Notification Configuration \(p. 148\)](#)
- [Verify That You Received Notification of the Scaling Event \(p. 148\)](#)
- [Delete the Notification Configuration \(p. 150\)](#)

SNS Notifications

Amazon EC2 Auto Scaling supports sending Amazon SNS notifications when the following events occur.

Event	Description
autoscaling:EC2_INSTANCE_LAUNCH	Successful instance launch
autoscaling:EC2_INSTANCE_LAUNCH_ERROR	Failed instance launch
autoscaling:EC2_INSTANCE_TERMINATE	Successful instance termination
autoscaling:EC2_INSTANCE_TERMINATE_ERROR	Failed instance termination

The message includes the following information:

- **Event** — The event.
- **AccountId** — The AWS account ID.
- **AutoScalingGroupName** — The name of the Auto Scaling group.
- **AutoScalingGroupARN** — The ARN of the Auto Scaling group.
- **EC2InstanceId** — The ID of the EC2 instance.

For example:

```
Service: AWS Auto Scaling
Time: 2016-09-30T19:00:36.414Z
RequestId: 4e6156f4-a9e2-4bda-a7fd-33f2ae528958
Event: autoscaling:EC2_INSTANCE_LAUNCH
AccountId: 123456789012
AutoScalingGroupName: my-asg
AutoScalingGroupARN: arn:aws:autoscaling:region:123456789012:autoScalingGroup...
ActivityId: 4e6156f4-a9e2-4bda-a7fd-33f2ae528958
Description: Launching a new EC2 instance: i-0598c7d356eba48d7
Cause: At 2016-09-30T18:59:38Z a user request update of AutoScalingGroup constraints to ...
StartTime: 2016-09-30T19:00:04.445Z
EndTime: 2016-09-30T19:00:36.414Z
StatusCode: InProgress
```

```
StatusMessage:  
Progress: 50  
EC2InstanceId: i-0598c7d356eba48d7  
Details: {"Subnet ID":"subnet-id","Availability Zone":"zone"}
```

Configure Amazon SNS

To use Amazon SNS to send email notifications, you must first create a *topic* and then subscribe your email addresses to the topic.

Create an Amazon SNS Topic

An SNS topic is a logical access point, a communication channel your Auto Scaling group uses to send the notifications. You create a topic by specifying a name for your topic.

For more information, see [Create a Topic](#) in the *Amazon Simple Notification Service Developer Guide*.

Subscribe to the Amazon SNS Topic

To receive the notifications that your Auto Scaling group sends to the topic, you must subscribe an endpoint to the topic. In this procedure, for **Endpoint**, specify the email address where you want to receive the notifications from Amazon EC2 Auto Scaling.

For more information, see [Subscribe to a Topic](#) in the *Amazon Simple Notification Service Developer Guide*.

Confirm Your Amazon SNS Subscription

Amazon SNS sends a confirmation email to the email address you specified in the previous step.

Make sure that you open the email from AWS Notifications and choose the link to confirm the subscription before you continue with the next step.

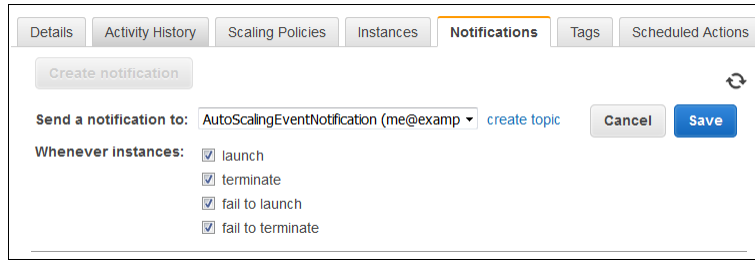
You will receive an acknowledgement message from AWS. Amazon SNS is now configured to receive notifications and send the notification as an email to the email address that you specified.

Configure Your Auto Scaling Group to Send Notifications

You can configure your Auto Scaling group to send notifications to Amazon SNS when a scaling event, such as launching instances or terminating instances, takes place. Amazon SNS sends a notification with information about the instances to the email address that you specified.

To configure Amazon SNS notifications for your Auto Scaling group (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Notifications** tab, choose **Create notification**.
5. On the **Create notifications** pane, do the following:
 - a. For **Send a notification to**, select your SNS topic.
 - b. For **Whenever instances**, select the events to send the notifications for.
 - c. Choose **Save**.



To configure Amazon SNS notifications for your Auto Scaling group (AWS CLI)

Use the following [put-notification-configuration](#) command.

```
aws autoscaling put-notification-configuration --auto-scaling-group-name my-asg --topic-arn arn --notification-types "autoscaling:EC2_INSTANCE_LAUNCH" "autoscaling:EC2_INSTANCE_TERMINATE"
```

Test the Notification Configuration

To generate a notification for a launch event, update the Auto Scaling group by increasing the desired capacity of the Auto Scaling group by 1. Amazon EC2 Auto Scaling launches the EC2 instance, and you receive an email notification within a few minutes.

To change the desired capacity (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Details** tab, choose **Edit**.
5. For **Desired**, increase the current value by 1. If this value exceeds **Max**, you must also increase the value of **Max** by 1.
6. Choose **Save**.
7. After a few minutes, you'll receive a notification email for the launch event. If you do not need the additional instance that you launched for this test, you can decrease **Desired** by 1. After a few minutes, you'll receive a notification email for the terminate event.

To change the desired capacity (AWS CLI)

Use the following [set-desired-capacity](#) command.

```
aws autoscaling set-desired-capacity --auto-scaling-group-name my-asg --desired-capacity 2
```

Verify That You Received Notification of the Scaling Event

Check your email for a message from Amazon SNS and open the email. After you receive notification of a scaling event for your Auto Scaling group, you can confirm the scaling event by looking at the description of your Auto Scaling group. You need information from the notification email, such as the ID of the instance that was launched or terminated.

To verify that your Auto Scaling group has launched new instance (console)

1. Select your Auto Scaling group.
2. On the **Activity History** tab, the **Status** column shows the current status of your instance. For example, if the notification indicates that an instance has launched, use the refresh button to verify that the status of the launch activity is **Successful**.
3. On the **Instances** tab, you can view the current **Lifecycle** state of the instance whose ID you received in the notification email. After a new instance starts, its lifecycle state changes to **InService**.

To verify that your Auto Scaling group has launched a new instance (AWS CLI)

Use the following [describe-auto-scaling-groups](#) command to confirm that the size of your Auto Scaling group has changed.

```
aws autoscaling describe-auto-scaling-groups --auto-scaling-group-name my-asg
```

The following example output shows that the group has two instances. Check for the instance whose ID you received in the notification email.

```
{
  "AutoScalingGroups": [
    {
      "AutoScalingGroupARN": "arn",
      "HealthCheckGracePeriod": 0,
      "SuspendedProcesses": [],
      "DesiredCapacity": 2,
      "Tags": [],
      "EnabledMetrics": [],
      "LoadBalancerNames": [],
      "AutoScalingGroupName": "my-asg",
      "DefaultCooldown": 300,
      "MinSize": 1,
      "Instances": [
        {
          "InstanceId": "i-d95eb0d4",
          "AvailabilityZone": "us-west-2b",
          "HealthStatus": "Healthy",
          "LifecycleState": "InService",
          "LaunchConfigurationName": "my-lc"
        },
        {
          "InstanceId": "i-13d7dc1f",
          "AvailabilityZone": "us-west-2a",
          "HealthStatus": "Healthy",
          "LifecycleState": "InService",
          "LaunchConfigurationName": "my-lc"
        }
      ],
      "MaxSize": 5,
      "VPCZoneIdentifier": null,
      "TerminationPolicies": [
        "Default"
      ],
      "LaunchConfigurationName": "my-lc",
      "CreatedTime": "2015-03-01T16:12:35.608Z",
      "AvailabilityZones": [
        "us-west-2b",
        "us-west-2a"
      ],
      "HealthCheckType": "EC2"
    }
  ]
}
```

```
}  
}
```

Delete the Notification Configuration

You can delete your Amazon EC2 Auto Scaling notification configuration at any time.

To delete Amazon EC2 Auto Scaling notification configuration (console)

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. On the navigation pane, under **Auto Scaling**, choose **Auto Scaling Groups**.
3. Select your Auto Scaling group.
4. On the **Notifications** tab, choose **Delete** next to the notification.

To delete Amazon EC2 Auto Scaling notification configuration (AWS CLI)

Use the following **delete-notification-configuration** command.

```
aws autoscaling delete-notification-configuration --auto-scaling-group-name my-asg --topic-arn arn
```

For information about deleting the Amazon SNS topic and all subscriptions associated with your Auto Scaling group, see [Clean Up](#) in the *Amazon Simple Notification Service Developer Guide*.

Logging Amazon EC2 Auto Scaling API Calls with AWS CloudTrail

Amazon EC2 Auto Scaling is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service using Amazon EC2 Auto Scaling. CloudTrail captures all API calls for Amazon EC2 Auto Scaling as events. The calls captured include calls from the Amazon EC2 Auto Scaling console and code calls to the Amazon EC2 Auto Scaling API.

If you create a trail, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Amazon EC2 Auto Scaling. If you don't configure a trail, you can still view the most recent events in the CloudTrail console in **Event history**. Using the information collected by CloudTrail, you can determine the request that was made to Amazon EC2 Auto Scaling, the IP address from which the request was made, who made the request, when it was made, and additional details.

To learn more about CloudTrail, see the [AWS CloudTrail User Guide](#).

Amazon EC2 Auto Scaling Information in CloudTrail

CloudTrail is enabled on your AWS account when you create the account. When activity occurs in Amazon EC2 Auto Scaling, that activity is recorded in a CloudTrail event along with other AWS service events in **Event history**. You can view, search, and download recent events in your AWS account. For more information, see [Viewing Events with CloudTrail Event History](#).

For an ongoing record of events in your AWS account, including events for Amazon EC2 Auto Scaling, create a trail. A *trail* enables CloudTrail to deliver log files to an Amazon S3 bucket. By default, when you create a trail in the console, the trail applies to all AWS Regions. The trail logs events from all Regions in the AWS partition and delivers the log files to the Amazon S3 bucket that you specify. Additionally, you

can configure other AWS services to further analyze and act upon the event data collected in CloudTrail logs. For more information, see the following:

- [Overview for Creating a Trail](#)
- [CloudTrail Supported Services and Integrations](#)
- [Configuring Amazon SNS Notifications for CloudTrail](#)
- [Receiving CloudTrail Log Files from Multiple Regions](#) and [Receiving CloudTrail Log Files from Multiple Accounts](#)

All Amazon EC2 Auto Scaling actions are logged by CloudTrail and are documented in the [Amazon EC2 Auto Scaling API Reference](#). For example, calls to the **CreateLaunchConfiguration**, **DescribeAutoScalingGroup**, and **UpdateAutoScalingGroup** actions generate entries in the CloudTrail log files.

Every event or log entry contains information about who generated the request. The identity information helps you determine the following:

- Whether the request was made with root or AWS Identity and Access Management (IAM) user credentials.
- Whether the request was made with temporary security credentials for a role or federated user.
- Whether the request was made by another AWS service.

For more information, see the [CloudTrail `userIdentity` Element](#).

Understanding Amazon EC2 Auto Scaling Log File Entries

A trail is a configuration that enables delivery of events as log files to an Amazon S3 bucket that you specify. CloudTrail log files contain one or more log entries. An event represents a single request from any source and includes information about the requested action, the date and time of the action, request parameters, and so on. CloudTrail log files aren't an ordered stack trace of the public API calls, so they don't appear in any specific order.

The following example shows a CloudTrail log entry that demonstrates the **CreateLaunchConfiguration** action.

```
{
  "eventVersion": "1.05",
  "userIdentity": {
    "type": "Root",
    "principalId": "123456789012",
    "arn": "arn:aws:iam::123456789012:root",
    "accountId": "123456789012",
    "accessKeyId": "AKIAIOSFODNN7EXAMPLE",
    "sessionContext": {
      "attributes": {
        "mfaAuthenticated": "false",
        "creationDate": "2018-08-21T17:05:42Z"
      }
    }
  },
  "eventTime": "2018-08-21T17:07:49Z",
  "eventSource": "autoscaling.amazonaws.com",
  "eventName": "CreateLaunchConfiguration",
  "awsRegion": "us-west-2",
  "sourceIPAddress": "192.0.2.0",
```

```
"userAgent": "Coral/Jakarta",
"requestParameters": {
  "ebsOptimized": false,
  "instanceMonitoring": {
    "enabled": false
  },
  "instanceType": "t2.micro",
  "keyName": "EC2-key-pair-oregon",
  "blockDeviceMappings": [
    {
      "deviceName": "/dev/xvda",
      "ebs": {
        "deleteOnTermination": true,
        "volumeSize": 8,
        "snapshotId": "snap-01676e0a2c3c7de9e",
        "volumeType": "gp2"
      }
    }
  ],
  "launchConfigurationName": "launch_configuration_1",
  "imageId": "ami-6cd6f714d79675a5",
  "securityGroups": [
    "sg-00c429965fd921483"
  ]
},
"responseElements": null,
"requestID": "0737e2ea-fb2d-11e3-bfd8-99133058e7bb",
"eventID": "3fcfb182-98f8-4744-bd45-b38835ab61cb",
"eventType": "AwsApiCall",
"recipientAccountId": "123456789012"
}
```

Controlling Access to Your Amazon EC2 Auto Scaling Resources

Access to Amazon EC2 Auto Scaling requires credentials that AWS can use to authenticate your requests. Those credentials must have [permissions](#) to perform Amazon EC2 Auto Scaling actions.

This topic provides details on how an IAM administrator can use AWS Identity and Access Management (IAM) to help secure your resources, by controlling who can perform Amazon EC2 Auto Scaling actions.

By default, a brand new IAM user has no permissions to do anything. To grant permissions to call Amazon EC2 Auto Scaling actions, you attach an IAM policy to the IAM users or groups that require the permissions it grants.

For the policy that grants a user full access to Amazon EC2 Auto Scaling, see [Predefined AWS Managed Policies \(p. 156\)](#). You can also create your own policies to specify allowed or denied actions and resources, as well as the conditions under which actions are allowed or denied. For examples, see [Customer Managed Policy Examples \(p. 157\)](#). However, we recommend that you first review the following introductory topics that explain the basic concepts and options for managing access to your Amazon EC2 Auto Scaling resources.

Specifying Actions in a Policy

You can specify any and all Amazon EC2 Auto Scaling actions in an IAM policy. For more information, see [Actions](#) in the *Amazon EC2 Auto Scaling API Reference*.

To specify a single policy, you can use the following prefix with the name of the action: `autoscaling:`.

```
"Action": "autoscaling:CreateAutoScalingGroup"
```

To specify multiple actions in a single policy, enclose them in square brackets and separate them with commas.

```
"Action": [  
  "autoscaling:CreateAutoScalingGroup",  
  "autoscaling:UpdateAutoScalingGroup"  
]
```

Wildcards are supported. For example, you can use `autoscaling:*` to specify all Amazon EC2 Auto Scaling actions.

```
"Action": "autoscaling:*"
```

You can also use `Describe*` to specify all actions whose names start with `Describe`.


```
"Action": "autoscaling:Describe"
```

Additional IAM Permissions

Users must have additional permissions from Amazon EC2 and IAM to perform certain actions. You specify the following actions in the Action element of an IAM policy statement.

Create an Auto Scaling group using a launch configuration

- `autoscaling:CreateAutoScalingGroup`
- `iam:CreateServiceLinkedRole`

Create an Auto Scaling group using a launch template

- `autoscaling:CreateAutoScalingGroup`
- `iam:CreateServiceLinkedRole`
- `ec2:RunInstances`

Update an Auto Scaling group that uses a launch template

- `autoscaling:UpdateAutoScalingGroup`
- `ec2:RunInstances`

Create a launch configuration

- `autoscaling:CreateLaunchConfiguration`
- `ec2:DescribeImages`
- `ec2:DescribeInstances`
- `ec2:DescribeInstanceAttribute`
- `ec2:DescribeKeyPairs`
- `ec2:DescribeSecurityGroups`
- `ec2:DescribeSpotInstanceRequests`
- `ec2:DescribeVpcClassicLink`

Users may require additional permissions to create or use Amazon EC2 resources, for example, to work with volumes and manage security groups from the console. For more information, see [Example Policies for Working in the Amazon EC2 Console](#) in the *Amazon EC2 User Guide for Linux Instances*. For information about permissions for creating and updating launch templates, see [Controlling the Use of Launch Templates](#) in the *Amazon EC2 User Guide for Linux Instances*.

There are also additional API actions for CloudWatch, Elastic Load Balancing, IAM, and Amazon SNS that may be required. For example, the `iam:PassRole` action is required to use an [instance profile](#) (p. 169).

Specifying the Resource

Access to resources can be controlled with an IAM policy. For actions that support resource-level permissions, you use an Amazon Resource Name (ARN) to identify the resource the policy applies to.

To specify an Auto Scaling group, you must specify its ARN as follows:

```
"Resource":  
  "arn:aws:autoscaling:region:123456789012:autoScalingGroup:uuid:autoScalingGroupName/asg-  
name"
```

To specify a launch configuration, you must specify its ARN as follows:

```
"Resource":  
  "arn:aws:autoscaling:region:123456789012:launchConfiguration:uuid:launchConfigurationName/  
lc-name"
```

To specify an Auto Scaling group with the `CreateAutoScalingGroup` action, you must replace the UUID with `*` as follows:

```
"Resource":  
  "arn:aws:autoscaling:region:123456789012:autoScalingGroup:*:autoScalingGroupName/asg-name"
```

To specify a launch configuration with the `CreateLaunchConfiguration` action, you must replace the UUID with `*` as follows:

```
"Resource":  
  "arn:aws:autoscaling:region:123456789012:launchConfiguration:*:launchConfigurationName/lc-  
name"
```

Alternatively, if you do not want to target a specific resource, you can use the `*` wildcard as the resource.

```
"Resource": "*" 
```

All Amazon EC2 Auto Scaling actions can be used in an IAM policy to either grant or deny users permission to use that action. However, not all Amazon EC2 Auto Scaling actions support resource-level permissions, which enable you to specify the resources on which an action can be performed.

For actions that don't support resource-level permissions, you must use `"*"` as the resource.

The following Amazon EC2 Auto Scaling actions do not support resource-level permissions:

- `DescribeAccountLimits`
- `DescribeAdjustmentTypes`
- `DescribeAutoScalingGroups`
- `DescribeAutoScalingInstances`
- `DescribeAutoScalingNotificationTypes`
- `DescribeLaunchConfigurations`
- `DescribeLifecycleHooks`
- `DescribeLifecycleHookTypes`
- `DescribeLoadBalancers`
- `DescribeLoadBalancerTargetGroups`
- `DescribeMetricCollectionTypes`
- `DescribeNotificationConfigurations`
- `DescribePolicies`

- `DescribeScalingActivities`
- `DescribeScalingProcessTypes`
- `DescribeScheduledActions`
- `DescribeTags`
- `DescribeTerminationPolicyTypes`

Specifying Conditions in a Policy

For actions that support resource-level permissions, you can control when users are allowed to use those actions based on conditions that have to be fulfilled.

When you grant permissions, you can use IAM policy language and predefined condition keys to specify the conditions.

The following condition keys are specific to Amazon EC2 Auto Scaling:

- `autoscaling:ImageId`
- `autoscaling:InstanceType`
- `autoscaling:InstanceTypes`
- `autoscaling:LaunchConfigurationName`
- `autoscaling:LaunchTemplateVersionSpecified`
- `autoscaling:LoadBalancerNames`
- `autoscaling:MaxSize`
- `autoscaling:MinSize`
- `autoscaling:ResourceTag/key`
- `autoscaling:SpotPrice`
- `autoscaling:TargetGroupARNs`
- `autoscaling:VPCZoneIdentifiers`

Important

For a complete list of constrainable API actions, the supported condition keys for each action, and the AWS-wide condition keys, see [Actions, Resources, and Condition Keys for Amazon EC2 Auto Scaling](#) and [AWS Global Condition Context Keys](#) in the *IAM User Guide*.

Predefined AWS Managed Policies

The managed policies created by AWS grant the required permissions for common use cases. You can attach these policies to your IAM users, based on the access that they need. Each policy grants access to all or some of the API actions for Amazon EC2 Auto Scaling.

The following are the AWS managed policies for Amazon EC2 Auto Scaling:

- `AutoScalingConsoleFullAccess` — Grants full access to Amazon EC2 Auto Scaling using the AWS Management Console.
- `AutoScalingConsoleReadOnlyAccess` — Grants read-only access to Amazon EC2 Auto Scaling using the AWS Management Console.

- `AutoScalingFullAccess` — Grants full access to Amazon EC2 Auto Scaling.
- `AutoScalingReadOnlyAccess` — Grants read-only access to Amazon EC2 Auto Scaling.

You can also use the `AmazonEC2FullAccess` policy to grant full access to all Amazon EC2 resources and related services.

Customer Managed Policy Examples

You can create your own custom IAM policies to allow or deny permissions for IAM users or groups to perform Amazon EC2 Auto Scaling actions. You can attach these custom policies to the IAM users or groups that require the specified permissions. The following examples show permissions for several common use cases.

Example: Restricting Which Service-Linked Role Can Be Passed (Using `PassRole`)

If your users require the ability to pass custom suffix service-linked roles to an Auto Scaling group, attach a policy to the users or roles, based on the access that they need. We recommend that you restrict this policy to only the service-linked roles that your users must access. For more information about custom suffix service-linked roles, see [Service-Linked Roles for Amazon EC2 Auto Scaling \(p. 163\)](#).

The following example is helpful for facilitating the security of your customer managed CMKs if you give different service-linked roles access to different keys. Depending on your needs, you might have a CMK for the development team, another for the QA team, and another for the finance team. First, create a service-linked role that has access to the required CMK, for example, a service-linked role named `AWSServiceRoleForAutoScaling_devteamkeyaccess`. Then, to grant permissions to pass that service-linked role to an Auto Scaling group, attach the policy to your IAM users as shown.

The policy in the following example grants users permissions to pass the `AWSServiceRoleForAutoScaling_devteamkeyaccess` role to create any Auto Scaling group whose name begins with `devteam`. If they try to specify a different service-linked role, they receive an error. If they choose not to specify a service-linked role, the default `AWSServiceRoleForAutoScaling` role is used instead.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "iam:PassRole",
      "Resource": "arn:aws:iam::123456789012:role/aws-service-role/autoscaling.amazonaws.com/AWSServiceRoleForAutoScaling_devteamkeyaccess",
      "Condition": {
        "StringEquals": {
          "iam:PassedToService": [
            "autoscaling.amazonaws.com"
          ]
        },
        "StringLike": {
          "iam:AssociatedResourceARN": [
            "arn:aws:autoscaling:region:123456789012:autoScalingGroup:*:autoScalingGroupName/devteam*"
          ]
        }
      }
    }
  ]
}
```

```
}  
]  
}
```

Example: Require a Launch Template

The following policy grants IAM users permissions to create and update Auto Scaling groups, provided that they use a launch template and specify the version of the launch template that the group uses to launch instances. Each instance uses the user-specified launch template version during launch. Users may access the Amazon EC2 resources specified in the launch template.

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "autoscaling:CreateAutoScalingGroup",  
        "autoscaling:UpdateAutoScalingGroup"  
      ],  
      "Resource": "*",  
      "Condition": {  
        "Bool": {  
          "autoscaling:LaunchTemplateVersionSpecified": "true"  
        }  
      }  
    },  
    {  
      "Effect": "Allow",  
      "Action": [  
        "ec2:*"  
      ],  
      "Resource": "*"   
    }  
  ]  
}
```

The `autoscaling:LaunchTemplateVersionSpecified` condition key accepts the following values:

- `true` - Ensures that a launch template version is specified.
- `false` - Ensures that either the `Latest` or `Default` launch template version is specified.
- `null` - Ensures that a launch template is not specified.

The `ec2:*` grants permission to call all Amazon EC2 API actions and access all Amazon EC2 resources.

Example: Create and Manage Launch Configurations

The following policy grants users permissions to use all Amazon EC2 Auto Scaling actions that include the string `LaunchConfiguration` in their names. Alternatively, you can list each action explicitly instead of using wildcards. However, the policy does not automatically apply to any new Amazon EC2 Auto Scaling actions with `LaunchConfiguration` in their names.

```
{  
  "Version": "2012-10-17",  
  "Statement": [{  
    "Effect": "Allow",  
    "Action": "autoscaling:*LaunchConfiguration*",  
  }]
```

```
    "Resource": "*"
  }]
}
```

The following policy grants users permissions to create a launch configuration if the instance type is `t2.micro` and the name of the launch configuration starts with `t2micro-`. They can specify a launch configuration for an Auto Scaling group only if its name starts with `t2micro-`.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "autoscaling:CreateLaunchConfiguration",
      "Resource": [
        "arn:aws:autoscaling:region:123456789012:launchConfiguration:*:launchConfigurationName/t2micro-*"
      ],
      "Condition": {
        "StringEquals": { "autoscaling:InstanceType": "t2.micro" }
      }
    },
    {
      "Effect": "Allow",
      "Action": [
        "autoscaling:CreateAutoScalingGroup",
        "autoscaling:UpdateAutoScalingGroup"
      ],
      "Resource": "*",
      "Condition": {
        "StringLikeIfExists": { "autoscaling:LaunchConfigurationName": "t2micro-*" }
      }
    }
  ]
}
```

Example: Create and Manage Auto Scaling Groups and Scaling Policies

The following policy grants users permissions to use all Amazon EC2 Auto Scaling actions that include the string `Scaling` in their names.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": ["autoscaling:*Scaling*"],
    "Resource": "*"
  }]
}
```

The following policy grants users permissions to use all Amazon EC2 Auto Scaling actions that include the string `Scaling` in their names, as long as the Auto Scaling group has the tag `purpose=webserver`. Because the `Describe` actions do not support resource-level permissions, you must specify them in a separate statement without conditions.

```
{
  "Version": "2012-10-17",
```

```
"Statement": [
{
  "Effect": "Allow",
  "Action": ["autoscaling:*Scaling*"],
  "Resource": "*",
  "Condition": {
    "StringEquals": { "autoscaling:ResourceTag/purpose": "webserver" }
  }
},
{
  "Effect": "Allow",
  "Action": "autoscaling:Describe*Scaling*",
  "Resource": "*"
}]
}
```

The following policy grants users permissions to use all Amazon EC2 Auto Scaling actions that include the string `Scaling` in their names, as long as they don't specify a minimum size less than 1 or a maximum size greater than 10. Because the `Describe` actions do not support resource-level permissions, you must specify them in a separate statement without conditions.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": ["autoscaling:*Scaling*"],
      "Resource": "*",
      "Condition": {
        "NumericGreaterThanEqualsIfExists": { "autoscaling:MinSize": 1 },
        "NumericLessThanEqualsIfExists": { "autoscaling:MaxSize": 10 }
      }
    },
    {
      "Effect": "Allow",
      "Action": "autoscaling:Describe*Scaling*",
      "Resource": "*"
    }
  ]
}
```

Example: Control Access Using Tags

To grant users permissions to create or tag an Auto Scaling group only if they specify specific tags, use the `aws:RequestTag` condition key. To allow only specific tag keys, use the `aws:TagKeys` condition key with the `ForAnyValue` modifier.

The following policy requires users to tag any Auto Scaling groups with the tags `purpose=webserver` and `cost-center=cc123`, and allows only the `purpose` and `cost-center` tags (no other tags can be specified).

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": [
      "autoscaling:CreateAutoScalingGroup",
      "autoscaling:CreateOrUpdateTags"
    ],
    "Resource": "*",
    "Condition": {
      "StringEquals": {
```

```
        "aws:RequestTag/purpose": "webserver",
        "aws:RequestTag/cost-center": "cc123"
    },
    "ForAllValues:StringEquals": { "aws:TagKeys": ["purpose", "cost-center"] }
}
}]
}
```

The following policy requires users to specify a tag with the key environment in the request.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": [
      "autoscaling:CreateAutoScalingGroup",
      "autoscaling:CreateOrUpdateTags"
    ],
    "Resource": "*",
    "Condition": {
      "StringLike": { "aws:RequestTag/environment": "*" }
    }
  }]
}
```

The following policy requires users to specify at least one tag in the request, and allows only the cost-center and owner keys.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": [
      "autoscaling:CreateAutoScalingGroup",
      "autoscaling:CreateOrUpdateTags"
    ],
    "Resource": "*",
    "Condition": {
      "ForAnyValue:StringEquals": { "aws:TagKeys": ["cost-center", "owner"] }
    }
  }]
}
```

The following policy grants users access to Auto Scaling groups with the tag allowed=true and allows them to apply only the tag environment=test. Because launch configurations do not support tags, and Describe actions do not support resource-level permissions, you must specify them in a separate statement without conditions.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": "autoscaling:*Scaling*",
    "Resource": "*",
    "Condition": {
      "StringEquals": { "autoscaling:ResourceTag/allowed": "true" },
      "StringEqualsIfExists": { "aws:RequestTag/environment": "test" },
      "ForAllValues:StringEquals": { "aws:TagKeys": "environment" }
    }
  },
  {

```



```
    "Effect": "Allow",
    "Action": [
        "autoscaling:*LaunchConfiguration*",
        "autoscaling:Describe*"
    ],
    "Resource": "*"
  }
}
```

Example: Change the Capacity of Auto Scaling Groups

The following policy grants users permissions to use the `SetDesiredCapacity` action to change the capacity of any Auto Scaling group.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": "autoscaling:SetDesiredCapacity",
    "Resource": "*"
  }]
}
```

The following policy grants users permissions to use the `SetDesiredCapacity` action to change the capacity of the specified Auto Scaling groups. Including the UUID ensures that access is granted to the specific Auto Scaling group. The UUID for a new group is different than the UUID for a deleted group with the same name.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": "autoscaling:SetDesiredCapacity",
    "Resource": [
      "arn:aws:autoscaling:region:123456789012:autoScalingGroup:7fe02b8e-7442-4c9e-8c8e-85fa99e9b5d9:autoScalingGroup-1",
      "arn:aws:autoscaling:region:123456789012:autoScalingGroup:9d8e8ea4-22e1-44c7-a14d-520f8518c2b9:autoScalingGroupName/group-2",
      "arn:aws:autoscaling:region:123456789012:autoScalingGroup:60d6b363-ae8b-467c-947f-f1d308935521:autoScalingGroupName/group-3"
    ]
  }]
}
```

The following policy grants users permissions to use the `SetDesiredCapacity` action to change the capacity of any Auto Scaling group whose name begins with `group-`.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": "autoscaling:SetDesiredCapacity",
    "Resource": [
      "arn:aws:autoscaling:region:123456789012:autoScalingGroup:*:autoScalingGroupName/group-*"
    ]
  }]
}
```

}

Service-Linked Roles for Amazon EC2 Auto Scaling

Amazon EC2 Auto Scaling uses service-linked roles for the permissions that it requires to call other AWS services on your behalf. A service-linked role is a unique type of IAM role that is linked directly to an AWS service.

Service-linked roles provide a secure way to delegate permissions to AWS services because only the linked service can assume a service-linked role. For more information, see [Using Service-Linked Roles](#) in the *IAM User Guide*. Service-linked roles also enable all API calls to be visible through AWS CloudTrail. This helps with monitoring and auditing requirements because you can track all actions that Amazon EC2 Auto Scaling performs on your behalf. For more information, see [Logging Amazon EC2 Auto Scaling API Calls with AWS CloudTrail](#) (p. 150).

The following sections describe how to create and manage Amazon EC2 Auto Scaling service-linked roles. Start by configuring permissions to allow an IAM entity (such as a user, group, or role) to create, edit, or delete a service-linked role. For more information, see [Using Service-Linked Roles](#) in the *IAM User Guide*.

Overview

There are two types of Amazon EC2 Auto Scaling service-linked roles:

- The default service-linked role for your account, named **AWSServiceRoleForAutoScaling**. This role is automatically assigned to your Auto Scaling groups unless you specify a different service-linked role.
- A service-linked role with a custom suffix that you specify when you create the role, for example, **AWSServiceRoleForAutoScaling_mysuffix**.

The permissions of a custom suffix service-linked role are identical to those of the default service-linked role. In both cases, you cannot edit the roles, and you also cannot delete them if they are still in use by an Auto Scaling group. The only difference is the role name suffix.

You can specify either role when you edit your AWS Key Management Service key policies to allow instances that are launched by Amazon EC2 Auto Scaling to be encrypted with your customer managed CMK. However, if you plan to give granular access to a specific customer managed CMK, you should use a custom suffix service-linked role. Using a custom suffix service-linked role provides you with:

- More control over the CMK.
- The ability to track which Auto Scaling group made an API call in your CloudTrail logs.

If you create customer managed CMKs that not all users should have access to, follow these steps to allow the use of a custom suffix service-linked role:

1. Create a service-linked role with a custom suffix. For more information, see [Create a Service-Linked Role \(Manual\)](#) (p. 164).
2. Give the service-linked role access to a customer managed CMK. For more information about the key policy that allows the CMK to be used by a service-linked role, see [Required CMK Key Policy for Use with Encrypted Volumes](#) (p. 166).
3. Give IAM users or roles access to the specified service-linked role. For more information about creating the IAM policy, see [Example: Restricting Which Service-Linked Role Can Be Passed \(Using PassRole\)](#) (p. 157).

Permissions Granted by the Service-Linked Role

Amazon EC2 Auto Scaling uses the **AWSServiceRoleForAutoScaling** service-linked role or your custom suffix service-linked role to make the following actions on the specified resources on your behalf:

- `ec2:AttachClassicLinkVpc`
- `ec2:CancelSpotInstanceRequests`
- `ec2:CreateFleet`
- `ec2:CreateTags`
- `ec2>DeleteTags`
- `ec2:Describe*`
- `ec2:DetachClassicLinkVpc`
- `ec2:ModifyInstanceAttribute`
- `ec2:RequestSpotInstances`
- `ec2:RunInstances`
- `ec2:TerminateInstances`
- `elasticloadbalancing:Register*`
- `elasticloadbalancing:Deregister*`
- `elasticloadbalancing:Describe*`
- `cloudwatch>DeleteAlarms`
- `cloudwatch:DescribeAlarms`
- `cloudwatch:PutMetricAlarm`
- `sns:Publish`

The role trusts the `autoscaling.amazonaws.com` service to assume it.

Create a Service-Linked Role (Automatic)

Amazon EC2 Auto Scaling creates the **AWSServiceRoleForAutoScaling** service-linked role for you the first time that you create an Auto Scaling group, unless you manually create a custom suffix service-linked role and specify it when creating the group.

Important

You must have IAM permissions to create the service-linked role. Otherwise, the automatic creation fails. For more information, see [Service-Linked Role Permissions](#) in the *IAM User Guide* or the information about [required user permissions \(p. 154\)](#) in this guide.

Amazon EC2 Auto Scaling began supporting service-linked roles in March 2018. If you created an Auto Scaling group before then, Amazon EC2 Auto Scaling created the **AWSServiceRoleForAutoScaling** role in your AWS account. For more information, see [A New Role Appeared in My AWS Account](#) in the *IAM User Guide*.

Create a Service-Linked Role (Manual)

To create a service-linked role (console)

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**, **Create role**.

3. For **Select type of trusted entity**, choose **AWS service**.
4. For **Choose the service that will use this role**, choose **EC2 Auto Scaling** and the **EC2 Auto Scaling** use case.
5. Choose **Next: Permissions**, **Next: Tags**, and then **Next: Review**. Note: You cannot attach tags to service-linked roles during creation.
6. On the **Review** page, leave **Role name** blank to create a service-linked role with the name `AWSServiceRoleForAutoScaling`, or enter a suffix to create a service-linked role with the name `AWSServiceRoleForAutoScaling_`*suffix*.
7. (Optional) For **Role description**, edit the description for the service-linked role.
8. Choose **Create role**.

To create a service-linked role (AWS CLI)

Use the following [create-service-linked-role](#) CLI command to create a service-linked role for Amazon EC2 Auto Scaling with the name `AWSServiceRoleForAutoScaling_`*suffix*.

```
aws iam create-service-linked-role --aws-service-name autoscaling.amazonaws.com --custom-suffix suffix
```

The output of this command includes the ARN of the service-linked role, which you can use to give the service-linked role access to your CMK.

```
{
  "Role": {
    "RoleId": "ABCDEF0123456789ABCDEF",
    "CreateDate": "2018-08-30T21:59:18Z",
    "RoleName": "AWSServiceRoleForAutoScaling_
```

suffix",
 "Arn": "arn:aws:iam::123456789012:role/aws-service-role/autoscaling.amazonaws.com/AWSServiceRoleForAutoScaling_*suffix*",
 "Path": "/aws-service-role/autoscaling.amazonaws.com/",
 "AssumeRolePolicyDocument": {
 "Version": "2012-10-17",
 "Statement": [
 {
 "Action": [
 "sts:AssumeRole"
],
 "Principal": {
 "Service": [
 "autoscaling.amazonaws.com"
]
 },
 "Effect": "Allow"
 }
]
 }
 }
}

For more information, see [Creating a Service-Linked Role](#) in the *IAM User Guide*.

Edit the Service-Linked Role

You cannot edit the service-linked roles that are created for Amazon EC2 Auto Scaling. After you create a service-linked role, you cannot change the name of the role or its permissions. However, you can edit the description of the role. For more information, see [Editing a Service-Linked Role](#) in the *IAM User Guide*.

Delete the Service-Linked Role

If you are not using an Auto Scaling group, we recommend that you delete its service-linked role. Deleting the role prevents you from having an entity that is not used or actively monitored and maintained.

You can delete a service-linked role only after first deleting the related AWS resources. This protects you from inadvertently revoking Amazon EC2 Auto Scaling permissions to your resources. If a service-linked role is used with multiple Auto Scaling groups, you must delete all Auto Scaling groups that use the service-linked role before you can delete it. For more information, see [Deleting Your Auto Scaling Infrastructure](#) (p. 70).

You can use IAM to delete a service-linked role. For more information, see [Deleting a Service-Linked Role](#) in the *IAM User Guide*.

If you delete the **AWSServiceRoleForAutoScaling** service-linked role, Amazon EC2 Auto Scaling creates the role again when you create an Auto Scaling group and do not specify a different service-linked role.

Supported Regions for Amazon EC2 Auto Scaling Service-Linked Roles

Amazon EC2 Auto Scaling supports using service-linked roles in all of the AWS Regions where the service is available. For more information, see [AWS Regions and Endpoints](#).

Required CMK Key Policy for Use with Encrypted Volumes

When creating an encrypted Amazon EBS snapshot or a launch template that specifies encrypted volumes, or enabling encryption by default, you can choose one of the following AWS Key Management Service customer master keys (CMK) to encrypt your data:

- **AWS managed CMK** — An encryption key in your account that Amazon EBS creates, owns, and manages. This is the default encryption key for a new account. The AWS managed CMK is used for encryption unless you specify a customer managed CMK.
- **Customer managed CMK** — A custom encryption key that you create, own, and manage. For more information, see [Creating Keys](#) in the *AWS Key Management Service Developer Guide*.

Amazon EC2 Auto Scaling does not need additional authorization to use the default AWS managed CMK to protect the encrypted volumes in your AWS account.

If you specify a customer managed CMK for Amazon EBS encryption, you (or your account administrator) must give the appropriate [service-linked role](#) (p. 163) access to the CMK, so that Amazon EC2 Auto Scaling can launch instances on your behalf. To do this, you must modify the CMK's key policy either when the CMK is created or at a later time.

Configuring Key Policies

Use the examples on this page to configure a key policy to give Amazon EC2 Auto Scaling access to your customer managed CMK. You must, at minimum, add two policy statements to your CMK's key policy for it to work with Amazon EC2 Auto Scaling.

- The first statement allows the IAM identity specified in the `Principal` element to use the CMK directly. It includes permissions to perform the AWS KMS `Encrypt`, `Decrypt`, `ReEncrypt*`, `GenerateDataKey*`, and `DescribeKey` operations on the CMK.
- The second statement allows the IAM identity specified in the `Principal` element to use grants to delegate a subset of its own permissions to AWS services that are integrated with AWS KMS or another principal. This allows them to use the CMK to create encrypted resources on your behalf.

When you add the new policy statements to your CMK policy, do not change any existing statements in the policy.

For each of the following examples, arguments that must be replaced, such as a key ID or the name of a service-linked role, are shown as *replaceable text in italics*. In most cases, you can replace the name of the service-linked role with the name of an Amazon EC2 Auto Scaling service-linked role. However, when using a launch configuration to launch Spot Instances, use the role named `AWSServiceRoleForEC2Spot`.

See the following resources:

- To create a CMK with the AWS CLI, see [create-key](#).
- To update a CMK policy with the AWS CLI, see [put-key-policy](#).
- To find a key ID and Amazon Resource Name (ARN), see [Finding the Key ID and ARN](#) in the *AWS Key Management Service Developer Guide*.
- For information about Amazon EC2 Auto Scaling service-linked roles, see [Service-Linked Roles for Amazon EC2 Auto Scaling](#) (p. 163).

Editing Key Policies in the Console

The examples in the following sections show only how to add statements to a key policy, which is just one way of changing a key policy. The easiest way to change a key policy is to use the IAM console's default view for key policies and make an IAM entity (user or role) one of the *key users* for the appropriate key policy. For more information, see [Using the AWS Management Console Default View](#) in the *AWS Key Management Service Developer Guide*.

Important

Be cautious. The console's default view policy statements include permissions to perform AWS KMS `Revoke` operations on the CMK. If you give an AWS account access to a CMK in your account, and you accidentally revoke the grant that gave them this permission, external users can no longer access their encrypted data or the key that was used to encrypt their data.

Example: CMK Key Policy Sections That Allow Access to the CMK

Add the following two policy statements to the key policy of the customer managed CMK, replacing the example ARN with the ARN of the appropriate service-linked role that is allowed access to the CMK. In this example, the policy sections give the service-linked role named `AWSServiceRoleForAutoScaling` permissions to use the customer managed CMK.

```
{
  "Sid": "Allow use of the CMK",
  "Effect": "Allow",
  "Principal": {
    "AWS": [
      "arn:aws:iam::123456789012:role/aws-service-role/autoscaling.amazonaws.com/AWSServiceRoleForAutoScaling"
    ]
  }
}
```

```
    },
    "Action": [
        "kms:Encrypt",
        "kms:Decrypt",
        "kms:ReEncrypt*",
        "kms:GenerateDataKey*",
        "kms:DescribeKey"
    ],
    "Resource": "*"
}
```

```
{
  "Sid": "Allow attachment of persistent resources",
  "Effect": "Allow",
  "Principal": {
    "AWS": [
      "arn:aws:iam::123456789012:role/aws-service-role/
autoscaling.amazonaws.com/AWSServiceRoleForAutoScaling"
    ]
  },
  "Action": [
    "kms:CreateGrant"
  ],
  "Resource": "*",
  "Condition": {
    "Bool": {
      "kms:GrantIsForAWSResource": true
    }
  }
}
```

Example: CMK Key Policy Sections That Allow Cross-Account Access to the CMK

If your customer managed CMK is in a different account than the Auto Scaling group, you must use a grant in combination with the key policy to allow access to the CMK. For more information, see [Using Grants](#) in the *AWS Key Management Service Developer Guide*.

First, add the following two policy statements to the CMK's key policy, replacing the example ARN with the ARN of the external account, and specifying the account in which the key can be used. The `GrantIsForAWSResource` condition is not included to allow an IAM user or role in the specified account to create the grant using the CLI command that follows.

```
{
  "Sid": "Allow use of the key in account 111122223333",
  "Effect": "Allow",
  "Principal": {
    "AWS": [
      "arn:aws:iam::111122223333:root"
    ]
  },
  "Action": [
    "kms:Encrypt",
    "kms:Decrypt",
    "kms:ReEncrypt*",
    "kms:GenerateDataKey*",
    "kms:DescribeKey"
  ],
  "Resource": "*"
}
```

```
{
  "Sid": "Allow attachment of persistent resources in account 111122223333",
  "Effect": "Allow",
  "Principal": {
    "AWS": [
      "arn:aws:iam::111122223333:root"
    ]
  },
  "Action": [
    "kms:CreateGrant"
  ],
  "Resource": "*"
}
```

Then, from the external account, create a grant that delegates the relevant permissions to the appropriate service-linked role. The `Grantee Principal` element of the grant is the ARN of the appropriate service-linked role. The `key-id` is the ARN of the CMK. The following is an example [create-a-grant](#) CLI command that gives the service-linked role named `AWSServiceRoleForAutoScaling` in account 111122223333 permissions to use the CMK in account 444455556666.

```
aws kms create-grant \
  --region us-west-2 \
  --key-id arn:aws:kms:us-west-2:444455556666:key/1a2b3c4d-5e6f-1a2b-3c4d-5e6f1a2b3c4d \
  --grantee-principal arn:aws:iam::111122223333:role/aws-service-role/
autoscaling.amazonaws.com/AWSServiceRoleForAutoScaling \
  --operations "Encrypt" "Decrypt" "ReEncryptFrom" "ReEncryptTo" "GenerateDataKey"
  "GenerateDataKeyWithoutPlaintext" "DescribeKey" "CreateGrant"
```

For this command to succeed, the user making the request must have permissions for the `CreateGrant` action. The following example IAM policy allows an IAM user or role in account 111122223333 to create a grant for the CMK in account 444455556666.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowCreationOfGrant",
      "Effect": "Allow",
      "Action": "kms:CreateGrant",
      "Resource": "arn:aws:kms:us-west-2:444455556666:key/1a2b3c4d-5e6f-1a2b-3c4d-5e6f1a2b3c4d"
    }
  ]
}
```

If you have any problems configuring the cross-account access to a customer managed CMK that is required to launch an instance with an encrypted volume, see the [troubleshooting section \(p. 177\)](#).

For more information, see [Amazon EBS Encryption](#) in the *Amazon EC2 User Guide for Linux Instances* and the [AWS Key Management Service Developer Guide](#).

IAM Role for Applications That Run on Amazon EC2 Instances

Applications that run on Amazon EC2 instances need credentials to access other AWS services. To provide these credentials in a secure way, use an IAM role. The role supplies temporary permissions that the

application can use when it accesses other AWS resources. The role's permissions determine what the application is allowed to do.

For more information, see [Using an IAM Role to Grant Permissions to Applications Running on Amazon EC2 Instances](#) in the *IAM User Guide*.

For instances in an Auto Scaling group, you must create a launch configuration or template and choose an instance profile to associate with the instances. An instance profile is a container for an IAM role that allows Amazon EC2 to pass the IAM role to an instance when the instance is launched. First, create an IAM role that has all of the permissions required to access the AWS resources. Then, create the instance profile and assign the role to it. For more information, see [Using Instance Profiles](#) in the *IAM User Guide*.

Note

When you use the IAM console to create a role for Amazon EC2, the console guides you through the steps for creating the role and automatically creates an instance profile with the same name as the IAM role.

Prerequisites

Create the IAM role that your application running on Amazon EC2 can assume. Choose the appropriate permissions, so that the application that is subsequently given the role can access the resources that it needs.

To create an IAM role

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**, **Create role**.
3. For **Select type of trusted entity**, choose **AWS service**.
4. For **Choose the service that will use this role**, choose **EC2** and the **EC2** use case. Choose **Next: Permissions**.
5. For **Attach permissions policies**, choose the AWS managed policies that contain the required permissions. Choose **Next: Tags** and then **Next: Review**.
6. On the **Review** page, enter a name for the role and choose **Create role**.

Create a Launch Configuration

When you create the launch configuration using the AWS Management Console, on the **Configure Details** page, select the role from **IAM role**. For more information, see [Creating a Launch Configuration](#) (p. 32).

When you create the launch configuration using the [create-launch-configuration](#) command from the AWS CLI, specify the name of the instance profile as follows:

```
aws autoscaling create-launch-configuration --launch-configuration-name my-lc-with-  
instance-profile \  
--image-id ami-01e24be29428c15b2 --instance-type t2.micro \  
--iam-instance-profile my-instance-profile
```

Create a Launch Template

When you create the launch template using the AWS Management Console, in the **Advanced Details** section, select the role from **IAM instance profile**. For more information, see [Creating a Launch Template for an Auto Scaling Group](#) (p. 24).

When you create the launch template using the [create-launch-template](#) command from the AWS CLI, specify the name of the instance profile as follows:

```
aws ec2 create-launch-template --launch-template-name my-lt-with-instance-profile --  
version-description version1 \  
--launch-template-data  
'{"ImageId": "ami-01e24be29428c15b2", "InstanceType": "t2.micro", "IamInstanceProfile":  
{"Name": "my-instance-profile"}}'
```

Troubleshooting Amazon EC2 Auto Scaling

Amazon EC2 Auto Scaling provides specific and descriptive errors to help you troubleshoot issues. You can find the error messages in the description of the scaling activities.

Contents

- [Retrieving an Error Message \(p. 172\)](#)
- [Troubleshooting Amazon EC2 Auto Scaling: EC2 Instance Launch Failures \(p. 174\)](#)
- [Troubleshooting Amazon EC2 Auto Scaling: AMI Issues \(p. 178\)](#)
- [Troubleshooting Amazon EC2 Auto Scaling: Load Balancer Issues \(p. 180\)](#)
- [Troubleshooting Auto Scaling: Capacity Limits \(p. 181\)](#)

Retrieving an Error Message

To retrieve an error message from the description of scaling activities, use the [describe-scaling-activities](#) command as follows:

```
aws autoscaling describe-scaling-activities --auto-scaling-group-name my-asg
```

The following is an example response, where `StatusCode` contains the current status of the activity and `StatusMessage` contains the error message:

```
{
  "Activities": [
    {
      "Description": "Launching a new EC2 instance: i-4ba0837f",
      "AutoScalingGroupName": "my-asg",
      "ActivityId": "f9f2d65b-f1f2-43e7-b46d-d86756459699",
      "Details": "{\"Availability Zone\":\"us-west-2c\"}",
      "StartTime": "2013-08-19T20:53:29.930Z",
      "Progress": 100,
      "EndTime": "2013-08-19T20:54:02Z",
      "Cause": "At 2013-08-19T20:53:25Z a user request created an
AutoScalingGroup...",
      "StatusCode": "Failed",
      "StatusMessage": "The image id 'ami-4edb0327' does not exist. Launching EC2
instance failed."
    }
  ]
}
```

The following tables list the types of error messages and provide links to the troubleshooting resources that you can use to troubleshoot issues.

EC2 Instance Launch Failures

Issue	Error Message
Auto Scaling group	AutoScalingGroup <Auto Scaling group name> not found. (p. 175)

Issue	Error Message
Availability Zone	The requested Availability Zone is no longer supported. Please retry your request... (p. 176)
AWS account	You are not subscribed to this service. Please see https://aws.amazon.com/ . (p. 176)
Block device mapping	Invalid device name upload. Launching EC2 instance failed. (p. 176)
Block device mapping	Value (<name associated with the instance storage device>) for parameter virtualName is invalid... (p. 177)
Block device mapping	EBS block device mappings not supported for instance-store AMIs. (p. 177)
Instance type and Availability Zone	Your requested instance type (<instance type>) is not supported in your requested Availability Zone (<instance Availability Zone>)... (p. 176)
Key pair	The key pair <key pair associated with your EC2 instance> does not exist. Launching EC2 instance failed. (p. 175)
Launch configuration	The requested configuration is currently not supported. (p. 175)
Placement group	Placement groups may not be used with instances of type 'm1.large'. Launching EC2 instance failed. (p. 177)
Security group	The security group <name of the security group> does not exist. Launching EC2 instance failed. (p. 175)
Service-linked role	Client.InternalError: Client error on launch. (p. 177)

AMI Issues

Issue	Error Message
AMI ID	The AMI ID <ID of your AMI> does not exist. Launching EC2 instance failed. (p. 179)
AMI ID	AMI <AMI ID> is pending, and cannot be run. Launching EC2 instance failed. (p. 179)
AMI ID	Value (<ami ID>) for parameter virtualName is invalid. (p. 179)
Architecture mismatch	The requested instance type's architecture (i386) does not match the architecture in the manifest for ami-6622f00f (x86_64). Launching ec2 instance failed. (p. 179)

Load Balancer Issues

Issue	Error Message
Cannot find load balancer	Cannot find Load Balancer <your launch environment>. Validating load balancer configuration failed. (p. 180)
Instances in VPC	EC2 instance <instance ID> is not in VPC. Updating load balancer configuration failed. (p. 181)

Issue	Error Message
No active load balancer	There is no ACTIVE Load Balancer named <load balancer name>. Updating load balancer configuration failed. (p. 180)
Security token	The security token included in the request is invalid. Validating load balancer configuration failed. (p. 181)

Capacity Limits

Issue	Error Message
Capacity limits	<number of instances> instance(s) are already running. Launching EC2 instance failed. (p. 182)
Insufficient capacity in Availability Zone	We currently do not have sufficient <instance type> capacity in the Availability Zone you requested (<requested Availability Zone>)... (p. 181)

Troubleshooting Amazon EC2 Auto Scaling: EC2 Instance Launch Failures

This page provides information about your EC2 instances that fail to launch, potential causes, and the steps you can take to resolve the issues.

To retrieve an error message, see [Retrieving an Error Message](#) (p. 172).

When your EC2 instances fail to launch, you might get one or more of the following error messages:

Error Messages

- The security group <name of the security group> does not exist. Launching EC2 instance failed. (p. 175)
- The key pair <key pair associated with your EC2 instance> does not exist. Launching EC2 instance failed. (p. 175)
- The requested configuration is currently not supported. (p. 175)
- AutoScalingGroup <Auto Scaling group name> not found. (p. 175)
- The requested Availability Zone is no longer supported. Please retry your request... (p. 176)
- Your requested instance type (<instance type>) is not supported in your requested Availability Zone (<instance Availability Zone>)... (p. 176)
- You are not subscribed to this service. Please see <https://aws.amazon.com/>. (p. 176)
- Invalid device name upload. Launching EC2 instance failed. (p. 176)
- Value (<name associated with the instance storage device>) for parameter virtualName is invalid... (p. 177)
- EBS block device mappings not supported for instance-store AMIs. (p. 177)
- Placement groups may not be used with instances of type 'm1.large'. Launching EC2 instance failed. (p. 177)
- Client.InternalError: Client error on launch. (p. 177)

The security group <name of the security group> does not exist. Launching EC2 instance failed.

- **Cause:** The security group specified in your launch configuration might have been deleted.
- **Solution:**
 1. Use the [describe-security-groups](#) command to get the list of the security groups associated with your account.
 2. From the list, select the security groups to use. To create a security group instead, use the [create-security-group](#) command.
 3. Create a new launch configuration.
 4. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

The key pair <key pair associated with your EC2 instance> does not exist. Launching EC2 instance failed.

- **Cause:** The key pair that was used when launching the instance might have been deleted.
- **Solution:**
 1. Use the [describe-key-pairs](#) command to get the list of the key pairs available to you.
 2. From the list, select the key pair to use. To create a key pair instead, use the [create-key-pair](#) command.
 3. Create a new launch configuration.
 4. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

The requested configuration is currently not supported.

- **Cause:** Some options in your launch configuration might not be currently supported.
- **Solution:**
 1. Create a new launch configuration.
 2. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

AutoScalingGroup <Auto Scaling group name> not found.

- **Cause:** The Auto Scaling group might have been deleted.
- **Solution:** Create a new Auto Scaling group.

The requested Availability Zone is no longer supported. Please retry your request...

- **Error Message:** The requested Availability Zone is no longer supported. Please retry your request by not specifying an Availability Zone or choosing <list of available Availability Zones>. Launching EC2 instance failed.
- **Cause:** The Availability Zone associated with your Auto Scaling group might not be currently available.
- **Solution:** Update your Auto Scaling group with the recommendations in the error message.

Your requested instance type (<instance type>) is not supported in your requested Availability Zone (<instance Availability Zone>)...

- **Error Message:** Your requested instance type (<instance type>) is not supported in your requested Availability Zone (<instance Availability Zone>). Please retry your request by not specifying an Availability Zone or choosing <list of Availability Zones that supports the instance type>. Launching EC2 instance failed.
- **Cause:** The instance type associated with your launch configuration might not be currently available in the Availability Zones specified in your Auto Scaling group.
- **Solution:** Update your Auto Scaling group with the recommendations in the error message.

You are not subscribed to this service. Please see <https://aws.amazon.com/>.

- **Cause:** Your AWS account might have expired.
- **Solution:** Go to <https://aws.amazon.com/> and choose **Sign Up Now** to open a new account.

Invalid device name upload. Launching EC2 instance failed.

- **Cause:** The block device mappings in your launch configuration might contain block device names that are not available or currently not supported.
- **Solution:**
 1. Use the [describe-volumes](#) command to see how the volumes are exposed to the instance.
 2. Create a new launch configuration using the device name listed in the volume description.
 3. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

Value (<name associated with the instance storage device>) for parameter virtualName is invalid...

- **Error Message:** Value (<name associated with the instance storage device>) for parameter virtualName is invalid. Expected format: 'ephemeralNUMBER'. Launching EC2 instance failed.
- **Cause:** The format specified for the virtual name associated with the block device is incorrect.
- **Solution:**
 1. Create a new launch configuration by specifying the device name in the `virtualName` parameter. For information about the device name format, see [Instance Store Device Names](#) in the *Amazon EC2 User Guide for Linux Instances*.
 2. Update your Auto Scaling group with the new launch configuration using the `update-auto-scaling-group` command.

EBS block device mappings not supported for instance-store AMIs.

- **Cause:** The block device mappings specified in the launch configuration are not supported on your instance.
- **Solution:**
 1. Create a new launch configuration with block device mappings supported by your instance type. For more information, see [Block Device Mapping](#) in the *Amazon EC2 User Guide for Linux Instances*.
 2. Update your Auto Scaling group with the new launch configuration using the `update-auto-scaling-group` command.

Placement groups may not be used with instances of type 'm1.large'. Launching EC2 instance failed.

- **Cause:** Your cluster placement group contains an invalid instance type.
- **Solution:**
 1. For information about valid instance types supported by the placement groups, see [Placement Groups](#) in the *Amazon EC2 User Guide for Linux Instances*.
 2. Follow the instructions detailed in the [Placement Groups](#) to create a new placement group.
 3. Alternatively, create a new launch configuration with the supported instance type.
 4. Update your Auto Scaling group with a new placement group or launch configuration using the `update-auto-scaling-group` command.

Client.InternalError: Client error on launch.

- **Cause:** This error can be caused when an Auto Scaling group attempts to launch an instance that has an encrypted EBS volume, but the service-linked role does not have access to the customer managed CMK used to encrypt it. For more information, see [Required CMK Key Policy for Use with Encrypted Volumes](#) (p. 166).
- **Solution:** Additional setup is required to allow the Auto Scaling group to launch instances. The following table summarizes the steps for resolving the error. For more information, see <https://forums.aws.amazon.com/thread.jspa?threadID=277523>.

Scenario	Next Steps
Scenario 1: CMK and Auto Scaling group are in the same AWS account	Allow the service-linked role to use the CMK as follows: <ol style="list-style-type: none"> 1. Determine which service-linked role to use for this Auto Scaling group. 2. Update the key policy on the CMK and allow the service-linked role to use the CMK. 3. Update the Auto Scaling group to use the service-linked role.
Scenario 2: CMK and Auto Scaling group are in different AWS accounts	There are two possible solutions: <p>Solution 1: Use a CMK in the same AWS account as the Auto Scaling group</p> <ol style="list-style-type: none"> 1. Copy and re-encrypt the snapshot with another CMK that belongs to the same account as the Auto Scaling group. 2. Allow the service-linked role to use the new CMK. See the steps for Scenario 1. <p>Solution 2: Continue to use the CMK in a different AWS account from the Auto Scaling group</p> <ol style="list-style-type: none"> 1. Determine which service-linked role to use for this Auto Scaling group. 2. Allow the Auto Scaling group account access to the CMK. 3. Define an IAM user or role in the Auto Scaling group account that can create a grant. 4. Create a grant to the CMK with the service-linked role as the grantee principal. 5. Update the Auto Scaling group to use the service-linked role.

Troubleshooting Amazon EC2 Auto Scaling: AMI Issues

This page provides information about the issues associated with your AMIs, potential causes, and the steps you can take to resolve the issues.

To retrieve an error message, see [Retrieving an Error Message \(p. 172\)](#).

When your EC2 instances fail to launch due to issues with your AMI, you might get one or more of the following error messages.

Error Messages

- [The AMI ID <ID of your AMI> does not exist. Launching EC2 instance failed. \(p. 179\)](#)
- [AMI <AMI ID> is pending, and cannot be run. Launching EC2 instance failed. \(p. 179\)](#)
- [Value \(<ami ID>\) for parameter virtualName is invalid. \(p. 179\)](#)
- [The requested instance type's architecture \(i386\) does not match the architecture in the manifest for ami-6622f00f \(x86_64\). Launching ec2 instance failed. \(p. 179\)](#)

The AMI ID <ID of your AMI> does not exist. Launching EC2 instance failed.

- **Cause:** The AMI might have been deleted after creating the launch configuration.
- **Solution:**
 1. Create a new launch configuration using a valid AMI.
 2. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

AMI <AMI ID> is pending, and cannot be run. Launching EC2 instance failed.

- **Cause:** You might have just created your AMI (by taking a snapshot of a running instance or any other way), and it might not be available yet.
- **Solution:** You must wait for your AMI to be available and then create your launch configuration.

Value (<ami ID>) for parameter virtualName is invalid.

- **Cause:** Incorrect value. The `virtualName` parameter refers to the virtual name associated with the device.
- **Solution:**
 1. Create a new launch configuration by specifying the name of the virtual device of your instance for the `virtualName` parameter.
 2. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

The requested instance type's architecture (i386) does not match the architecture in the manifest for ami-6622f00f (x86_64). Launching ec2 instance failed.

- **Cause:** The architecture of the `InstanceType` mentioned in your launch configuration does not match the image architecture.
- **Solution:**
 1. Create a new launch configuration using the AMI architecture that matches the architecture of the requested instance type.
 2. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

Troubleshooting Amazon EC2 Auto Scaling: Load Balancer Issues

This page provides information about issues caused by the load balancer associated with your Auto Scaling group, potential causes, and the steps you can take to resolve the issues.

To retrieve an error message, see [Retrieving an Error Message \(p. 172\)](#).

When your EC2 instances fail to launch due to issues with the load balancer associated with your Auto Scaling group, you might get one or more of the following error messages.

Error Messages

- [Cannot find Load Balancer <your launch environment>. Validating load balancer configuration failed. \(p. 180\)](#)
- [There is no ACTIVE Load Balancer named <load balancer name>. Updating load balancer configuration failed. \(p. 180\)](#)
- [EC2 instance <instance ID> is not in VPC. Updating load balancer configuration failed. \(p. 181\)](#)
- [EC2 instance <instance ID> is in VPC. Updating load balancer configuration failed. \(p. 181\)](#)
- [The security token included in the request is invalid. Validating load balancer configuration failed. \(p. 181\)](#)

Cannot find Load Balancer <your launch environment>. Validating load balancer configuration failed.

- **Cause 1:** The load balancer has been deleted.
- **Solution 1:**
 1. Check to see if your load balancer still exists. You can use the [describe-load-balancers](#) command.
 2. If you see your load balancer listed in the response, see **Cause 2**.
 3. If you do not see your load balancer listed in the response, you can either create a new load balancer and then create a new Auto Scaling group or you can create a new Auto Scaling group without the load balancer.
- **Cause 2:** The load balancer name was not specified in the right order when creating the Auto Scaling group.
- **Solution 2:** Create a new Auto Scaling group and specify the load balancer name at the end.

There is no ACTIVE Load Balancer named <load balancer name>. Updating load balancer configuration failed.

- **Cause:** The specified load balancer might have been deleted.
- **Solution:** You can either create a new load balancer and then create a new Auto Scaling group or create a new Auto Scaling group without the load balancer.

EC2 instance <instance ID> is not in VPC. Updating load balancer configuration failed.

- **Cause:** The specified instance does not exist in the VPC.
- **Solution:** You can either delete your load balancer associated with the instance or create a new Auto Scaling group.

EC2 instance <instance ID> is in VPC. Updating load balancer configuration failed.

- **Cause:** The load balancer is in EC2-Classic but the Auto Scaling group is in a VPC.
- **Solution:** Ensure that the load balancer and the Auto Scaling group are in the same network (EC2-Classic or a VPC).

The security token included in the request is invalid. Validating load balancer configuration failed.

- **Cause:** Your AWS account might have expired.
- **Solution:** Check whether your AWS account is valid. Go to <https://aws.amazon.com/> and choose **Sign Up Now** to open a new account.

Troubleshooting Auto Scaling: Capacity Limits

This page provides information about issues with the capacity limits of your Auto Scaling group, potential causes, and the steps you can take to resolve the issues. For more information about Amazon EC2 Auto Scaling limits, see [Amazon EC2 Auto Scaling Limits](#) (p. 9).

To retrieve an error message, see [Retrieving an Error Message](#) (p. 172).

If your EC2 instances fail to launch due to issues with the capacity limits of your Auto Scaling group, you might get one or more of the following error messages.

Error Messages

- [We currently do not have sufficient <instance type> capacity in the Availability Zone you requested \(<requested Availability Zone>\)...](#) (p. 181)
- [<number of instances> instance\(s\) are already running. Launching EC2 instance failed.](#) (p. 182)

We currently do not have sufficient <instance type> capacity in the Availability Zone you requested (<requested Availability Zone>)..

- **Error Message:** We currently do not have sufficient <instance type> capacity in the Availability Zone you requested (<requested Availability Zone>). Our system will be working on provisioning additional capacity. You can currently get <instance type> capacity by not specifying an Availability Zone in your

request or choosing <list of Availability Zones that currently supports the instance type>. Launching EC2 instance failed.

- **Cause:** At this time, Auto Scaling cannot support your instance type in your requested Availability Zone.
- **Solution:**
 1. Create a new launch configuration by following the recommendations in the error message.
 2. Update your Auto Scaling group with the new launch configuration using the [update-auto-scaling-group](#) command.

<number of instances> instance(s) are already running. Launching EC2 instance failed.

- **Cause:** The Auto Scaling group has reached the limit set by the `DesiredCapacity` parameter.
- **Solution:**
 - Update your Auto Scaling group by providing a new value for the `--desired-capacity` parameter using the [update-auto-scaling-group](#) command.
 - If you've reached your limit for the number of EC2 instances, you can request an increase. For more information, see [AWS Service Limits](#).

Auto Scaling Resources

The following related resources can help you as you work with this service.

- [Amazon EC2 Auto Scaling](#) – The primary web page for information about Amazon EC2 Auto Scaling.
- [Amazon EC2 Technical FAQ](#) – The answers to questions customers ask about Amazon EC2 and Amazon EC2 Auto Scaling.
- [Amazon EC2 Discussion Forum](#) – Get help from the community.
- [AWS Auto Scaling User Guide](#) – The AWS Auto Scaling console makes it easier for you to use the scaling features of multiple services. With AWS Auto Scaling, you can also simplify the process of defining dynamic scaling policies for your Auto Scaling groups and use predictive scaling to scale your Amazon EC2 capacity in advance of predicted traffic changes.

The following additional resources are available to help you learn more about AWS.

- [Classes & Workshops](#) – Links to role-based and specialty courses as well as self-paced labs to help sharpen your AWS skills and gain practical experience.
- [AWS Developer Tools](#) – Links to developer tools, SDKs, IDE toolkits, and command line tools for developing and managing AWS applications.
- [AWS Whitepapers](#) – Links to a comprehensive list of technical AWS whitepapers, covering topics such as architecture, security, and economics and authored by AWS Solutions Architects or other technical experts.
- [AWS Support Center](#) – The hub for creating and managing your AWS Support cases. Also includes links to other helpful resources, such as forums, technical FAQs, service health status, and AWS Trusted Advisor.
- [AWS Support](#) – The primary web page for information about AWS Support, a one-on-one, fast-response support channel to help you build and run applications in the cloud.
- [Contact Us](#) – A central contact point for inquiries concerning AWS billing, account, events, abuse, and other issues.
- [AWS Site Terms](#) – Detailed information about our copyright and trademark; your account, license, and site access; and other topics.

Document History

The following table describes important additions to the Amazon EC2 Auto Scaling documentation, beginning in July 2018. For notification about updates to this documentation, you can subscribe to the RSS feed.

update-history-change	update-history-description	update-history-date
Minimum number of instance types (p. 184)	You no longer have to specify additional instance types for groups of Spot, On-Demand, and Reserved Instances. For all Auto Scaling groups, the minimum is now one instance type. For more information, see Auto Scaling Groups with Multiple Instance Types and Purchase Options in the <i>Amazon EC2 Auto Scaling User Guide</i> .	September 16, 2019
Support for new Spot allocation strategy (p. 184)	Amazon EC2 Auto Scaling now supports a new Spot allocation strategy "capacity-optimized" that fulfills your request using Spot Instance pools that are optimally chosen based on the available Spot capacity. For more information, see Auto Scaling Groups with Multiple Instance Types and Purchase Options in the <i>Amazon EC2 Auto Scaling User Guide</i> .	August 12, 2019
Guide changes (p. 184)	Improved Amazon EC2 Auto Scaling documentation in the Service-Linked Roles and Required CMK Key Policy for Use with Encrypted Volumes topics.	August 1, 2019
Support for tagging enhancement (p. 184)	Amazon EC2 Auto Scaling now adds tags to Amazon EC2 instances as part of the same API call that launches the instances. For more information, see Tagging Auto Scaling Groups and Instances .	July 26, 2019
Guide changes (p. 184)	Improved Amazon EC2 Auto Scaling documentation in the Suspending and Resuming Scaling Processes topic. Updated Customer Managed Policy Examples to include an example policy that allows users to pass	June 13, 2019

	only specific custom suffix service-linked roles to Amazon EC2 Auto Scaling.	
Support for new Amazon EBS feature (p. 184)	Added support for new Amazon EBS feature in the launch template topic. Change the encryption state of an EBS volume while restoring from a snapshot. For more information, see Creating a Launch Template for an Auto Scaling Group in the <i>Amazon EC2 Auto Scaling User Guide</i> .	May 13, 2019
Guide changes (p. 184)	Improved Amazon EC2 Auto Scaling documentation in the following sections: Controlling Which Auto Scaling Instances Terminate During Scale In , Auto Scaling Groups , Auto Scaling Groups with Multiple Instance Types and Purchase Options , and Dynamic Scaling for Amazon EC2 Auto Scaling .	March 12, 2019
Support for combining instance types and purchase options (p. 184)	Provision and automatically scale instances across purchase options (Spot, On-Demand, and Reserved Instances) and instance types within a single Auto Scaling group. For more information, see Auto Scaling Groups with Multiple Instance Types and Purchase Options in the <i>Amazon EC2 Auto Scaling User Guide</i> .	November 13, 2018
Updated topic for scaling based on Amazon SQS (p. 184)	Updated guide to explain how you can use custom metrics to scale an Auto Scaling group in response to changing demand from an Amazon SQS queue. For more information, see Scaling Based on Amazon SQS in the <i>Amazon EC2 Auto Scaling User Guide</i> .	July 26, 2018

The following table describes important changes to the Amazon EC2 Auto Scaling documentation before July 2018.

Feature	Description	Release Date
Support for target tracking scaling policies	Set up dynamic scaling for your application in just a few steps. For more information, see Target Tracking Scaling Policies for Amazon EC2 Auto Scaling .	12 July 2017

Feature	Description	Release Date
Support for resource-level permissions	Create IAM policies to control access at the resource level. For more information, see Controlling Access to Your Amazon EC2 Auto Scaling Resources .	15 May 2017
Monitoring improvements	Auto Scaling group metrics no longer require that you enable detailed monitoring. You can now enable group metrics collection and view metrics graphs from the Monitoring tab in the console. For more information, see Monitoring Your Auto Scaling Groups and Instances Using Amazon CloudWatch .	18 August 2016
Support for Application Load Balancers	Attach one or more target groups to a new or existing Auto Scaling group. For more information, see Attaching a Load Balancer to Your Auto Scaling Group .	11 August 2016
Events for lifecycle hooks	Amazon EC2 Auto Scaling sends events to CloudWatch Events when it executes lifecycle hooks. For more information, see Getting CloudWatch Events When Your Auto Scaling Group Scales .	24 February 2016
Instance protection	Prevent Amazon EC2 Auto Scaling from selecting specific instances for termination when scaling in. For more information, see Instance Protection .	07 December 2015
Step scaling policies	Create a scaling policy that enables you to scale based on the size of the alarm breach. For more information, see Scaling Policy Types .	06 July 2015
Update load balancer	Attach a load balancer to or detach a load balancer from an existing Auto Scaling group. For more information, see Attaching a Load Balancer to Your Auto Scaling Group .	11 June 2015
Support for ClassicLink	Link EC2-Classic instances in your Auto Scaling group to a VPC, enabling communication between these linked EC2-Classic instances and instances in the VPC using private IP addresses. For more information, see Linking EC2-Classic Instances to a VPC .	19 January 2015
Lifecycle hooks	Hold your newly launched or terminating instances in a pending state while you perform actions on them. For more information, see Amazon EC2 Auto Scaling Lifecycle Hooks .	30 July 2014
Detach instances	Detach instances from an Auto Scaling group. For more information, see Detach EC2 Instances from Your Auto Scaling Group .	30 July 2014
Put instances into a Standby state	Put instances that are in an <code>InService</code> state into a <code>Standby</code> state. For more information, see Temporarily Removing Instances from Your Auto Scaling Group .	30 July 2014
Manage tags	Manage your Auto Scaling groups using the AWS Management Console. For more information, see Tagging Auto Scaling Groups and Instances .	01 May 2014
Support for Dedicated Instances	Launch Dedicated Instances by specifying a placement tenancy attribute when you create a launch configuration. For more information, see Instance Placement Tenancy .	23 April 2014

Feature	Description	Release Date
Create a group or launch configuration from an EC2 instance	Create an Auto Scaling group or a launch configuration using an EC2 instance. For information about creating a launch configuration using an EC2 instance, see Creating a Launch Configuration Using an EC2 Instance . For information about creating an Auto Scaling group using an EC2 instance, see Creating an Auto Scaling Group Using an EC2 Instance .	02 January 2014
Attach instances	Enable automatic scaling for an EC2 instance by attaching the instance to an existing Auto Scaling group. For more information, see Attach EC2 Instances to Your Auto Scaling Group .	02 January 2014
View account limits	View the limits on Auto Scaling resources for your account. For more information, see Auto Scaling Limits .	02 January 2014
Console support for Amazon EC2 Auto Scaling	Access Amazon EC2 Auto Scaling using the AWS Management Console. For more information, see Getting Started with Amazon EC2 Auto Scaling .	10 December 2013
Assign a public IP address	Assign a public IP address to an instance launched into a VPC. For more information, see Launching Auto Scaling Instances in a VPC .	19 September 2013
Instance termination policy	Specify an instance termination policy for Amazon EC2 Auto Scaling to use when terminating EC2 instances. For more information, see Controlling Which Auto Scaling Instances Terminate During Scale In .	17 September 2012
Support for IAM roles	Launch EC2 instances with an IAM instance profile. You can use this feature to assign IAM roles to your instances, allowing your applications to access other AWS services securely. For more information, see Launch Auto Scaling Instances with an IAM Role .	11 June 2012
Support for Spot Instances	Request Spot Instances in Auto Scaling groups by specifying a Spot Instance bid price in your launch configuration. For more information, see Launching Spot Instances in Your Auto Scaling Group .	7 June 2012
Tag groups and instances	Tag Auto Scaling groups and specify that the tag also applies to EC2 instances launched after the tag was created. For more information, see Tagging Auto Scaling Groups and Instances .	26 January 2012

Feature	Description	Release Date
Support for Amazon SNS	<p>Use Amazon SNS to receive notifications whenever Amazon EC2 Auto Scaling launches or terminates EC2 instances. For more information, see Getting SNS Notifications When Your Auto Scaling Group Scales.</p> <p>Amazon EC2 Auto Scaling also added the following new features:</p> <ul style="list-style-type: none"> • The ability to set up recurring scaling activities using cron syntax. For more information, see the PutScheduledUpdateGroupAction API command. • A new configuration setting that allows you to scale out without adding the launched instance to the load balancer (LoadBalancer). For more information, see the ProcessType API data type. • The <code>ForceDelete</code> flag in the <code>DeleteAutoScalingGroup</code> operation that tells Amazon EC2 Auto Scaling to delete the Auto Scaling group with the instances associated to it without waiting for the instances to be terminated first. For more information, see the DeleteAutoScalingGroup API operation. 	20 July 2011
Scheduled scaling actions	Added support for scheduled scaling actions. For more information, see Scheduled Scaling for Amazon EC2 Auto Scaling .	2 December 2010
Support for Amazon VPC	Added support for Amazon VPC. For more information, see Launching Auto Scaling Instances in a VPC .	2 December 2010
Support for HPC clusters	Added support for high performance computing (HPC) clusters.	2 December 2010
Support for health checks	Added support for using Elastic Load Balancing health checks with Amazon EC2 Auto Scaling-managed EC2 instances. For more information, see Using ELB Health Checks with Auto Scaling .	2 December 2010
Support for CloudWatch alarms	Removed the older trigger mechanism and redesigned Amazon EC2 Auto Scaling to use the CloudWatch alarm feature. For more information, see Dynamic Scaling for Amazon EC2 Auto Scaling .	2 December 2010
Suspend and resume scaling	Added support to suspend and resume scaling processes.	2 December 2010
Support for IAM	Added support for IAM. For more information, see Controlling Access to Your Amazon EC2 Auto Scaling Resources .	2 December 2010